

Do LLM Agents Have Regret?

A Case Study in Online Learning and Games

Chanwoo Park^{1,*}, Xiangyu Liu^{2,*}, Asuman Ozdaglar¹, and Kaiqing Zhang²

¹Massachusetts Institute of Technology

²University of Maryland, College Park

*Equal Contribution

Abstract

Large language models (LLMs) have been increasingly employed for (interactive) decision-making, via the development of LLM-based autonomous agents. Despite their emerging successes, the performance of LLM agents in decision-making has not been fully investigated through quantitative metrics, especially in the multi-agent setting when they interact with each other, a typical scenario in real-world LLM-agent applications. To better understand the limits of LLM agents in these interactive environments, we propose to study their interactions in benchmark decision-making settings of online learning and game theory, through the performance metric of *regret*. We first empirically study the no-regret behaviors of LLMs in canonical non-stochastic online learning problems, as well as the emergence of equilibria when multiple of them interact through playing repeated games. We then provide some theoretical insights into the sublinear regret growth in the cases we observed, under certain assumptions on (supervised) pre-training and the data generation model. Notably, we also identify (simple) cases where advanced LLMs such as GPT-4 fail to be no-regret. To further promote the no-regret behaviors, we propose a novel *unsupervised* training loss, the *regret-loss*, which, in contrast to the supervised pre-training loss, does *not* require the labels of (optimal) actions. Finally, we establish the *statistical* guarantee of generalization bound for regret-loss minimization, and more importantly, the *optimization* guarantee that minimizing such a loss can lead to known no-regret learning algorithms, when single-layer self-attention models are used. Our further experiments demonstrate the effectiveness of our regret-loss, especially in addressing the above “regrettable” cases.

1 Introduction

Large language models (LLMs) have recently exhibited remarkable emerging capabilities (Bubeck et al., 2023; Achiam et al., 2023; Wei et al., 2022b; Yao et al., 2023a). As a consequence, a burgeoning body of work has been investigating the employment of LLMs as central controllers for (interactive) decision-making, through the construction of *LLM-based autonomous agents* (Hao et al., 2023; Shen et al., 2023; Yao et al., 2023b; Shinn et al., 2023; Wang et al., 2023d; Significant Gravitass, 2023). Specifically, the LLM agent interacts with the (physical) world in a *dynamic/sequential* way: it uses LLMs as an oracle for reasoning and planning, then acts in the environment based on the reasoning/planning and the feedback it perceives over time. LLM agent has achieved impressive successes in embodied AI (Ahn et al., 2022; Huang et al., 2022a; Wang et al., 2023a), natural science (Wu et al., 2023; Swan et al., 2023), and social science (Park et al., 2022, 2023) applications.

Besides being *dynamic*, another increasingly captivating feature of LLM-based decision-making is the involvement of *strategic* interactions, oftentimes among multiple LLM agents. For example,

it has been reported that the reasoning capability of LLMs can be improved by interacting with each other through negotiation and/or debate games (Fu et al., 2023; Du et al., 2023); LLM agents have now been widely used to *simulate* the strategic behaviors for social and economic studies, to understand the emerging behaviors in interactive social systems (Aher et al., 2023; Park et al., 2023). Moreover, LLMs have also exhibited remarkable potential in solving various games (Bakhtin et al., 2022; Mukobi et al., 2023), and in fact, a rapidly expanding literature has employed *repeated games* as a fundamental benchmark to understand the strategic behaviors of LLMs (Brookins & DeBacker, 2023; Akata et al., 2023; Fan et al., 2023). These exciting empirical successes call for a rigorous examination and understanding through a theoretical lens of decision-making.

Regret, on the other hand, has been a core metric in (online) decision-making. It measures how “sorry” the decision-maker is, in retrospect, not to have followed the best prediction in hindsight (Shalev-Shwartz, 2012). It provides not only a sensible way to *evaluate* the sophistication level of on-line decision-makers, but also a quantitative way to measure their *robustness* against arbitrary (and possibly adversarial) environments. More importantly, it inherently offers a connection to modeling and analyzing *strategic behaviors*: the long-run interaction of no-regret learners leads to certain *equilibrium* when they repeatedly play games (Cesa-Bianchi & Lugosi, 2006). In fact, *no-regret* learning has served as a natural model for predicting and explaining human behaviors in strategic decision-making, with experimental evidence (Erev & Roth, 1998; Nekipelov et al., 2015; Balseiro & Gur, 2019). It has thus been posited as an important model of “rational behaviors” in playing games (Blum et al., 2008; Roughgarden, 2015). Hence, it is natural to ask:

Can we examine and better understand the online and strategic decision-making behaviors of LLMs through the lens of regret?

Acknowledging that LLM(-agents) are extremely complicated to analyze, to gain some insights into the question, we focus on benchmark decision-making settings: online learning with convex (linear) loss functions, and playing repeated games. We defer a detailed literature review to ??, and summarize our contributions as follows.

Contributions. First, we carefully examine the performance of several representative pre-trained LLMs in the aforementioned benchmark online decision-making settings, in terms of *regret*. We observe that LLM agents can achieve regret sublinear in time in (non-stochastic) online learning settings, where the loss functions change over time either arbitrarily, or by following some patterns with bounded variation, and in playing both representative and randomly generated repeated games. For the latter, equilibria will emerge as the long-term behavior of the multi-LLM interactions. Second, we provide some theoretical insights into the observed sublinear regret behaviors, based on certain assumptions on the *supervised pre-training* procedure, a common practice in training large models for decision-making, and some hypothetical models for training data generation. In particular, we make a connection of the pre-trained LLMs to the known no-regret algorithm of *follow-the-perturbed-leader* (FTPL) under these assumptions. Third, we also identify (simple) cases where advanced LLMs such as GPT-4 fail to be no-regret. We thus propose a novel *unsupervised* training loss, *regret-loss*, which, in contrast to the supervised pre-training loss, does not require the *labels* of (optimal) actions. We then establish both *statistical* and *optimization* guarantees for regret-loss minimization, which, in particular, show that minimizing such a loss can *automatically* lead to the known no-regret learning algorithm of *follow-the-regularized leader* (FTRL), under single-layer self-attention parameterization. Our further experiments demonstrate the effectiveness of our new loss, especially in addressing the above “regrettable” cases. With the fast development of LLMs, we emphasize that our goal is not to assert whether (current) LLMs are no-regret learners or not, especially given both the positive and negative observations above. Instead, our hope is to introduce and inspire more rigorous metrics and

principles into the current evaluation and development of LLM agents, for online and multi-agent strategic decision-making.

1.1 Related Work

LLM(-agent) for decision-making. The impressive capability of LLMs for *reasoning* (Bubeck et al., 2023; Achiam et al., 2023; Wei et al., 2022b,a; Srivastava et al., 2023; Yao et al., 2023a) has inspired a growing line of research on *LLM for (interactive) decision-making*, i.e., an LLM-based autonomous agent interacts with the environment by taking actions repeatedly/sequentially, based on the feedback it perceives. Some promises have been shown from a *planning* perspective (Hao et al., 2023; Valmeekam et al., 2023; Huang et al., 2022b; Shen et al., 2023). In particular, for embodied AI applications, e.g., robotics, LLMs have achieved impressive performance when used as the controller for decision-making (Ahn et al., 2022; Yao et al., 2023b; Shinn et al., 2023; Wang et al., 2023d; Driess et al., 2023; Significant Gravititas, 2023). However, the performance of decision-making has not been rigorously characterized via the regret metric in these works. Very recently, Liu et al. (2023e) has proposed a principled architecture for LLM-agent, with provable regret guarantees in stationary and stochastic decision-making environments, under the Bayesian adaptive Markov decision processes framework. In contrast, our work focuses on online learning and game-theoretic settings, in potentially adversarial and non-stationary environments. Moreover, (first part of) our work focuses on *evaluating* the intelligence level of LLM per se in decision-making (in terms of the regret metric), while Liu et al. (2023e) focused on *developing* a new architecture that uses LLM as an oracle for reasoning, together with memory and specific planning/acting subroutines, to *achieve* sublinear (Bayesian) regret, in stationary and stochastic environments.

LLMs in multi-agent environments. The interaction of multiple LLM agents has garnered significant attention lately. For example, Fu et al. (2023) showed that LLMs can autonomously improve each other in a negotiation game by playing and criticizing each other. Similarly, (Du et al., 2023; Liang et al., 2023; Xiong et al., 2023; Chan et al., 2024; Li et al., 2023c) showed that multi-LLM *debate* can improve the reasoning and evaluation capabilities of the LLMs. Qian et al. (2023); Schick et al. (2023); Wu et al. (2023) demonstrated the potential of multi-LLM interactions and collaboration in software development, writing, and problem-solving, respectively. Zhang et al. (2024) exhibited a similar potential in embodied cooperative environments. Recently, Park et al. (2025) trained multi-agent LLM systems using multi-agent reinforcement learning. More formally, multi-LLM interactions have also been investigated under a *game-theoretic* framework, to characterize the *strategic* decision-making of LLM agents. Bakhtin et al. (2022); Mukobi et al. (2023) and Xu et al. (2023b,a) have demonstrated the promise of LLMs in playing Diplomacy and Werewolf games, respectively, which are both language-based games with a mixture of competitive and cooperative agents. Note that these works utilized LLM to solve a specific rather than a general game. Related to our work, Brookins & DeBacker (2023); Akata et al. (2023); Lorè & Heydari (2023); Brookins & DeBacker (2023); Fan et al. (2023) have also used (repeated) matrix games as a benchmark to evaluate the reasoning capability and rationality of LLM agents. In contrast to our work, these empirical studies have not formally investigated LLM agents using the metric of *regret*, nor through the lenses of *online learning* and *equilibrium-computation*, which are all fundamental in modeling and analyzing strategic multi-agent interactions. Moreover, our work also provides theoretical results to explain and further enhance the no-regret property of LLM agents.

LLMs & Human/Social behavior. LLMs have also been used to *simulate* the behavior of human beings, for social science and economics studies (Engel et al., 2023). The extent of LLMs simulat-

ing human behavior has been claimed as a way to evaluate the level of its intelligence in a controlled environment (Aher et al., 2023; Tsai et al., 2023). For example, Li et al. (2023b); Hong et al. (2024); Zhao et al. (2023) showed that by specifying different “roles” to LLM agents, certain collaborative/competitive behaviors can emerge. Argyle et al. (2023) showed that LLMs can emulate response distributions from diverse human subgroups, illustrating their adaptability. Horton (2023) argued that an LLM, as a computational model of humans, can be used as *homo economicus* when given endowments, information, preferences, etc., to gain new economic insights by simulating its interaction with other LLMs. Park et al. (2022, 2023) proposed scalable simulators that can generate realistic social behaviors emerging in populated and interactive social systems, and the emerging behaviors of LLM agents in society have also been consistently observed in Chen et al. (2024, 2023). Li et al. (2023d,a) studied the behavioral dynamics of LLM agents on social networks. These empirical results have inspired our work, which can be viewed as an initial attempt towards quantitatively understanding the *emerging behavior* of LLMs as computational human models, given the known justification of *equilibrium* being a long-run emerging behavior of *learning dynamics* (Fudenberg & Levine, 1998) and strategic interactions (Young, 2004; Camerer, 2011).

Transformers & In-context-learning. LLMs nowadays are predominantly built upon the architecture of Transformers (Vaswani et al., 2017). Transformers have exhibited a remarkable capacity of *in-context-learning* (ICL), which can construct new predictors from sequences of labeled examples as input, without further parameter updates. This has enabled the *few-shot learning* capability of Transformers (Brown et al., 2020; Garg et al., 2022; Min et al., 2022). The empirical successes have inspired burgeoning theoretical studies on ICL. Xie et al. (2022) used a Bayesian inference framework to explain how ICL works, which has also been adopted in Wang et al. (2023b); Jiang (2023). Akyürek et al. (2023); Von Oswald et al. (2023); Dai et al. (2023); Giannou et al. (2023) showed (among other results) that ICL comes from that Transformers can implement the gradient descent (GD) algorithm. Bai et al. (2023) further established that Transformers can implement a broad class of machine learning algorithms in context. Moreover, Ahn et al. (2023); Zhang et al. (2023a); Mahankali et al. (2023) proved that a *minimizer* of the certain training loss among single-layer Transformers is equivalent to a single step of GD for linear regression. Li et al. (2023e) established generalization bounds of ICL from a multi-task learning perspective. Zhang et al. (2023b) argued that ICL implicitly implements Bayesian model averaging, and can be approximated by the attention mechanism. They also established a result on some *regret* metric. However, the regret notion is not defined for (online) decision-making, and is fundamentally different from ours that is standard in online learning and games. Also, we provide extensive experiments to validate the no-regret behavior by our definition. More recently, the ICL property has also been generalized to decision-making settings. Laskin et al. (2023); Lee et al. (2023); Lin et al. (2024) investigated the in-context reinforcement learning (RL) property of Transformers under supervised pre-training, for solving stochastic bandits and Markov decision processes. In contrast, our work focuses on online learning settings with an arbitrary and *potentially adversarial* nature, as well as *game-theoretic* settings. We also provide a new *unsupervised* loss to promote the no-regret behavior in our settings.

Online learning and games. Online learning has been extensively studied to model the decision-making of an agent who interacts with the environment sequentially, with a potentially arbitrary sequence of loss functions (Shalev-Shwartz, 2012; Hazan, 2016), and has a deep connection to game theory (Cesa-Bianchi & Lugosi, 2006). In particular, regret, the difference between the incurred accumulated loss and the best-in-hindsight accumulated loss, has been the core performance metric, and a good online learning algorithm should have regret at most sublinear in time T , which is referred to as being *no-regret*. Many well-known algorithms can achieve no-regret against *arbitrary* loss

sequences, e.g., multiplicative weight updates (MWU)/Hedge (Freund & Schapire, 1997; Arora et al., 2012b), EXP3 (Auer et al., 2002), and more generally follow-the-regularized-leader (FTRL) (Shalev-Shwartz & Singer, 2007) and follow-the-perturbed-leader (FTPL) (Kalai & Vempala, 2005). In the bandit literature (Lattimore & Szepesvári, 2020; Bubeck et al., 2012), such a setting without any statistical assumptions on the losses is also referred to as the *adversarial/non-stochastic* setting. Following the conventions in this literature, the online settings we focus on shall not be confused with the stationary and *stochastic*(-bandit)/(-reinforcement learning) settings that have been explored in several other recent works on *Transformers for decision-making* (Lee et al., 2023; Lin et al., 2024). Centering around the regret metric, our work has also explored the non-stationary bandit setting (Besbes et al., 2014), as well as the repeated game setting where the environment itself consists of strategic agents (Cesa-Bianchi & Lugosi, 2006).

2 Preliminaries

Notation. We use \mathbb{N} and \mathbb{N}^+ to denote the sets of non-negative and positive integers, respectively. For a finite set \mathcal{S} , we use $\Delta(\mathcal{S})$ to denote the simplex over \mathcal{S} . For $d \in \mathbb{N}^+$, we define $[d] := \{1, 2, \dots, d\}$. For two vectors $x, y \in \mathbb{R}^d$, we use $\langle x, y \rangle$ to denote the inner product of x and y . We define $\mathbf{0}_d$ and $\mathbf{1}_d$ as a d -dimensional zero or one vector, and $\mathbf{O}_{d \times d}$ and $I_{d \times d}$ as a $d \times d$ -dimensional zero matrix and identity matrix, respectively. We omit d when it is clear from the context. We define e_i as a unit vector (with proper dimension) whose i -th coordinate equal to 1. For $p \in \mathbb{R}^d, R > 0$ and $C \subseteq \mathbb{R}^d$ is a convex set, define $B(p, R, \|\cdot\|) := \{x \in \mathbb{R}^d \mid \|x - p\| \leq R\}$, $\text{Proj}_{C, \|\cdot\|}(p) = \arg \min_{x \in C} \|x - p\|$ (which is well defined as C is a convex set), and $\text{clip}_R(x) := [\text{Proj}_{B(0, R, \|\cdot\|_2), \|\cdot\|_2}(x_i)]_{i \in [d]}$. Define $\text{Softmax}(x) := \left(\frac{e^{x_i}}{\sum_{i \in [d]} e^{x_i}} \right)_{i \in [d]}$ and $\text{ReLU}(x) = \max(0, x)$ for $x \in \mathbb{R}^d$. For $A \in \mathbb{R}^{m \times n}$ with A_i denoting its i -th column, we define $\|A\|_{\text{op}} := \max_{\|x\|_2 \leq 1} \|Ax\|_2$, $\|A\|_{2, \infty} := \sup_{i \in [n]} \|A_i\|_2$, $\|A\|_F$ as the Frobenius norm, and $A_{-1} := A_n$ to denote the last column vector of A . We define $\mathbb{R}^+ := \{x \mid x \geq 0\}$. For a set Π , define $\text{diam}(\Pi, \|\cdot\|) := \sup_{\pi_1, \pi_2 \in \Pi} \|\pi_1 - \pi_2\|$. We define $\mathbb{1}(\mathcal{E}) := 1$ if \mathcal{E} is true, and $\mathbb{1}(\mathcal{E}) := 0$ otherwise. For a random variable sequence $(X_n)_{n \in \mathbb{N}}$ and random variables X, Y , we denote F_X as the cumulative distribution function of a random variable X , $X_n \xrightarrow{p} X$ if $\forall \epsilon > 0, \lim_{n \rightarrow \infty} \mathbb{P}(|X_n - X| > \epsilon) = 0$, $X_n \xrightarrow{d} X$ if $\lim_{n \rightarrow \infty} F_{X_n}(x) = F_X(x)$ for all x where $F_X(x)$ is continuous, $X \stackrel{d}{=} Y$ if $F_X(x) = F_Y(x)$ for all x , $X_n \xrightarrow{a.s.} X$ if $\mathbb{P}(\lim_{n \rightarrow \infty} X_n = X) = 1$, and $\text{esssup}(X) := \inf\{M \in \mathbb{R} : \mathbb{P}(X > M) = 0\}$. For a random variable X , we use $\text{supp}(X)$ to denote its support. For functions $f, g : \mathbb{R} \rightarrow \mathbb{R}$, we define $g(x) = \mathcal{O}(f(x))$ if there exist $x_0, M < \infty$ such that $|g(x)| \leq M|f(x)|$ for all $x > x_0$. We use f' to denote the derivative of f . Let $F : \Omega \rightarrow \mathbb{R}$ be a continuously-differentiable, strictly convex function defined on a convex set Ω . The Bregman divergence associated with F for points p, q is defined as $D_F(p, q) := F(p) - F(q) - \langle \nabla F(q), p - q \rangle$. For a sequence $(\ell_t)_{t \in [T]}$ for some $T \in \mathbb{N}^+$, we define $\ell_{a:b} := (\ell_a, \dots, \ell_b)$ for $1 \leq a \leq b \leq T$. If $a > b$, we define $\ell_{a:b} = \emptyset$.

2.1 Online Learning & Repeated Games

Online learning. We consider the online learning setting where an agent interacts with the environment for T rounds, by iteratively making decisions based on the feedback she receives. Specifically, at each time step t , the agent chooses her decision policy $\pi_t \in \Pi$ for some bounded domain Π , and after her commitment to π_t , a bounded loss function $f_t : \Pi \rightarrow [-B, B]$ for some constant $B > 0$ is chosen by the environment, potentially in an adversarial fashion. The agent thus incurs a loss of $f_t(\pi_t)$, and will update her decision to π_{t+1} using the feedback. We focus on the most basic setting where the agent chooses actions from a finite set \mathcal{A} every round, which is also referred to as the *Experts Problem* (Cover, 1966; Vovk, 1990; Littlestone & Warmuth, 1994; Hazan, 2016), without loss

of much generality (c.f. Appendix B.3 for a detailed discussion). In this case, Π becomes the simplex over \mathcal{A} , i.e., $\Pi = \Delta(\mathcal{A})$, and $f_t(\pi_t) = \langle \ell_t, \pi_t \rangle$ for some loss vector $\ell_t \in \mathbb{R}^d$ that may change over time, where $d := |\mathcal{A}|$.

At time step $t \in [T]$, the agent may receive either the full vector ℓ_t , or only the realized loss ℓ_{ta_t} (we sometimes also interchangeably write it as $\ell_t(a_t)$), the a_t th element of ℓ_t , for some $a_t \sim \pi_t(\cdot)$, as feedback, which will be referred to as online learning with *full-information feedback*, and that with *bandit feedback*, respectively. The latter is also referred to as the *adversarial/non-stochastic bandit* problem in the multi-armed bandit (MAB) literature. Note that hereafter, we will by default refer to this setting that does *not* make any assumptions on the loss sequence $(\ell_t)_{t \in [T]}$ simply as *online learning*. Moreover, if the loss functions change over time (usually with certain bounded variation), we will refer to it as *non-stationary online learning* for short, whose bandit-feedback version is also referred to as the *non-stationary bandit* problem.

Repeated games. The online learning setting above has an intimate connection to game theory. Consider a normal-form game $\mathcal{G} = \langle N, \{\mathcal{A}_n\}_{n \in [N]}, \{r_n\}_{n \in [N]} \rangle$, where N is the number of players, \mathcal{A}_n and $r_n : \mathcal{A}_1 \times \dots \times \mathcal{A}_N \rightarrow [-B, B]$ are the action set and the payoff function of player n , respectively. The N players repeatedly play the game for T rounds, each player n maintains a strategy $\pi_{n,t} \in \Delta(\mathcal{A}_n)$ at time t , and takes action $a_{n,t} \sim \pi_{n,t}(\cdot)$. The *joint* action $a_t = (a_{1,t}, \dots, a_{N,t})$ determines the payoff of each player at time t , $\{r_n(a_t)\}_{n \in [N]}$. From a single-player's (e.g., player n 's) perspective, she encounters an online learning problem with (expected) loss function $\ell_t := -\mathbb{E}_{a_{-n,t} \sim \pi_{-n,t}}[r_n(\cdot, a_{-n,t})]$ at time t , where $-n$ denotes the index for all the players other than player n . We will refer to it as the *game setting* for short, and use the terms of “agent” and “player” interchangeably hereafter. The key difference between online learning and repeated games is in their interaction dynamics: online learning involves an agent facing a potentially adversarial, changing environment (or sequence of loss functions), while in repeated games, agents interact by playing the same game repeatedly, which might be less adversarial when they follow specific learning algorithms.

2.2 Self-attention & Transformers

(Linear) Self-attention. One key component in Transformers (Vaswani et al., 2017), the backbone of modern language models, is the (*self*-)attention mechanism. For simplicity, we here focus on introducing the *single-layer* self-attention architecture. The mechanism takes a sequence of vectors $Z = [z_1, \dots, z_t] \in \mathbb{R}^{d \times t}$ as input, and outputs some sequence of $[\widehat{z}_1, \dots, \widehat{z}_t] \in \mathbb{R}^{d \times t}$. For each $i \in [t]$ where $i > 1$, the output is generated by $\widehat{z}_i = (Vz_{1:i-1})\sigma((Kz_{1:i-1})^\top(Qz_i))$, where $z_{1:i-1}$ denotes the 1 to $i-1$ columns of Z , σ is either the Softmax or ReLU activation function, and for the initial output, $\widehat{z}_1 = \mathbf{0}_d$. Here, $V, Q, K \in \mathbb{R}^{d \times d}$ are referred to as the *Value*, *Query*, and *Key* matrices, respectively. Following the theoretical framework in Von Oswald et al. (2023); Mahankali et al. (2023), we exclude the attention score for a token z_i in relation to itself. For theoretical analysis, we also consider the *linear* self-attention model, where $\widehat{z}_i = (Vz_{1:i-1})((Kz_{1:i-1})^\top(Qz_i))$. We write this (linear) self-attention layer's output as $(L)SA_{(V,Q,K)}(Z)$. We define an M -head self-attention layer with $\theta = \{(V_m, Q_m, K_m)\}_{m \in [M]}$ as $M-(L)SA_\theta(Z) := \sum_{m=1}^M (L)SA_{(V_m, Q_m, K_m)}(Z)$. We define $\|\cdot\|_{M-(L)SA}$ as $\|\theta\|_{M-(L)SA} := \max_{m \in [M]} \{\|Q_m\|_{\text{op}}, \|K_m\|_{\text{op}}\} + \sum_{m=1}^M \|V_m\|_{\text{op}}$.

Transformers. For a multi-layer perceptron (MLP) layer, it takes $Z = [z_1, \dots, z_t] \in \mathbb{R}^{d \times t}$ as input, with parameter $\theta = (W_1, W_2) \in \mathbb{R}^{d' \times d} \times \mathbb{R}^{d \times d'}$ such that for each $i \in [t]$, the output is $\widehat{z}_i := W_2 \sigma(W_1 z_i)$ where σ is either Softmax or ReLU. We write the output of an MLP layer with parameter θ as $\text{MLP}_\theta(Z)$. Defining $\|\cdot\|_{\text{MLP}}$ as $\|\theta\|_{\text{MLP}} := \|W_1\|_{\text{op}} + \|W_2\|_{\text{op}}$ and $\text{ResNet}(f, Z) := Z + f(Z)$, we can define an L -layer

Transformer with parameter $\theta = (\theta^{(lm)}, \theta^{(la)})_{l \in [L]}$ as

$$\text{TF}_\theta(Z) := Z^{(L)},$$

where the output $Z^{(L)}$ is defined iteratively from $Z^{(0)} = \text{clip}_R(Z) := \min(-R, \max(R, Z))$ and

$$Z^{(l)} = \text{clip}_R\left(\text{ResNet}\left(\text{MLP}_{\theta^{(la)}}, \text{ResNet}\left(\text{M-(L)SA}_{\theta^{(lm)}}, Z^{(l-1)}\right)\right)\right),$$

for some $R > 0$. We define a class of Transformers with certain parameters as $\Theta_{d,L,M,d',B_{\text{TF}}} := \{\theta = (\theta^{(lm)}, \theta^{(la)})_{l \in [L], m \in [M]} : \|\theta\|_{\text{TF}} \leq B_{\text{TF}}\}$, where M is the number of heads of self-attention,

$$\|\theta\|_{\text{TF}} := \max_{l \in [L]} \left\{ \|\theta^{(la)}\|_{\text{M-(L)SA}} + \|\theta^{(lm)}\|_{\text{MLP}} \right\}, \quad (2.1)$$

and $B_{\text{TF}} > 0$ is some constant. When it is clear from the context, we may omit the subscripts and write it as Θ for simplicity. We assume R to be sufficiently large such that clip does not take effect on any of our approximation results.

2.3 Performance Metric: Regret

We now introduce *regret*, the core performance metric used in online learning and games. For a given algorithm \mathcal{A} , let $\pi_{\mathcal{A},t}$ denote the decision policy of the agent at time t generated by \mathcal{A} . Then, the regret, which is the difference between the accumulated (expected) loss incurred by implementing \mathcal{A} and that incurred by the best-in-hindsight fixed decision, can be defined as

$$\text{Regret}_{\mathcal{A}}((f_t)_{t \in [T]}) := \sum_{t=1}^T f_t(\pi_{\mathcal{A},t}) - \inf_{\pi \in \Pi} \sum_{t=1}^T f_t(\pi).$$

In the Experts Problem, the definition is instantiated as $\text{Regret}_{\mathcal{A}}((\ell_t)_{t \in [T]}) := \sum_{t=1}^T \langle \ell_t, \pi_{\mathcal{A},t} \rangle - \inf_{\pi \in \Pi} \sum_{t=1}^T \langle \ell_t, \pi \rangle$. With bandit-feedback, a common metric may also take further expectation for $\text{Regret}_{\mathcal{A}}$, over the randomness of the policies $(\pi_{\mathcal{A},t})_{t \in [T]}$. An algorithm \mathcal{A} is referred to as being *no-regret*, if $\max_{(f_t)_{t \in [T]}} \text{Regret}_{\mathcal{A}}((f_t)_{t \in [T]}) \sim o(T)$, i.e., the worse-case regret grows sublinearly in T . Known no-regret algorithms include follow-the-regularized-leader (Shalev-Shwartz & Singer, 2007), follow-the-perturbed-leader (Kalai & Vempala, 2005) (see Appendix B.2 for more details).

In non-stationary online learning, one may also use the metric of *dynamic regret* (Zinkevich, 2003), where the *comparator* in the definition also changes over time, as the best decision policy at each time t : $\text{D-Regret}_{\mathcal{A}}((f_t)_{t \in [T]}) := \sum_{t=1}^T f_t(\pi_{\mathcal{A},t}) - \sum_{t=1}^T \inf_{\pi \in \Pi} f_t(\pi)$, which is a stronger notion than $\text{Regret}_{\mathcal{A}}((f_t)_{t \in [T]})$ in that $\text{Regret}_{\mathcal{A}}((f_t)_{t \in [T]}) \leq \text{D-Regret}_{\mathcal{A}}((f_t)_{t \in [T]})$.

3 Do Pre-Trained LLMs Have Regret? Experimental Validation

In this section, we explore the no-regret behaviors of representative LLMs (i.e., mainly GPT-4 Turbo and GPT-4, together with GPT-3.5 Turbo, Mixtral-8x7b-instruct, and Llama-3-70B-instruct), in the context of online learning and games. All experiments with LLMs are conducted using the public OpenAI (Openai, 2023) or LLM Engine (LLM Engine, 2023) Python API. We provide some hypothetical intuitions as to why pre-trained LLM might be no-regret in Appendix C.1, which will be made concrete next.

Interaction protocol. To enable the sequential interaction with LLMs, we first describe the setup and objective of our experimental study. At each round, we incorporate the entire history of loss vectors of past interactions into our prompts, as concatenated texts, and ask the LLM agent to determine a policy that guides the decision-making for the next round. Note that since we hope to *evaluate* the sophistication level of pre-trained LLMs through online learning or games, we only provide simple prompts that she should utilize the history information, without providing explicit rules of *how* to make use of the history information, nor asking her to *minimize regret* (in any sense). A detailed description and an ablation study of the prompts are deferred to Appendix C.7, and an illustration of the protocol for playing repeated games is given in Figure 3.1.

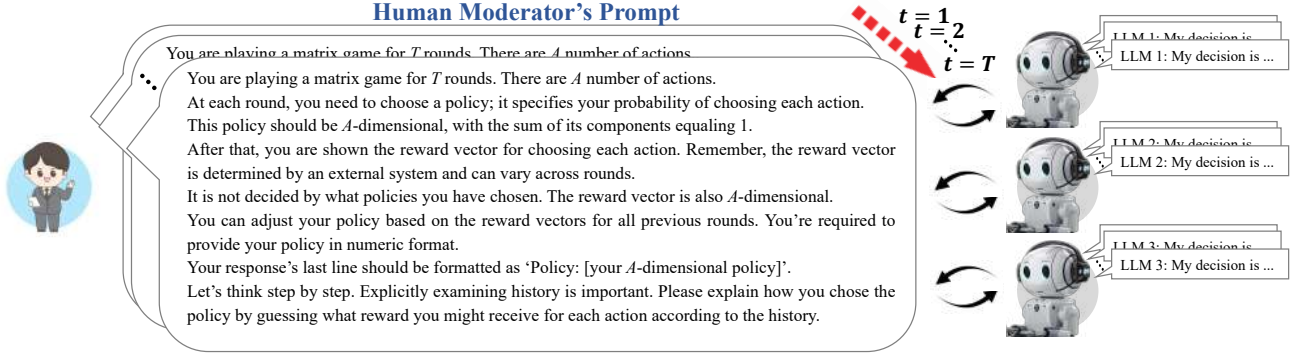


Figure 3.1: Demonstration of the prompts and interaction protocol for multi-player repeated games. A human moderator does not provide the game’s payoff matrices to the LLMs. Instead, at each round, the human moderator provides each player’s own payoff vector history.

3.1 Framework for Sublinear Regret Behavior Validation

Before delving into the results, we note that to the best of our knowledge, we are not aware of any principled framework for validating sublinear growth of the regret with *finite-time* experimental data. Therefore, we propose two frameworks below to rigorously validate the no-regret behaviors of algorithms over a *finite* T , which might be of independent interest. More details can be found in Appendix C.2.

Trend-checking framework. We propose a statistical hypothesis test aligned with our objectives:

H_0 : The sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t \in [T]}$ does not exhibit a decreasing pattern

H_1 : The sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t \in [T]}$ shows a decreasing pattern.

Ideally, one should check if $\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t$ approaches zero (or a negative value) as t goes to infinity. With a finite T value, testing these hypotheses provides a method to quantify this – whether we reject H_0 offers a way to measure it. To this end, one needs to count the number of $\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t - \text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t+1]}\right)/(t+1) > 0$, for which we use Proposition 1 below. We will report the p -value of H_0 , denoted as p_{trend} , as the output of this framework.

Proposition 1. (p -value of the null hypothesis). *Define the event*

$$\mathcal{E}(s, T) := \left\{ \text{The number of } \frac{\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)}{t} - \frac{\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t+1]}\right)}{t+1} > 0 \text{ for } t = 1, \dots, T \text{ is at least } s \geq \frac{T-1}{2} \right\}.$$

Under the assumption that the null hypothesis H_0 holds, the probability of this event happening is bounded as $\mathbb{P}_{H_0}(\mathcal{E}(s, T)) \leq \frac{1}{2^{T-1}} \sum_{t=s}^{T-1} \binom{T-1}{t}$.

Regression-based framework. We propose an alternative approach by fitting the data with regression. In particular, one can use the data $\left\{ \left(t, \log \text{Regret}_{\mathcal{A}} \left((f_{\tau})_{\tau \in [t]} \right) \right) \right\}_{t \in [T]}$ to fit a function $g(t) = \beta_0 \log t + \beta_1$, where the estimate of β_0 , i.e., $\widehat{\beta}_0$, satisfying $\widehat{\beta}_0 < 1$ may be used to indicate the no-regret behavior, i.e., the *sublinear* growth of $\text{Regret}_{\mathcal{A}} \left((f_{\tau})_{\tau \in [t]} \right)$ over time. While being simple, it cannot be directly used when $\text{Regret}_{\mathcal{A}} \left((f_{\tau})_{\tau \in [t]} \right) < 0$. Hence, we set $\log \text{Regret}_{\mathcal{A}} \left((f_{\tau})_{\tau \in [t]} \right)$ as -10 if this happens. We define p_{reg} as the p -value of the regression parameter $\widehat{\beta}_0$, and will report the pair of $(\widehat{\beta}_0, p_{reg})$ as the output of this framework.

3.2 Results: Online Learning

We now present the experimental results of pre-trained LLMs in online learning in: 1) (arbitrarily) changing environments, 2) non-stationary environments, and 3) bandit-feedback environments. Results for 2) and 3) are deferred to Appendices C.3 and C.4.

Changing environments. We first consider the setting with (arbitrarily) changing environments, which are instantiated as follows: 1) *Randomly-generated loss sequences.* At every timestep, we generate a random loss vector $\ell_t \sim \text{Unif}(\times_{i=1}^d [\min\{x_i, y_i\}, \max\{x_i, y_i\}])$ for $\{x_i, y_i \sim \text{Unif}(0, 10)\}_{i \in [d]}$ or $\ell_t \sim \mathcal{N}(\boldsymbol{\mu}_d, I)$ with clipping to $[0, 10]$ to ensure boundedness of the loss, where $\boldsymbol{\mu}_d \sim \text{Unif}([0, 10]^d)$. Note that we use this as a way to *systematically* generate potentially arbitrary loss sequences, and also note that our regret was defined for each *realization* of the *random loss vectors* (instead of their expectations as in the definition of regret in *stochastic bandit* problems), which can be arbitrarily different across timesteps. 2) *Loss sequences with certain trends.* Although many real-world environments may change, they often change by following certain patterns. Therefore, we consider two representative trends, the *linear* trend and the *periodic* (sinusoid) trend. We sample $a, b \sim \text{Unif}([0, 10]^d)$ and let $\ell_t = (b-a)\frac{t}{T} + a$ for the linear trend and $\ell_t = 5(1 + \sin(at+b))$ for the periodic trend. In the experiments, we choose $d = 2$. The average regret (over multiple randomly generated instances) performance is presented in Figure 3.2¹, where we compare GPT-4 with well-known no-regret algorithms, FTRL with entropy regularization and FTPL with Gaussian perturbations (with tuned parameters). It is seen that these pre-trained LLMs can achieve sublinear regret in a large portion of the instances, and have sometimes even lower regret values than baselines.

Behavioral patterns of LLMs. To understand how LLMs make decisions at each time step, we provided example outputs of LLMs *reasoning* how they generate their policies in Appendix C.9. We find that LLMs tend to use the history of the reward vectors by looking at their *sum/average*, and tend to introduce *randomization* in decision-making. These are known to be the keys to achieving no-regret behaviors in online learning (Hazan, 2016; Cesa-Bianchi & Lugosi, 2006).

Longer-horizon results. We also test the robustness and scalability of our empirical findings in more challenging environments. We extend the problem horizon to $T = 100$ for the two settings where loss vectors are generated in a stationary way (i.e., *Uniform* and *Gaussian*), and $T = 200$ for the

¹We emphasize that the error bars in the figures are *not* associated with the randomness/variance of the algorithms/LLM-agents, but with the randomness/variance of the generation of environment instances.




Figure 3.2: Regret of pre-trained LLMs for online learning with full-information feedback. Notably, both commercial and open-source LLMs can achieve sublinear regret as validated by our frameworks and the comparison with FTRL/FTPL, though the performances of weaker models of GPT-3.5 and open-source ones are worse. Interestingly, the GPT-4 model can even outperform well-known no-regret learning algorithms, FTRL and FTPL.

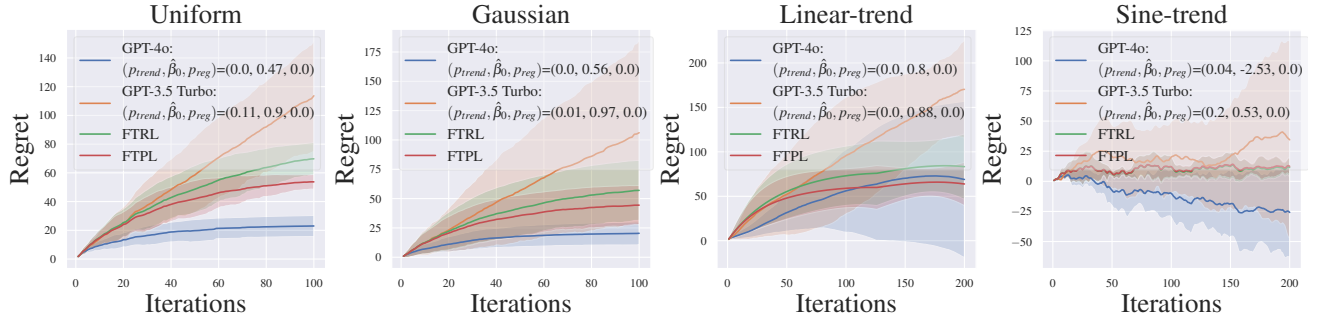


Figure 3.3: Regret of pre-trained LLMs for online learning with full-information feedback, with longer horizons of $T = 100$ and $T = 200$. In most cases, the LLMs can achieve sublinear regret as validated by our frameworks and the comparison with FTRL/FTPL, though the performances of the weaker model of GPT-3.5 is worse.

other two non-stationary settings (i.e., *Linear-trend* and *Sine-trend*). Note that since in each round, we need to feed all the previous history to the LLMs, the API costs in fact scale *quadratically* with respect to the horizon T . Therefore, we replace GPT-4 by its cheaper (and more recent) version of GPT-4o. To further scale to even longer-horizon cases with $T = 500$, we *summarize* the history to reduce the prompt length by providing LLMs with the summation of the history loss associated with each action. Similar summary-based input was also used in the concurrent work [Krishnamurthy et al. \(2024\)](#), where both the *averaged reward* and the *action selection count* of each action were summarized for the (i.i.d.) stochastic bandit setting. The corresponding results are provided in Figure 3.3 and Table 1, where the LLMs can exhibit no-regret behaviors as validated by our frameworks and the comparison with FTRL/FTPL.

3.3 Results: Multi-Player Repeated Games

We now consider the setting when multiple LLMs make online decisions in a *shared* environment repeatedly. Specifically, at each round, the loss vectors each agent receives are determined by both her payoff matrix and the strategies of all other agents. Note that the payoff matrix is not directly revealed to the LLM agent, but she has to make decisions in a completely online fashion based on the payoff vector marginalized by the opponents' strategies (see Figure 3.1 for an example of the

$(p_{trend}, \hat{\beta}_0, p_{reg})$	GPT-4o	FTRL	FTPL
Uniform	(0.0, 0.85, 0.0)	(0.0, 0.6, 0.0)	(0.0, 0.52, 0.0)
Gaussian	(0.0, 0.86, 0.0)	(0.0, 0.64, 0.0)	(0.0, 0.68, 0.0)
Linear-trend	(0.02, 0.83, 0.5)	(0.02, 0.76, 0.1)	(0.01, 0.79, 0.0)
Sine-trend	(0.09, 0.28, 0.0)	(0.01, 0.24, 0.0)	(0.01, 0.26, 0.0)

Table 1: Longer-horizon ($T = 500$). GPT-4o model can still exhibit sublinear regret behaviors as validated by our frameworks and the comparison with FTRL/FTPL.

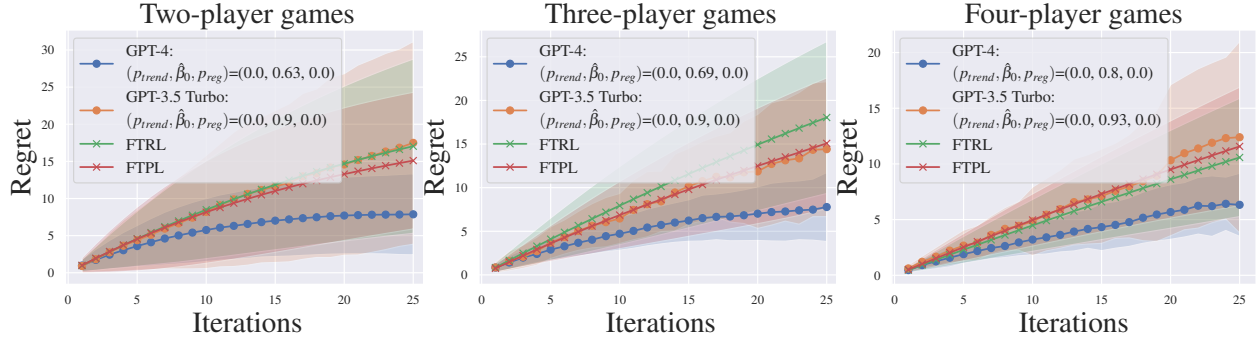


Figure 3.4: Regret of pre-trained LLMs for repeated games of different sizes, in most cases, both commercial and open-source LLMs can achieve sublinear regret as validated by our frameworks and the comparison with FTRL/FTPL. We report the regret of one agent for ease of presentation.

prompt). This is a typical scenario in learning in (repeated) games (Cesa-Bianchi & Lugosi, 2006).

Representative games. We first test LLMs on 6 representative general-sum games (*win-win*, *prisoner’s dilemma*, *unfair*, *cyclic*, *biased*, and *second best*) studied in Robinson & Goforth (2005) (c.f. Appendix B.4). For each type of the game, we conduct 20 repeated experiments.

Randomly generated games. To further validate the no-regret behaviors of LLMs, we also test on 50 randomly generated three-player general-sum games, and 50 randomly generated four-player general-sum games, where each entry of the payoff matrix is sampled randomly from $\text{Unif}([0, 10])$. These are larger and more challenging settings than the structured and representative ones above.

We summarize the experimental results in Figure 3.4, which are similar to the above in the online setting: for all types of games, pre-trained LLMs can achieve sublinear regret, which is often lower than that obtained by FTRL/FTPL for most games. We provide six instances of three-player general-sum games and six instances of four-player general-sum games in Figure C.3 and Figure C.4, respectively. Occasionally, GPT-4 even provides a negative regret value.

3.4 Pre-Trained LLM Agents Can Still Have Regret

The experiments above may suggest the no-regret behaviors of LLMs in online learning and game playing. However, is this capability *universal*? We show that the no-regret property can break for LLM agents if the loss vectors are generated in a more adversarial way.

Canonical counterexamples for follow-the-leader. First, we consider two well-known examples that the *follow-the-leader* (FTL) algorithm (Shalev-Shwartz, 2012) suffers from *linear regret*.




Figure 3.5: (left) Regret of GPT-4 (Turbo) under the canonical counterexample for FTL (Hazan, 2016, Chapter 5). (mid, right) Failure of GPT-4 (Turbo) on two scenarios with regrettable behaviors, while Transformers trained by our new regret-loss ($N = 1$) in Section 5 can achieve sublinear regret.

Example 1: $\ell_1(1) = 5, \ell_1(2) = 0$ and $\ell_t(2 - t\%2) = 10, \ell_t(1 + t\%2) = 0$ for $t \geq 2$ (Hazan, 2016).

Example 2: $\ell_t(2 - t\%2) = 10, \ell_t(1 + t\%2) = 0$ for $1 \leq t \leq c$ and $\ell_t(1) = 10, \ell_t(2) = 0$ for $c+1 \leq t \leq T (= 500)$, for some integer c satisfying $0 < c < T$ (Feder et al., 1992).

Here, % denotes the modulo operation. Interestingly, for **Example 1**, GPT-4 agent can easily identify the pattern for the loss sequence that the optimal action *alternates*, thus accurately predicting the loss it will receive and achieving low regret in Figure 3.5. For **Example 2**, the GPT-4 agent with *raw history* input also provides an impressively lower (negative) regret than FTRL and FTPL (Figure C.5). The GPT-4 agent with *summarized history* input, in contrast, suffers from much larger regret than FTRL and FTPL. We defer the detailed comparison between using raw history and summarized history to Figure C.5, and an explanation of LLMs’ behaviors via predicting the *trend* of the loss instances to Appendix C.6. In summary, the GPT-4 agent may predict such worst-case sequences well, and does not fail in the same way as FTL, which is known to suffer from a lack of randomness in decisions.

Additionally, the results on **Example 2** also imply that summary-based history input can perform worse than the raw-history-based one in the adversarial setting we consider, while the former was claimed to be the key in succeeding in the i.i.d. stochastic bandit setting (Krishnamurthy et al., 2024). The regret values with these two types of input differ significantly, with a p -value of 1.2×10^{-157} under a one-sided independent t-test. These results further illustrate the fundamental differences between the settings considered in Krishnamurthy et al. (2024) and ours.

Noisy alternating loss sequence. Inspired by the above, we design a new loss sequence that is *similar but less predictable*, by adding some noise to the canonical counterexample. Specifically, we construct the following (simple) loss sequence with 2 actions such that $\ell_t(1 + t\%2) = \min(25/t, 10), \ell_t(2 - t\%2) \sim \text{Unif}([9, 10])$ for $t \in [25]$.

Adaptive loss sequence. We also develop a simpler but more *adaptive* loss sequence that takes the full power of the adversary in our online learning setup. After the GPT-4 agent provides π_t , we choose ℓ_t with $\ell_t(\arg \max_i \pi_{ti}) = 10$ and $\ell_t(3 - \arg \max_i \pi_{ti}) = 0$.

We also report the average regret over 20 repeated experiments for the later two settings using GPT-4 and more advanced GPT-4 Turbo in Figure 3.5, where we cannot reject the hypothesis that GPT-4 (Turbo) has linear regret by either our trend-checking or regression-based framework. These observations have thus motivated us to design new approaches to further promote the no-regret behaviors of the models, with additional training, as to be detailed in Section 5. Before it, we first provide some theoretical insights into the observed sublinear regret behaviors.

4 Why Do Pre-Trained LLMs (Not) Have Regret? A Hypothetical Model and Some Theoretical Insights

We now provide some plausible explanations about the observed no-regret behaviors of pre-trained LLMs, which are highly *hypothetical* by nature, since to the best of our knowledge, the details of pre-training these popular LLMs (e.g., GPT-3.5 Turbo and GPT-4), regarding data distribution, training algorithm, etc., have not been revealed. We instead make the explanations based on some existing assumptions in the literature for modeling human behaviors, and the recent literature on understanding LLMs and Transformers.

4.1 A (Human) Decision-Making Model: Quantal Response

A seminal model for human decision-making behaviors is the *quantal response* model, which assumes that humans are often imperfect decision-makers, and their *bounded rationality* can be modeled through unseen *latent variables* that influence the decision-making process (McFadden, 1976; McKelvey & Palfrey, 1995), for which we defer the formal definition and introduction to Appendix D.2. In online decision-making, given the *history* information with *multiple* loss vectors, we adopt the following generalization of the quantal response model.

Definition 4.1 (Quantal response against multiple losses). *Given a set of losses $(\ell_i)_{i \in [t]}$, a noise distribution $\epsilon \sim P_{\text{noise}}$, and $\eta_t > 0$, the generalized quantal response against $(\ell_i)_{i \in [t]}$ is defined as*

$$P_{\text{quantal}}^{\eta_t}(a | (\ell_i)_{i \in [t]}) := P_{\text{quantal}}^{\eta_t}\left(a \mid \sum_{i=1}^t \ell_i\right) = \mathbb{P}\left(a \in \arg \min_{a' \in \mathcal{A}} z(a')\right), \text{ where } z = \eta_t \epsilon + \sum_{i=1}^t \ell_i.$$

In simpler terms, the generalized quantal response is defined as the standard quantal response against the *summation* of the losses. Such a model has been investigated in the learning-in-games and behavioral economics literature (see Appendix D.2 for more details). Such a definition is also aligned with our empirical findings on LLMs’ behavioral patterns in Section 3.2: i) evaluating the summation/average; ii) introducing randomization in decision-making. To gain more insights into these empirical findings, we next analyze a case where pre-training under certain assumptions provably leads to the quantal response behaviors and further yields no-regret guarantees.

4.2 Case Study: Pre-Training under Canonical Data Distribution

Pre-training of LLMs is predominantly based on *next-token prediction*. When applying LLMs to sequential decision-making, the model receives the context of the decision-making task as (x_1, x_2, \dots, x_N) and then generates (x_{N+1}, \dots, x_M) encoding the *action* for some $N, M \in \mathbb{N}^+$ and $N < M$, where each $x_i \in \mathcal{V}$ represents one *natural language token* for $i \in [M]$, and \mathcal{V} is the finite token set. This process can be conceptualized as *predicting the optimal action* in the form of the next token prediction (Yao et al., 2023b; Shinn et al., 2023; Liu et al., 2023a,e). Note that this training procedure may also appear in the form of *supervised fine-tuning (SFT)* for downstream tasks of decision-making or question-answering, where optimal action labels may be easier to obtain (Cobbe et al., 2021; Li et al., 2022; Lewkowycz et al., 2022). Meanwhile, large models are often (pre-)trained under several *fixed/stationary* environments (Laskin et al., 2023; Lin et al., 2024; Lee et al., 2023; Reed et al., 2022), which may limit their ability to handle *arbitrary/non-stationary/adversarial* loss sequences in online learning. Thus, it is natural to ask: *Is it possible to have no-regret behaviors emerging as a consequence of this (optimal) action prediction, under only a fixed pre-training distribution of the environments?*

Here we analyze a standard pre-training objective on a token sequence distribution $x_{1:N_{t+1}} \sim P_t^{text}$ for given $t \in [T]$, which is the expected log-likelihood maximization for next-token prediction over Θ , the parameter space of the LLM:

$$\max_{\theta \in \Theta} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{text}} \sum_{j=1}^{N_{t+1}} \log \text{LLM}_{\theta}(x_j | x_{1:j-1}), \quad (4.1)$$

where we define $\text{LLM}_{\theta}(x_1 | x_{1:0}) = \text{LLM}_{\theta}(x_1)$.

For the pre-training distribution, we model it as follows: there exists a latent variable z , representing the loss for the underlying *static* decision-making problem. The pre-training dataset, however, only contains *partial observations* $x_{1:N_t}$ (a natural language representation of $\ell_{1:t}$) of z due to imperfect data collection, which could be attributed to the fact that z is private to the data-generator (human), representing the actual intention of the human/data-generator. Hence, LLM will only be pre-trained with partial and noisy information about z . Meanwhile, we assume that some high-quality action label $x_{N_t+1:N_{t+1}}$ (a natural language representation of a) with respect to the underlying loss vector z is also available in the dataset, which could come from user surveys, personal blogs, or data annotation. We formalize such an assumption:

Assumption 1 (Pre-training distribution). *Given $T \in \mathbb{N}^+$, $t \in [T]$, $N_{t+1} \in \mathbb{N}^+$, there are latent variables $(z, \ell_{1:t})$, $N_1, \dots, N_t \in [N_{t+1}]$, $N_0 = 0$, such that $\mathbb{P}(z, \ell_{1:t}, x_{1:N_{t+1}}) = \mathbb{P}(z, \ell_{1:t})\mathbb{P}(x_{1:N_t} | \ell_{1:t})\mathbb{P}(x_{N_t+1:N_{t+1}} | z)$, and $P_t^{text}(x_{1:N_{t+1}}) := \mathbb{P}(x_{1:N_{t+1}}) = \int_z \int_{\ell_{1:t}} \mathbb{P}(z, \ell_{1:t}, x_{1:N_{t+1}}) d\ell_{1:t} dz$. Intuitively, tokens $\{x_{N_{i-1}+1:N_i}\}_{i \in [t]}$ encode the context, i.e., information for $\ell_{1:t}$, and the user will decode action a from $x_{N_t+1:N_{t+1}}$.*

To further understand our assumption, we provide an example in Appendix D.3, showing how a natural text corpus may satisfy it. Similar assumptions that suppose the existence of such latent variables in generating the pre-training datasets have also been made recently in Lee et al. (2023); Lin et al. (2024); Liu et al. (2023e), for understanding the in-context decision-making behaviors of LLMs/Transformers through posterior sampling, for which we defer a detailed comparison to Appendix D.8. In particular, we show in Theorem 4.1 that if the noise, i.e., $\ell_i - z$ is modeled as Gaussian distributions and $x_{N_t+1:N_{t+1}}$ encodes the optimal action for z , the pre-trained LLM provably recovers the prominent human behavior model in Section 4.1, the quantal response model.

Theorem 4.1 (Informal: Emergence of no-regret behavior). *Suppose Assumption 1 holds with both the prior distribution of z and the conditional distribution of $\{\ell_i | z\}_{i \in [t]}$ being Gaussian, and $x_{N_t+1:N_{t+1}}$ encodes the optimal action for z . Then, with the function class of LLM_{θ} being expressive enough, and θ^* being a maximizer of Equation (4.1), the behavior of LLM_{θ^*} follows Definition 4.1. Furthermore, the use of LLM_{θ^*} can achieve no (dynamic) regret for (non-stationary) online learning with full-information/bandit feedback for arbitrary loss vectors (with bounded variation).*

The formal statement and proof are deferred to Appendix D.6. The results show that even when pre-training is conducted solely with loss vectors generated from *stationary* distributions ($\ell_{1:t}$ are i.i.d. conditioned on z), it can still enable the *emergence of no-regret behaviors* in online learning against *potentially adversarial losses*. Key in the proof is the connection of pre-trained LLM models to the online learning algorithm of FTPL. Furthermore, Assumption 1 can be relaxed to better match the actual LLMs' pre-training data distributions from diverse sources (c.f. Appendix D.7), and the prior distribution of z could also be replaced by a general distribution (c.f. Theorem D.2). Finally, we point out its implications for playing games in Appendix D.6.1.

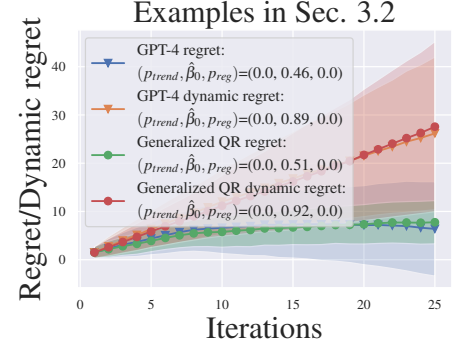


Figure 4.1: Comparison of GPT-4 with the generalized QR model, where the model can very well capture the behavior of the GPT-4 agent for examples in Section 3.2.

How well can our hypothetical model class predict actual LLMs’ behaviors? To further verify our theoretically-justified model in Theorem 4.1, we propose to *estimate* the parameters of $\{\eta_t\}_{t=0}^{T-1}$ in Definition 4.1 using the interaction data with actual LLMs, and use the estimated model to predict LLMs’ behaviors on some test set. In Figure 4.1, we show the averaged regret for the LLMs and our estimated model, where the generalized quantal response can *very well capture* the behavior of the LLM agent for all problem instances in Section 3.2, on which the LLMs oftentimes achieve sublinear regret, justifying the applicability of our hypothetical model and assumptions.

Finally, we acknowledge that for existing pre-trained LLMs like GPT-4, the canonical assumptions above, though may be further relaxed (c.f. Remark D.3), may not hold in general. More importantly, the *supervision labels*, i.e., the optimal action given z , may be sometimes imperfect or unavailable in the dataset. These caveats motivate the study in our next section.

5 Provably Promoting No-Regret Behavior by a New Loss

In light of the observations in Section 3, we ask the question:

*Is there a way to enhance the no-regret property of the models **without** (optimal) action labels?*

To address this question, we propose to train models with a new *unsupervised learning* loss that naturally provides no-regret behaviors. We will particularly focus on the *Transformer* architecture (Vaswani et al., 2017) under this new loss, a common architecture used in most existing LLMs.

5.1 A New Unsupervised Training Loss: *Regret-Loss*

Intuitively, our new training loss is designed to enforce the trained models to minimize regret under an arbitrary sequence of loss vectors. Specifically, we define the training loss as

$$\mathcal{L}(\theta) := \max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \quad (5.1)$$

where $\|\ell_t\|_\infty \leq B$ for $t \in [T]$. As discussed in Kirschner et al. (2023), directly minimizing the max regret can be computationally challenging, except for superficially simple problems. Moreover, Equation (5.1) is not necessarily differentiable with respect to the parameter θ , if it does not satisfy the condition of Danskin’s Theorem (Danskin, 1966); or even if it is differentiable (i.e., the maximizer of $(\ell_t)_{t \in [T]}$ is unique), computation of derivatives can be challenging since we need to calculate $\arg \max_{(\ell_t)_{t \in [T]}} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})$ while there is an inf in the definition of regret. Therefore, we provide a general class of surrogate losses to approximate Equation (5.1):

$$\mathcal{L}(\theta, k, N) := \mathbb{E} \left[\frac{\sum_{j \in [N]} h(\text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) f(\text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)}{\sum_{j \in [N]} f(\text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)} \right], \quad (5.2)$$

where $k \in \mathbb{N}^+$, $N \in \mathbb{N}^+$, $h : \mathbb{R} \rightarrow \mathbb{R}^+$ is a continuous function, with continuous derivative h' , and $f(\cdot, k) : \mathbb{R} \rightarrow \mathbb{R}^+$ is a continuous function for each $k \in \mathbb{N}^+$, satisfying $\lim_{k \rightarrow \infty} \frac{f(R_1, k)}{f(R_2, k)} = \infty \cdot \mathbb{1}(R_1 > R_2) + \mathbb{1}(R_1 = R_2)$, where we use the convention of $\infty \cdot 0 = 0$. These conditions on h, f will be assumed throughout the paper. Examples of such an f include $f(x, k) = x^k$ and $\exp(kx)$. We will sample N trajectories of loss sequences $(\ell_t^{(j)})_{t \in [T], j \in [N]}$ from some continuous probability distribution supported on $[-B, B]^{T \times N}$ (without other additional statistical assumptions), and the expectation in Equation (5.2) is thus taken

with respect to this distribution. In Appendix E.2, we prove that under certain regularity conditions of f and h , we have

$$\lim_{N,k \rightarrow \infty} \mathcal{L}(\theta, k, N) = h \left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right),$$

and the uniform convergence of $\mathcal{L}(\theta, k, N)$: $\lim_{N,k \rightarrow \infty} \sup_{\theta \in \Theta} \left| h \left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right) - \mathcal{L}(\theta, k, N) \right| = 0$, where Θ is a compact set of the model parameters. Hence, one can expect that minimizing the loss function in Equation (5.2) with large enough k and N may promote the trained models to have a small regret value. We will hereafter refer to Equation (5.2) as the *regret-loss*.

5.2 Generalization and Regret Guarantees of Regret-Loss Minimization

We first establish a *statistical* guarantee under general parameterizations of LLM_θ that are Lipschitz with respect to θ , including the Transformer-based models as used in GPT-4 and most existing LLMs (see Proposition 2 for an example with a formal statement). This guarantee focuses on their *generalization ability* when trained to minimize the empirical regret loss (c.f. Equation (E.3)), denoted as $\widehat{\mathcal{L}}(\theta, k, N, N_T)$, by replacing the expectation \mathbb{E} in Equation (5.2) with the empirical mean using N_T samples. We denote $\widehat{\theta}_{k,N,N_T} \in \arg \min_{\theta \in \Theta} \widehat{\mathcal{L}}(\theta, k, N, N_T)$, and present the generalization guarantee in Theorem E.1. Thanks to the uniform convergence of $\mathcal{L}(\theta, k, N)$ (c.f. Appendix E.2), we further obtain the following theorem on the regret guarantee of $\text{LLM}_{\widehat{\theta}_{k,N,N_T}}$:

Theorem 5.1. (Regret). *Suppose² for any $k \in \mathbb{N}^+$, $h, f(\cdot, k)$ are non-decreasing, and $\log f$ is a supermodular function (i.e., $\log f(R_1, k_1) - \log f(R_1, k_2) \geq \log f(R_2, k_1) - \log f(R_2, k_2)$ for $R_1 \geq R_2$ and $k_1 \geq k_2$). Then, with high probability, we have*

$$h \left(\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_{\widehat{\theta}_{k,N,N_T}}}((\ell_t)_{t \in [T]}) \right) \leq h \left(\inf_{\theta \in \Theta} \max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right) + \widetilde{\mathcal{O}} \left(\sqrt{\frac{d_\theta}{N_T}} \right).$$

We defer the proof of the theorem to Appendix E.4. Therefore, if additionally, the model parameterization (e.g., Transformers) can *realize* a no-regret algorithm (as to be shown next), then Theorem 5.1 means that with a large enough N_T , the learned $\text{LLM}_{\widehat{\theta}_{k,N,N_T}}$ becomes a *no-regret* learner, i.e., $\text{Regret}_{\text{LLM}_{\widehat{\theta}_{k,N,N_T}}}((\ell_t)_{t \in [T]}) = o(T)$. Finally, as a consequence, it is folklore that when multiple such LLMs interact, a coarse correlated equilibrium will emerge in the long-term (c.f. Corollary 1).

5.3 Regret-Loss Trained Transformers Can be Online Learning Algorithms

Despite the generality of the previous results, one cannot use an *infinitely large* N and k in practice. Hence, we now provide results when N is finite, for the architecture of *Transformer* models (Vaswani et al., 2017). We focus on single-layer (linear) self-attention models, as in most recent theoretical studies of Transformers (Ahn et al., 2023; Zhang et al., 2023a; Mahankali et al., 2023), and $N = 1$. Note that in this case, the choice of f (and thus k) is not relevant. Thus, throughout this subsection, we drop superscript (j) in Equation (5.2). We sample ℓ_t for $t \in [T]$ as realizations of some random variable Z , where we assume that Z is symmetric about zero, and $\text{Var}(Z) = \Sigma > 0$.

We consider the single-layer *linear* self-attention model as follows, for which we can show that the *global optimizer* of our regret-loss can automatically lead to a no-regret learning algorithm:

$$g(Z_t; V, K, Q, v_c, k_c, q_c) = \sum_{i=1}^t (V \ell_i + v_c) ((K \ell_i + k_c)^\top \cdot (Qc + q_c)). \quad (5.3)$$

²Note that these conditions on h, f are in addition to those specified after Equation (5.2).

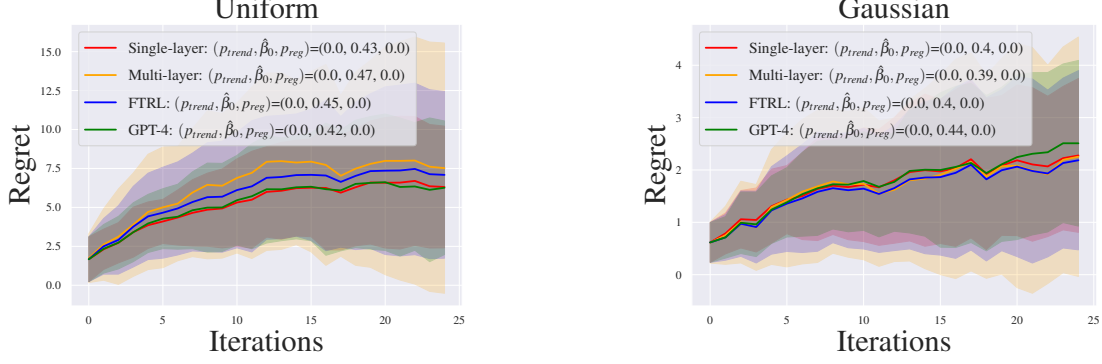


Figure 5.1: Regret performance for the randomly generated loss sequences that are generated by Gaussian with truncation and uniform distribution. No-regret behaviors of single-layer and multi-layer self-attention models are validated by both of our frameworks (low p -values and $\widehat{\beta}_0 < 1$).

Theorem 5.2. Consider the policy space $\Pi = B(0, R_\Pi, \|\cdot\|)$ for some $R_\Pi > 0$. The configuration of a single-layer linear self-attention model in Equation (5.3) (V, K, Q, v_c, k_c, q_c) such that $K^\top(Qc + q_c) = v_c = \mathbf{0}_d$ and $V = -2R_\Pi \Sigma^{-1} \mathbb{E}(\|\sum_{t=1}^T \ell_t\| \ell_1 \ell_2^\top) \Sigma^{-1}$ is a global optimal solution of Equation (5.2) with $N = 1$, $h(x) = x^2$. Moreover, every global optimal configuration of Equation (5.2) within the parameterization class of Equation (5.3) has the same output function g . Additionally, if Σ is a diagonal matrix, then plugging any global optimal configuration into Equation (5.3), and projecting the output with $\text{Proj}_{\Pi, \|\cdot\|}$ is equivalent to FTRL with an L_2 -regularizer.

Theorem 5.2 not only shows the *capacity* of self-attention models: it can realize online learning algorithms, but also shows, more importantly, that minimizing our new regret-loss may *automatically* produce it. In particular, one does not need to hard-code the parameters of the Transformer to implement no-regret algorithms. Under single-layer self-attention parameterization (with softmax), we can also show that a *stationary point* of the loss function (Equation (5.2)) can lead to FTRL (c.f. Appendix E.5). Some potential generalizations of the results are also discussed in Appendix E.9.

5.4 Experimental Results for Regret-Loss Trained Transformers

We now provide experimental results for minimizing our *regret-loss* with the Transformer models, and evaluate in the following environments: 1) randomly-generated loss sequences (Figure 5.1); 2) loss sequences with certain trends (Figure 5.2); 3) repeated games (Figure 5.3); and 4) counterexamples for pre-trained LLMs to be regrettable (Figure 3.5). Training setup can be found in Appendix E.11. We also provide an ablation study for optimizing Equation (5.2) in Appendix E.12.

Randomly generated loss sequences. We use the same loss vectors as those in Section 3.2 for randomly generated loss functions, and compare the results with that using GPT-4. The results show that with regret-loss, both the trained single-layer self-attention model and the trained Transformers with multi-layer self-attention structures can achieve comparable regrets as FTRL and GPT-4. The results can be found in Figure 5.1.

Loss sequences with certain trends. We investigate the case where the loss sequences have predictable trends such as linear-trend or sine-trend. One might expect that the performance of the

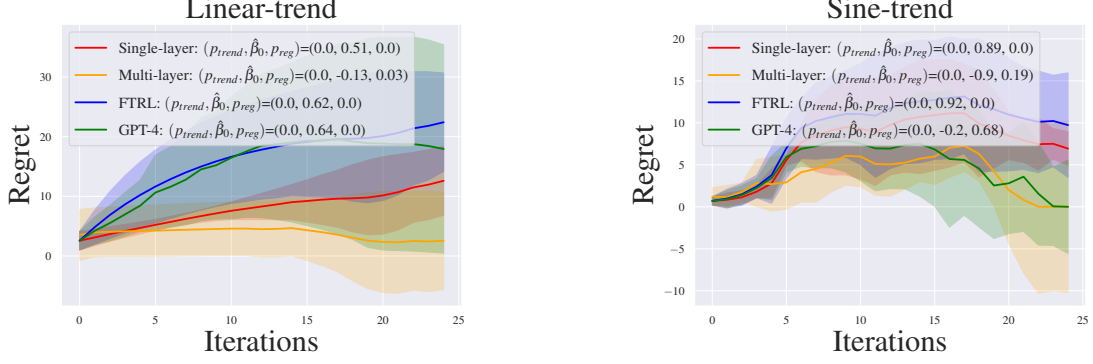


Figure 5.2: Regret performance for the randomly generated loss sequences that are generated by linear-trend and sine-trend. No-regret behaviors of single-layer and multi-layer self-attention models are validated by both of our frameworks (low p -values and $\hat{\beta}_0 < 1$).

trained Transformer would surpass the performance of traditional no-regret learning algorithms such as FTRL, since they may not be an optimal algorithm for the loss sequence with a predictable trend. We modify the training distribution by changing the distribution of random variable Z (which generates the loss vectors ℓ_t) to follow two kinds of trends: linear and sine functions. The results, as illustrated in Figure 5.2, show that the trained single-layer self-attention model and the trained Transformer with multi-layer self-attention structures with regret-loss outperformed GPT-4 and FTRL in terms of regret, when the loss sequence is a linear trend. Similarly, Figure 5.2 shows that the trained Transformer with multi-layer self-attention structures with regret-loss is comparable to GPT-4 and outperformed FTRL in terms of regret, when the loss sequence is a sine-trend. Note that the training dataset does not contain the sequence of losses. Nonetheless, by focusing on the overall trend during training, we can attain performance that is either superior to or on par with that of FTRL and GPT-4.

Repeated games. We then investigate the case of multi-player repeated games. We study 2x2, 3x3x3, 3x3x3x3 games, where each entry of the payoff matrix is sampled randomly from $\text{Unif}([0, 10])$. The results, as illustrated in Figure 5.3, show that the trained single-layer self-attention model and the trained Transformer with multi-layer self-attention structures with regret-loss have a similar performance as that of FTRL. However, GPT-4 still outperforms the trained single-layer self-attention model and the trained Transformer with multi-layer self-attention structures in terms of regret. Since for repeated games (in which the environment faced by the agent can be less adversarial than that in the online setting), there might be a better algorithm than FTRL (see e.g., [Daskalakis et al. \(2021\)](#)), while our self-attention models have a similar structure as FTRL (Theorem E.2 or Theorem 5.2). Also, in practical training (with the empirical loss in Equation (E.3)), we possibly did not find the exact global minimum or stationary point of the *expected* loss in Equation (5.2). Hence, it is possible that GPT-4 may have lower regret than our trained models with the regret-loss.

Two scenarios that caused regrettable behaviors of GPT-4. Finally, we investigate the cases that have caused GPT-4 to have regrettable performance in Section 3.2. The results, which can be found in Figure 3.5, show that both the trained single-layer self-attention model and the trained Transformer with regret-loss can achieve comparable no-regret performance as FTRL, and outperforms that of GPT-4. This validates that our new unsupervised training loss can address the regrettable cases, as our theory in Section 5.2 and 5.3 has predicted.

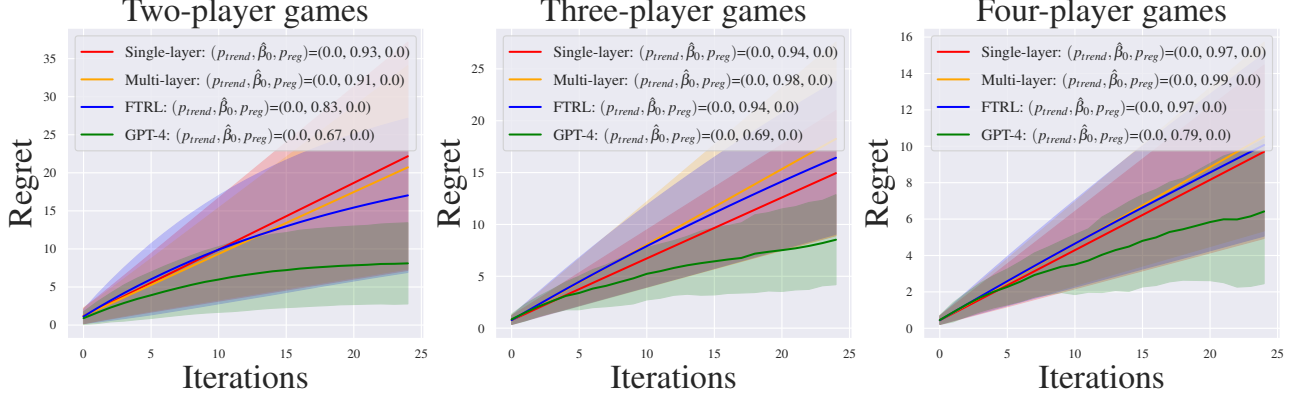


Figure 5.3: Regret performance for the game with two players, three players, and four players general-sum games. No-regret behaviors of single-layer and multi-layer self-attention models are validated by both of our frameworks (low p -values and $\hat{\beta}_0 < 1$).

Remark on performance discrepancy between single-agent and multi-agent settings. In certain cases above, LLMs can outperform FTRL/FTPL algorithms and single/multi-layer models. This phenomenon is primarily observed when the loss sequence exhibits discernible trends, as seen in the single-agent setting. In Section 3.4, we explored this behavior using canonical counterexamples for the follow-the-leader algorithm. Specifically, when the loss sequences display obvious or predictable patterns, LLMs can effectively infer the next loss vector based on historical data, enabling near-optimal decisions. This phenomenon can be further formalized through the lens of in-context learning. Conversely, FTRL/FTPL algorithms, constrained by their update rules, tend to produce near-uniform policies in such cases, as do single/multi-layer Transformer models. In Appendix C.6, we provide ablation studies to support these observations, demonstrating that LLMs leverage trends in the loss sequences by comparing their performance when provided with raw versus summarized historical data. When the loss sequences are summarized (e.g., through aggregation), the resulting loss vectors no longer reflect the trend, leading to significantly diminished performance by the LLMs. In contrast, in multi-agent or game settings, the loss sequence trends depend on the behavior of other agents, rendering them inherently less predictable as all agents continually update their behavior policies. This increased unpredictability likely accounts for the comparable or inferior performance of LLMs relative to the FTRL/FTPL algorithms and single/multi-agent-trained Transformer models in such settings.

6 Concluding Remarks

In this paper, we studied the online decision-making and strategic behaviors of LLMs quantitatively, through the metric of regret. We first examined the no-regret behaviors of several representative pre-trained LLMs in benchmark settings of online learning and games. We then provided some theoretical insights into the sublinear regret growth in the cases we observed, by connecting pre-trained LLMs to the follow-the-perturbed-leader algorithm in online learning, under certain assumptions. We also identified (simple) cases where pre-trained LLMs failed to be no-regret, and thus proposed a new unsupervised training loss, the *regret-loss*, to provably promote the no-regret behavior of Transformers without the labels of (optimal) actions. We established both experimental and theoretical evidence for the effectiveness of our regret-loss.

As a first attempt toward rigorously understanding the online and strategic decision-making

behaviors of LLMs through the metric of regret, We identify the following limitations and list some potential directions for future research:

- There are more than one definitions of (dynamic-)regret in the online learning literature, and we mainly focused on the so-called *external-regret* in the literature. It would be interesting to incorporate other regret metrics, e.g., swap-regret (Blum & Mansour, 2007), which may lead to stronger equilibrium notions in playing repeated games, and policy regret (Arora et al., 2012a), which accounts for adaptive adversaries, into understanding LLM agents’ decision-making behaviors.
- On a related note, no-regret behaviors can sometimes lead to better outcomes in terms of social efficiency (Blum et al., 2008; Roughgarden, 2015; Nekipelov et al., 2015). It would thus be interesting to further validate the efficiency of no-regret LLM agents in these scenarios, as well as identify new prompts and training losses for LLMs to promote the efficiency of the outcomes.
- Our Theorem 4.1 towards explaining why LLMs achieved sublinear regret is highly hypothetical. Given that LLMs are such complex, random, and black-box systems, there are definitely behaviors that our Theorem 4.1 cannot fully capture, and there do exist other possible explanations. For example, an alternative in-context-learning-based explanation may be used to account for the *improved* performance of LLMs on specific loss sequences *with trends* (see Appendix C.6 for more details). However, this explanation may not generalize to the loss sequences *without* obvious trends, complementing our explanations based on the connection to no-regret learning algorithms, which apply to general loss sequences. Hence, it would be interesting to propose and validate other hypotheses for the observed behaviors of LLMs.
- Our new regret-loss has exhibited promises in our experiments for training modest-scale Transformers. One limitation is that we have not trained on other larger-scale *language* models and Foundation Models for decision-making, which is an important ongoing effort.
- To evaluate the performance quantitatively, we focused on online learning and games with *numeric valued* payoffs. It would be interesting to connect our no-regret-based and game-theoretic framework with existing multi-LLM frameworks, e.g., debate, collaborative problem-solving, and human/social behavior simulation, with potentially new notions of regret (defined in different spaces) as performance metrics.

Acknowledgement

The authors thank Constantinos Daskalakis, Kristian Georgiev, Noah Golowich, Dingwen Kong, Akshay Krishnamurthy, and Aleksander Madry for their helpful feedback. In particular, the authors thank Dingwen Kong for discussing the truncation idea in proving Lemma 8, and thank Akshay Krishnamurthy for bringing up a concurrent work that inspired our new experiments for the stochastic bandit setting that strengthened our paper. We sincerely thank Jisu Jang for their valuable feedback on Figure 3.1. X.L. and K.Z. acknowledge the support from the U.S. Army Research Laboratory and the U.S. Army Research Office under grant number W911NF-24-1-0085 and NSF CAREER Award-2443704. C.P. acknowledges support from the Takeda Fellowship and the Korea Foundation for Advanced Studies.

References

- Jacob Abernethy, Chansoo Lee, Abhinav Sinha, and Ambuj Tewari. Online linear optimization via smoothing. In *Conference on Learning Theory*, pp. 807–823. PMLR, 2014.
- Jacob Abernethy, Chansoo Lee, and Ambuj Tewari. Fighting bandits with a new kind of smoothness. *Advances in Neural Information Processing Systems*, 28, 2015.
- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Gati V Aher, Rosa I Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. In *International Conference on Machine Learning*, pp. 337–371. PMLR, 2023.
- Kwangjun Ahn, Xiang Cheng, Hadi Daneshmand, and Suvrit Sra. Transformers learn to implement preconditioned gradient descent for in-context learning. *Advanced in Neural Information Processing Systems*, 2023.
- Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, et al. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*, 2022.
- Mohammad Ahsanullah, Valery B Nevzorov, and Mohammad Shakil. *An introduction to order statistics*, volume 8. Springer, 2013.
- Elif Akata, Lion Schulz, Julian Coda-Forno, Seong Joon Oh, Matthias Bethge, and Eric Schulz. Playing repeated games with large language models. *arXiv preprint arXiv:2305.16867*, 2023.
- Ekin Akyürek, Dale Schuurmans, Jacob Andreas, Tengyu Ma, and Denny Zhou. What learning algorithm is in-context learning? investigations with linear models. *International Conference on Learning Representations*, 2023.
- Lisa P Argyle, Ethan C Busby, Nancy Fulda, Joshua R Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023.
- Raman Arora, Ofer Dekel, and Ambuj Tewari. Online bandit learning against an adaptive adversary: from regret to policy regret. In *Proceedings of the 29th International Conference on Machine Learning*, pp. 1747–1754, 2012a.
- Sanjeev Arora, Elad Hazan, and Satyen Kale. The multiplicative weights update method: a meta-algorithm and applications. *Theory of computing*, 8(1):121–164, 2012b.
- Peter Auer, Nicolo Cesa-Bianchi, Yoav Freund, and Robert E Schapire. The nonstochastic multiarmed bandit problem. *SIAM journal on computing*, 32(1):48–77, 2002.
- Yu Bai, Fan Chen, Huan Wang, Caiming Xiong, and Song Mei. Transformers as statisticians: Provable in-context learning with in-context algorithm selection. *Advanced in Neural Information Processing Systems*, 2023.
- Anton Bakhtin, Noam Brown, Emily Dinan, Gabriele Farina, Colin Flaherty, Daniel Fried, Andrew Goff, Jonathan Gray, Hengyuan Hu, et al. Human-level play in the game of diplomacy by combining language models with strategic reasoning. *Science*, 378(6624):1067–1074, 2022.

- Santiago R Balseiro and Yonatan Gur. Learning in repeated auctions with budgets: Regret minimization and equilibrium. *Management Science*, 65(9):3952–3968, 2019.
- Claude Berge. *Topological spaces: Including a treatment of multi-valued functions, vector spaces and convexity*. Oliver & Boyd, 1877.
- Omar Besbes, Yonatan Gur, and Assaf Zeevi. Stochastic multi-armed-bandit problem with non-stationary rewards. *Advances in neural information processing systems*, 27, 2014.
- Avrim Blum and Yishay Mansour. From external to internal regret. *Journal of Machine Learning Research*, 8(6), 2007.
- Avrim Blum, MohammadTaghi Hajiaghayi, Katrina Ligett, and Aaron Roth. Regret minimization and the price of total anarchy. In *Proceedings of the fortieth annual ACM symposium on Theory of computing*, pp. 373–382, 2008.
- Philip Brookins and Jason Matthew DeBacker. Playing games with GPT: What can we learn about a large language model from canonical strategic games? *Available at SSRN 4493398*, 2023.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901, 2020.
- Sébastien Bubeck, Nicolo Cesa-Bianchi, et al. Regret analysis of stochastic and nonstochastic multi-armed bandit problems. *Foundations and Trends® in Machine Learning*, 5(1):1–122, 2012.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*, 2023.
- Colin F Camerer. *Behavioral game theory: Experiments in strategic interaction*. Princeton University Press, 2011.
- Nicolo Cesa-Bianchi and Gábor Lugosi. *Prediction, Learning, and Games*. Cambridge University Press, 2006.
- Nicolo Cesa-Bianchi, Philip M Long, and Manfred K Warmuth. Worst-case quadratic loss bounds for prediction using linear functions and gradient descent. *IEEE Transactions on Neural Networks*, 7(3):604–619, 1996.
- Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shanghang Zhang, Jie Fu, and Zhiyuan Liu. Chateval: Towards better llm-based evaluators through multi-agent debate. *International Conference on Learning Representations*, 2024.
- Weize Chen, Yusheng Su, Jingwei Zuo, Cheng Yang, Chenfei Yuan, Chen Qian, Chi-Min Chan, Yujia Qin, Yaxi Lu, Ruobing Xie, et al. Agentverse: Facilitating multi-agent collaboration and exploring emergent behaviors in agents. *International Conference on Learning Representations*, 2024.
- Yiting Chen, Tracy Xiao Liu, You Shan, and Songfa Zhong. The emergence of economic rationality of gpt. *Proceedings of the National Academy of Sciences*, 120(51):e2316205120, 2023.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

- Thomas M Cover. *Behavior of sequential predictors of binary sequences*. Number 7002. Stanford University, Stanford Electronics Laboratories, Systems Theory ... , 1966.
- Damai Dai, Yutao Sun, Li Dong, Yaru Hao, Shuming Ma, Zhifang Sui, and Furu Wei. Why can GPT learn in-context? language models secretly perform gradient descent as meta-optimizers. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (eds.), *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 4005–4019, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-acl.247. URL <https://aclanthology.org/2023.findings-acl.247>.
- John M Danskin. The theory of max-min, with applications. *SIAM Journal on Applied Mathematics*, 14(4):641–664, 1966.
- Constantinos Daskalakis, Maxwell Fishelson, and Noah Golowich. Near-optimal no-regret learning in general games. *Advances in Neural Information Processing Systems*, 34:27604–27616, 2021.
- Jingying Ding, Yifan Feng, and Ying Rong. Myopic quantal response policy: Thompson sampling meets behavioral economics. *arXiv preprint arXiv:2207.01028*, 2022.
- Danny Driess, Fei Xia, Mehdi SM Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, et al. Palm-e: An embodied multimodal language model. *International Conference on Machine Learning*, 2023.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2023.
- Christoph Engel, Max RP Grossmann, and Axel Ockenfels. Integrating machine behavior into human subject experiments: A user-friendly toolkit and illustrations. *Available at SSRN*, 2023.
- Ido Erev and Alvin E Roth. Predicting how people play games: Reinforcement learning in experimental games with unique, mixed strategy equilibria. *American Economic Review*, pp. 848–881, 1998.
- Caoyun Fan, Jindou Chen, Yaohui Jin, and Hao He. Can large language models serve as rational players in game theory? a systematic analysis. *arXiv preprint arXiv:2312.05488*, 2023.
- Meir Feder, Neri Merhav, and Michael Gutman. Universal prediction of individual sequences. *IEEE Transactions on Information Theory*, 38(4):1258–1270, 1992.
- Yoav Freund and Robert E Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of computer and system sciences*, 55(1):119–139, 1997.
- Yao Fu, Hao Peng, Tushar Khot, and Mirella Lapata. Improving language model negotiation with self-play and in-context learning from ai feedback. *arXiv preprint arXiv:2305.10142*, 2023.
- Drew Fudenberg and David M Kreps. Learning mixed equilibria. *Games and Economic Behavior*, 5(3):320–367, 1993.
- Drew Fudenberg and David K Levine. *The theory of learning in games*, volume 2. MIT Press, 1998.
- Bolin Gao and Lacra Pavel. On the properties of the softmax function with application in game theory and reinforcement learning. *arXiv preprint arXiv:1704.00805*, 2017.

- Shivam Garg, Dimitris Tsipras, Percy S Liang, and Gregory Valiant. What can transformers learn in-context? a case study of simple function classes. *Advances in Neural Information Processing Systems*, 35:30583–30598, 2022.
- Angeliki Giannou, Shashank Rajput, Jy-yong Sohn, Kangwook Lee, Jason D Lee, and Dimitris Papailiopoulou. Looped transformers as programmable computers. *International Conference on Machine Learning*, 2023.
- Shibo Hao, Yi Gu, Haodi Ma, Joshua Hong, Zhen Wang, Daisy Wang, and Zhiting Hu. Reasoning with language model is planning with world model. In Houda Bouamor, Juan Pino, and Kalika Bali (eds.), *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 8154–8173, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.507. URL <https://aclanthology.org/2023.emnlp-main.507>.
- Elad Hazan. Introduction to online convex optimization. *Foundations and Trends® in Optimization*, 2(3-4):157–325, 2016.
- Josef Hofbauer and William H Sandholm. On the global convergence of stochastic fictitious play. *Econometrica*, 70(6):2265–2294, 2002.
- Sirui Hong, Xiawu Zheng, Jonathan Chen, Yuheng Cheng, Ceyao Zhang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, et al. Metagpt: Meta programming for multi-agent collaborative framework. *International Conference on Learning Representations*, 2024.
- John J Horton. Large language models as simulated economic agents: What can we learn from homo silicus? Technical report, National Bureau of Economic Research, 2023.
- Arnljot Hoyland and Marvin Rausand. *System reliability theory: Models and statistical methods*. John Wiley & Sons, 2009.
- Wenlong Huang, Pieter Abbeel, Deepak Pathak, and Igor Mordatch. Language models as zero-shot planners: Extracting actionable knowledge for embodied agents. In *International Conference on Machine Learning*, pp. 9118–9147. PMLR, 2022a.
- Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, et al. Inner monologue: Embodied reasoning through planning with language models. *arXiv preprint arXiv:2207.05608*, 2022b.
- Hui Jiang. A latent space theory for emergent abilities in large language models. *arXiv preprint arXiv:2304.09960*, 2023.
- Adam Kalai and Santosh Vempala. Efficient algorithms for online decision problems. *Journal of Computer and System Sciences*, 71(3):291–307, 2005.
- Mikołaj J Kasprzak, Ryan Giordano, and Tamara Broderick. How good is your gaussian approximation of the posterior? finite-sample computable error bounds for a variety of useful divergences. *arXiv preprint arXiv:2209.14992*, 2022.
- Johannes Kirschner, Alireza Bakhtiari, Kushagra Chandak, Volodymyr Tkachuk, and Csaba Szepesvari. Regret minimization via saddle point optimization. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- Akshay Krishnamurthy, Keegan Harris, Dylan J Foster, Cyril Zhang, and Aleksandrs Slivkins. Can large language models explore in-context? *arXiv preprint arXiv:2403.15371*, 2024.

- Michael Laskin, Luyu Wang, Junhyuk Oh, Emilio Parisotto, Stephen Spencer, Richie Steigerwald, DJ Strouse, Steven Hansen, Angelos Filos, Ethan Brooks, et al. In-context reinforcement learning with algorithm distillation. *International Conference on Learning Representations*, 2023.
- Tor Lattimore and Csaba Szepesvári. *Bandit algorithms*. Cambridge University Press, 2020.
- Jonathan N Lee, Annie Xie, Aldo Pacchiano, Yash Chandak, Chelsea Finn, Ofir Nachum, and Emma Brunskill. Supervised pretraining can learn in-context reinforcement learning. *Neural Information Processing Systems*, 2023.
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, et al. Solving quantitative reasoning problems with language models. *Advances in Neural Information Processing Systems*, 35: 3843–3857, 2022.
- Chao Li, Xing Su, Chao Fan, Haoying Han, Cong Xue, and Chunmo Zheng. Quantifying the impact of large language models on collective opinion dynamics. *arXiv preprint arXiv:2308.03313*, 2023a.
- Guohao Li, Hasan Abed Al Kader Hammoud, Hani Itani, Dmitrii Khizbullin, and Bernard Ghanem. Camel: Communicative agents for” mind” exploration of large scale language model society. *Neural Information Processing Systems*, 2023b.
- Ruosen Li, Teerth Patel, and Xinya Du. Prd: Peer rank and discussion improve large language model based evaluations. *arXiv preprint arXiv:2307.02762*, 2023c.
- Shuang Li, Xavier Puig, Chris Paxton, Yilun Du, Clinton Wang, Linxi Fan, Tao Chen, De-An Huang, Ekin Akyürek, Anima Anandkumar, et al. Pre-trained language models for interactive decision-making. *Advances in Neural Information Processing Systems*, 35:31199–31212, 2022.
- Siyu Li, Jin Yang, and Kui Zhao. Are you in a masquerade? exploring the behavior and impact of large language model driven social bots in online social networks. *arXiv preprint arXiv:2307.10337*, 2023d.
- Yingcong Li, Muhammed Emrullah Ildiz, Dimitris Papailiopoulos, and Samet Oymak. Transformers as algorithms: Generalization and stability in in-context learning. *International Conference on Machine Learning*, 2023e.
- Zifan Li and Ambuj Tewari. Beyond the hazard rate: More perturbation algorithms for adversarial multi-armed bandits. *J. Mach. Learn. Res.*, 18:183–1, 2017.
- Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate. *arXiv preprint arXiv:2305.19118*, 2023.
- Licong Lin, Yu Bai, and Song Mei. Transformers as decision makers: Provable in-context reinforcement learning via supervised pretraining. *International Conference on Learning Representations*, 2024.
- Nick Littlestone and Manfred K Warmuth. The weighted majority algorithm. *Information and computation*, 108(2):212–261, 1994.
- Hao Liu, Carmelo Sferrazza, and Pieter Abbeel. Chain of hindsight aligns language models with feedback. *arXiv preprint arXiv:2302.02676*, 3, 2023a.

- Junling Liu, Chao Liu, Peilin Zhou, Renjie Lv, Kang Zhou, and Yan Zhang. Is chatgpt a good recommender? a preliminary study. *arXiv preprint arXiv:2304.10149*, 2023b.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-eval: Nlg evaluation using gpt-4 with better human alignment. *arXiv preprint arXiv:2303.16634*, 2023c.
- Yueyang Liu, Benjamin Van Roy, and Kuang Xu. Nonstationary bandit learning via predictive sampling. In *International Conference on Artificial Intelligence and Statistics*, pp. 6215–6244. PMLR, 2023d.
- Zhihan Liu, Hao Hu, Shenao Zhang, Hongyi Guo, Shuqi Ke, Boyi Liu, and Zhaoran Wang. Reason for future, act for now: A principled architecture for autonomous llm agents. In *NeurIPS 2023 Foundation Models for Decision Making Workshop*, 2023e.
- LLM Engine. LLM Engine, 2023. URL <https://llm-engine.scale.com>.
- Nunzio Lorè and Babak Heydari. Strategic behavior of large language models: Game structure vs. contextual framing. *arXiv preprint arXiv:2309.05898*, 2023.
- Arvind Mahankali, Tatsunori B Hashimoto, and Tengyu Ma. One step of gradient descent is provably the optimal in-context learner with one layer of linear self-attention. *International Conference on Learning Representations*, 2023.
- Weichao Mao, Kaiqing Zhang, Ruihao Zhu, David Simchi-Levi, and Tamer Başar. Model-free non-stationary RL: Near-optimal regret and applications in multi-agent RL and inventory control. *arXiv preprint arXiv:2010.03161*, 2020.
- Daniel L McFadden. Quantal choice analysis: A survey. *Annals of Economic and Social Measurement*, Volume 5, number 4, pp. 363–390, 1976.
- Richard D McKelvey and Thomas R Palfrey. Quantal response equilibria for normal form games. *Games and economic behavior*, 10(1):6–38, 1995.
- Sewon Min, Xinxu Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. Rethinking the role of demonstrations: What makes in-context learning work? In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (eds.), *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 11048–11064, Abu Dhabi, United Arab Emirates, December 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.emnlp-main.759. URL <https://aclanthology.org/2022.emnlp-main.759>.
- Gabriel Mukobi, Hannah Erlebach, Niklas Lauffer, Lewis Hammond, Alan Chan, and Jesse Clifton. Welfare diplomacy: Benchmarking language model cooperation. *arXiv preprint arXiv:2310.08901*, 2023.
- Denis Nekipelov, Vasilis Syrgkanis, and Eva Tardos. Econometrics for learning agents. In *ACM Conference on Economics and Computation*, pp. 1–18, 2015.
- Openai. Gpt-4 technical report. 2023.
- Ian Osband, Daniel Russo, and Benjamin Van Roy. (more) efficient reinforcement learning via posterior sampling. *Advances in Neural Information Processing Systems*, 26, 2013.
- Chanwoo Park, Seungju Han, Xingzhi Guo, Asuman Ozdaglar, Kaiqing Zhang, and Joo-Kyung Kim. Maporl: Multi-agent post-co-training for collaborative large language models with reinforcement learning. *arXiv preprint arXiv:2502.18439*, 2025.

- Joon Sung Park, Lindsay Popowski, Carrie Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Social simulacra: Creating populated prototypes for social computing systems. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology*, pp. 1–18, 2022.
- Joon Sung Park, Joseph O’Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, UIST ’23, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9798400701320. doi: 10.1145/3586183.3606763. URL <https://doi.org/10.1145/3586183.3606763>.
- Chen Qian, Xin Cong, Cheng Yang, Weize Chen, Yusheng Su, Juyuan Xu, Zhiyuan Liu, and Maosong Sun. Communicative agents for software development. *arXiv preprint arXiv:2307.07924*, 2023.
- Scott Reed, Konrad Zolna, Emilio Parisotto, Sergio Gómez Colmenarejo, Alexander Novikov, Gabriel Barth-maroon, Mai Giménez, Yury Sulsky, Jackie Kay, Jost Tobias Springenberg, Tom Eccles, Jake Bruce, Ali Razavi, Ashley Edwards, Nicolas Heess, Yutian Chen, Raia Hadsell, Oriol Vinyals, Mahyar Bordbar, and Nando de Freitas. A generalist agent. *Transactions on Machine Learning Research*, 2022. ISSN 2835-8856. URL <https://openreview.net/forum?id=1ikK0kHjvj>. Featured Certification, Outstanding Certification.
- David Robinson and David Goforth. *The topology of the 2x2 games: a new periodic table*, volume 3. Psychology Press, 2005.
- Tim Roughgarden. Intrinsic robustness of the price of anarchy. *Journal of the ACM (JACM)*, 62(5): 1–42, 2015.
- Timo Schick, Jane Dwivedi-Yu, Zhengbao Jiang, Fabio Petroni, Patrick Lewis, Gautier Izacard, Qingfei You, Christoforos Nalmpantis, Edouard Grave, and Sebastian Riedel. Peer: A collaborative language model. *International Conference on Learning Representations*, 2023.
- Shai Shalev-Shwartz. *Online learning: Theory, algorithms, and applications*. Hebrew University, 2007.
- Shai Shalev-Shwartz. Online learning and online convex optimization. *Foundations and Trends® in Machine Learning*, 4(2):107–194, 2012.
- Shai Shalev-Shwartz and Yoram Singer. A primal-dual perspective of online learning algorithms. *Machine Learning*, 69:115–142, 2007.
- Yongliang Shen, Kaitao Song, Xu Tan, Dongsheng Li, Weiming Lu, and Yueting Zhuang. Hugging-gpt: Solving AI tasks with chatgpt and its friends in huggingface. *Neural Information Processing Systems*, 2023.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik R Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- Significant Gravitas. Autogpt, 2023. URL <https://github.com/Significant-Gravitas/AutoGPT>.
- Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research*, 2023.

- Melanie Swan, Takashi Kido, Eric Roland, and Renato P dos Santos. Math agents: Computational infrastructure, mathematical embedding, and genomics. *arXiv preprint arXiv:2307.02502*, 2023.
- Chen Feng Tsai, Xiaochen Zhou, Sierra S Liu, Jing Li, Mo Yu, and Hongyuan Mei. Can large language models play text games well? current state-of-the-art and open questions. *arXiv preprint arXiv:2304.02868*, 2023.
- Karthik Valmeekam, Matthew Marquez, Alberto Olmo, Sarath Sreedharan, and Subbarao Kambhampati. Planbench: An extensible benchmark for evaluating large language models on planning and reasoning about change. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023.
- Aad W Van der Vaart. *Asymptotic statistics*, volume 3. Cambridge university press, 2000.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- Johannes Von Oswald, Eyvind Niklasson, Ettore Randazzo, João Sacramento, Alexander Mordvintsev, Andrey Zhmoginov, and Max Vladymyrov. Transformers learn in-context by gradient descent. In *International Conference on Machine Learning*, pp. 35151–35174. PMLR, 2023.
- Volodimir G Vovk. Aggregating strategies. In *Proceedings of the third Annual Workshop on Computational Learning Theory*, pp. 371–386, 1990.
- Martin J Wainwright. *High-dimensional statistics: A non-asymptotic viewpoint*, volume 48. Cambridge university press, 2019.
- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023a.
- Xinyi Wang, Wanrong Zhu, and William Yang Wang. Large language models are implicitly topic models: Explaining and finding good demonstrations for in-context learning. *International Conference on Machine Learning 2023 Workshop ES-FoMO*, 2023b.
- Yancheng Wang, Ziyang Jiang, Zheng Chen, Fan Yang, Yingxue Zhou, Eunah Cho, Xing Fan, Xiaojiang Huang, Yanbin Lu, and Yingzhen Yang. Recmind: Large language model powered agent for recommendation. *arXiv preprint arXiv:2308.14296*, 2023c.
- Zihao Wang, Shaofei Cai, Anji Liu, Xiaojian Ma, and Yitao Liang. Describe, explain, plan and select: Interactive planning with large language models enables open-world multi-task agents. *Advances in neural information processing systems*, 2023d.
- Chen-Yu Wei and Haipeng Luo. Non-stationary reinforcement learning without prior knowledge: An optimal black-box approach. In *Conference on learning theory*, pp. 4300–4354. PMLR, 2021.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022a.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022b.

- Likang Wu, Zhi Zheng, Zhaopeng Qiu, Hao Wang, Hongchao Gu, Tingjia Shen, Chuan Qin, Chen Zhu, Hengshu Zhu, Qi Liu, Hui Xiong, and Enhong Chen. A survey on large language models for recommendation, 2024a. URL <https://arxiv.org/abs/2305.19860>.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Shaokun Zhang, Erkang Zhu, Beibin Li, Li Jiang, Xiaoyun Zhang, and Chi Wang. Autogen: Enabling next-gen llm applications via multi-agent conversation framework. *arXiv preprint arXiv:2308.08155*, 2023.
- Yue Wu, Xuan Tang, Tom Mitchell, and Yuanzhi Li. Smartplay: A benchmark for llms as intelligent agents. In *The Twelfth International Conference on Learning Representations*, 2024b.
- Fanzeng Xia, Hao Liu, Yisong Yue, and Tongxin Li. Beyond numeric awards: In-context dueling bandits with llm agents. *arXiv preprint arXiv:2407.01887*, 2024.
- Sang Michael Xie, Aditi Raghunathan, Percy Liang, and Tengyu Ma. An explanation of in-context learning as implicit bayesian inference. *International Conference on Learning Representations*, 2022.
- Kai Xiong, Xiao Ding, Yixin Cao, Ting Liu, and Bing Qin. Examining inter-consistency of large language models collaboration: An in-depth analysis via debate. In Houda Bouamor, Juan Pino, and Kalika Bali (eds.), *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 7572–7590, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.508. URL <https://aclanthology.org/2023.findings-emnlp.508>.
- Yuzhuang Xu, Shuo Wang, Peng Li, Fuwen Luo, Xiaolong Wang, Weidong Liu, and Yang Liu. Exploring large language models for communication games: An empirical study on werewolf. *arXiv preprint arXiv:2309.04658*, 2023a.
- Zelai Xu, Chao Yu, Fei Fang, Yu Wang, and Yi Wu. Language agents with reinforcement learning for strategic play in the werewolf game. *arXiv preprint arXiv:2310.18940*, 2023b.
- Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. *Advances in Neural Information Processing Systems*, 2023a.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *International Conference on Learning Representations*, 2023b.
- H Peyton Young. *Strategic learning and its limits*. OUP Oxford, 2004.
- Hongxin Zhang, Weihua Du, Jiaming Shan, Qinhong Zhou, Yilun Du, Joshua B Tenenbaum, Tianmin Shu, and Chuang Gan. Building cooperative embodied agents modularly with large language models. *International Conference on Learning Representations*, 2024.
- Ruiqi Zhang, Spencer Frei, and Peter L Bartlett. Trained transformers learn linear models in-context. *arXiv preprint arXiv:2306.09927*, 2023a.
- Yufeng Zhang, Fengzhuo Zhang, Zhuoran Yang, and Zhaoran Wang. What and how does in-context learning learn? bayesian model averaging, parameterization, and generalization. *arXiv preprint arXiv:2305.19420*, 2023b.
- Qinlin Zhao, Jindong Wang, Yixuan Zhang, Yiqiao Jin, Kaijie Zhu, Hao Chen, and Xing Xie. Competeai: Understanding the competition behaviors in large language model-based agents. *arXiv preprint arXiv:2310.17512*, 2023.

Julian Zimmert and Yevgeny Seldin. Tsallis-inf: An optimal algorithm for stochastic and adversarial bandits. *The Journal of Machine Learning Research*, 22(1):1310–1358, 2021.

Martin Zinkevich. Online convex programming and generalized infinitesimal gradient ascent. In *International Conference on Machine Learning*, pp. 928–936, 2003.

Supplementary Materials for “Do LLM Agents Have Regret? A Case Study in Online Learning and Games”

Contents

1	Introduction	1
1.1	Related Work	3
2	Preliminaries	5
2.1	Online Learning & Repeated Games	5
2.2	Self-attention & Transformers	6
2.3	Performance Metric: Regret	7
3	Do Pre-Trained LLMs Have Regret? Experimental Validation	7
3.1	Framework for Sublinear Regret Behavior Validation	8
3.2	Results: Online Learning	9
3.3	Results: Multi-Player Repeated Games	10
3.4	Pre-Trained LLM Agents Can Still Have Regret	11
4	Why Do Pre-Trained LLMs (Not) Have Regret? A Hypothetical Model and Some Theoretical Insights	13
4.1	A (Human) Decision-Making Model: Quantal Response	13
4.2	Case Study: Pre-Training under Canonical Data Distribution	13
5	Provably Promoting No-Regret Behavior by a New Loss	15
5.1	A New Unsupervised Training Loss: <i>Regret-Loss</i>	15
5.2	Generalization and Regret Guarantees of Regret-Loss Minimization	16
5.3	Regret-Loss Trained Transformers Can be Online Learning Algorithms	16
5.4	Experimental Results for Regret-Loss Trained Transformers	17
6	Concluding Remarks	19
A	Comparison with Concurrent Work Krishnamurthy et al. (2024)	33
B	Deferred Background	34
B.1	In-Context Learning	34
B.2	Online Learning Algorithms	35
B.3	Why Focusing on Linear Loss Function?	36
B.4	Six Representative General-Sum Games	37
C	Deferred Results and Proofs in Section 3	37
C.1	Intuition Why Pre-Trained Language Models Might Exhibit No-Regret Behavior	37
C.2	Frameworks for No-Regret Behavior Validation	38
C.3	Deferred Experiments for Non-stationary Environments in Section 3.2	39
C.4	Deferred Experiments for Bandit-feedback Environments in Section 3.2	41
C.5	Additional Figures for Section 3.3	42
C.6	Additional Results for Section 3.4	43
C.7	Ablation Study on the Prompt	44
C.8	Results for GPT-4 Turbo	47

C.9	LLM Agents’ Explanation on Their Output Policies	48
C.10	Case Studies on Real-world Applications	49
C.10.1	Sequential Recommendation	49
C.10.2	Interactive Negotiation	50
D	Deferred Results and Proofs in Section 4	53
D.1	Pre-Trained LLMs Have Similar Regret as Humans (Who Generate Data)	53
D.2	Background and Motivations for (Generalized) Quantal Response	54
D.3	The Example Instantiating Assumption 1	55
D.4	Alignment of Assumption 1 with Quantal Response	55
D.5	Relationship between FTPL and Definition 4.1	57
D.6	Formal Statement and Proof of Theorem 4.1	57
D.6.1	Implications of Theorem 4.1 for Repeated Games	61
D.7	Extending Theorem 4.1 with Relaxed Assumptions	61
D.7.1	Relaxation under More General Data Distributions	61
D.7.2	Relaxation under Decision-Irrelevant Pre-Training Data	63
D.8	Comparison with Lee et al. (2023) ; Lin et al. (2024) ; Liu et al. (2023e)	64
D.9	Details of Estimating the Parameters of Our Hypothetical Model	64
E	Deferred Results and Proofs in Section 5	64
E.1	Basic Lemmas	64
E.2	Deferred Proof for the Arguments in Section 5.1	65
E.3	Definition of the Empirical Loss Function	70
E.4	Deferred Proofs of Theorem E.1 and Theorem 5.1	70
E.5	Detailed Explanation of Optimizing Equation (5.2) with Single-layer Self-attention Model	74
E.6	Deferred Proof of Theorem E.2	75
E.7	Deferred Proof of Theorem 5.2	78
E.8	Empirical Validation of Theorem E.2 and Theorem 5.2	83
E.8.1	Empirical Validation of Theorem E.2	83
E.8.2	Empirical Validation of Theorem 5.2	84
E.9	Discussions on the Production of FTRL with Entropy Regularization	84
E.9.1	Numerical Analysis of Step 2 and Step 4	88
E.9.2	Empirical Validation	88
E.10	Comparison with In-Context-Learning Analyses in Supervised Learning	89
E.11	Training Details of Section 5.4	89
E.12	Ablation Study on Training Equation (5.2)	89

A Comparison with Concurrent Work [Krishnamurthy et al. \(2024\)](#)

After submitting the first version of our manuscript, we were aware of a concurrent work [Krishnamurthy et al. \(2024\)](#), which considered using LLMs to solve multi-arm *stochastic* bandit problems entirely in-context, with a focus on the *exploration* behaviors of LLMs. Specifically, [Krishnamurthy et al. \(2024\)](#) claimed that LLMs may not show robust exploratory behaviors under a variety of prompt configurations, although there does exist some successful prompt configuration that enabled satisfactory exploratory behaviors. We here provide a detailed comparison between [Krishnamurthy et al. \(2024\)](#) and the first experimental part of our paper, i.e., Section 3 and related appendices.

- **(Focused settings).** We mainly considered the *full-information* online learning setting with potentially *adversarial* loss vectors, as well as the multi-agent *repeated-game* setting. In contrast, [Krishnamurthy et al. \(2024\)](#) focused on the *stochastic* setting with *bandit* feedback, where the loss vectors at different rounds are drawn i.i.d. from a *fixed* distribution. Therefore, both the *metrics* and most *results* are not directly comparable. For example, i) some failure cases in [Krishnamurthy et al. \(2024\)](#) for stochastic bandits did not appear in our setting (as will be detailed next); ii) for some adversarial loss instances (e.g., those from [Feder et al. \(1992\)](#), see the introduction in Section 3.4), the *summarized history input* that was claimed essential in [Krishnamurthy et al. \(2024\)](#) is not very effective in our settings, while a *raw-history input* as in our experiments can be more effective (see Section 3.4 and Figure C.5); iii) as studied in [Krishnamurthy et al. \(2024\)](#), *uniform-like* behaviors constitute one of the main failures in stochastic bandits. However, uniform-like policies do not necessarily correspond to failure cases in our setting, especially when the loss vectors are highly adversarial (c.f. examples in Section 3.4). In particular, such a metric may be irrelevant/inapplicable to validating the no-regret behaviors in our full-information non-stochastic/adversarial settings. These results/facts demonstrated the fundamentally different features in addressing the distinct settings in both works.
- **(Configuration/Prompt design choices).** Despite the negative results under many prompt configurations, [Krishnamurthy et al. \(2024\)](#) still found one successful prompt configuration that can lead to robust exploratory behaviors in stochastic bandits, which in fact shares many similarities with our default prompt configurations. For example, [Krishnamurthy et al. \(2024\)](#) found that asking the LLMs to output a *distribution* over the action space (instead of one *single action*) can address the *suffix failure* for stochastic bandits, which was indeed the default prompt we used in our settings. Moreover, as a standard technique, our default prompt asked the model to have the Chain-of-Thought (CoT) reasoning, while [Krishnamurthy et al. \(2024\)](#)’s successful prompt also emphasized the importance of CoT. [Krishnamurthy et al. \(2024\)](#) also showed the importance of *summarizing* the history, i.e., summarizing the mean reward associated with each arm, while we found that when we feed the LLMs with (raw) *full-information feedback in the vector form*, the LLMs may automatically choose to summarize the history and make decisions based on the summarized statistics (c.f. the output examples in Appendix C.9).
- **(Horizons v.s. No-regret behaviors).** In light of the findings from [Krishnamurthy et al. \(2024\)](#) that LLMs may fail when the problem horizon is long, we conduct experiments on problems with comparable horizons as in [Krishnamurthy et al. \(2024\)](#). Our results show that for the full-information non-stochastic setting we focused on, LLMs can still be no-regret with longer horizons (Figure 3.3 and Table 1), under the loss sequences we studied.
- **(Results in bandit setting & Failure cases).** As an extension and sanity check of our full-information-setting results, we have also experimented with the (adversarial) bandit setting. This extension setting is more comparable to that in [Krishnamurthy et al. \(2024\)](#). However,

Hard MAB instance of Krishnamurthy et al. (2024)	TS	UCB	Successful case of Krishnamurthy et al. (2024)	Ours (GPT-4)	Naive (GPT-4)	Ours (GPT-4o)	Naive (GPT-4o)
Median reward (higher is better)	0.47	0.55	0.47	0.46	0.46	0.475	0.455
SuffFailFreq($T/2$) (lower is better)	0.01	0.02	0.00	0.00	0.00	0.00	0.2
$n \cdot \text{MinFrac}$ (lower is better)	0.28	0.18	0.33	0.27	0.38	0.1	0.09

Table 2: Comparing Thompson Sampling (TS), Upper Confidence Bound (UCB), and the successful prompt configuration of Krishnamurthy et al. (2024) (from Figure 4 therein) with our approaches (named *Ours* in the table), on the *hard* MAB instance therein. We also conducted ablation studies by *removing* our re-weighting technique (named *Naive* in the table). Note that both *Ours* and *Naive* use *distributional output*, as it is the default prompt configuration we used throughout our paper. Specifically, as introduced in Krishnamurthy et al. (2024), for this hard instance, rewards associated with each arm follow a Bernoulli distribution, the horizon is $T = 100$, the number of actions is $n = 5$, and the reward gap is 0.2. For GPT-4, the model adopted by Krishnamurthy et al. (2024), we have observed similar results with their case using the *distributional* output, where although the median reward is comparable with the successful cases, *Naive* suffers from the uniform-like failure as indicated by a high $n \cdot \text{MinFrac}$ value. For GPT-4o, the model not studied by Krishnamurthy et al. (2024), we have a slightly different observation that *Naive* (with distributional output as in our default configurations) seems to still suffer from suffix failure, indicated by a slightly high SuffFailFreq($T/2$), while Krishnamurthy et al. (2024) reported that distributional output *can avoid* such a failure for GPT-4. In contrast to *Naive*, our re-weighting technique enabled the LLMs to avoid *both* the suffix and the uniform-like failures in this (stochastic) bandit-feedback case, *without* external history summarization, and achieve comparable rewards.

different from the focus therein, we did not ask the LLMs to *directly explore* in context. Instead, we manually input a *re-weighting* estimate of the full-information loss vector, a standard technique in online learning (Auer et al., 2002; Hazan, 2016; Lattimore & Szepesvári, 2020), to balance exploration and exploitation. We viewed this approach as a natural way to exploit the no-regret behaviors of LLMs in the full-information setting. In fact, with such a re-weighting, we show in Table 2 that the failure cases in Krishnamurthy et al. (2024) for the bandit setting may not appear, in the exact hard instance proposed therein, even under a relatively *long horizon* of $T = 100$. Complementing Krishnamurthy et al. (2024), our bandit-setting results may suggest that such *human-intervened input* may enhance LLMs’ decision-making capabilities. This is perhaps also in line with the observation in Krishnamurthy et al. (2024) that some additional “human intervention” (i.e., the *summarized* history input therein) may be critical in the (stochastic) bandit setting. Specifically, in Table 2, we validate that although LLMs may fail in bandit-feedback settings *without interventions*, such a simple re-weighting technique may be useful to handle exploration tasks by leveraging LLMs’ performance in the full-information setting.

B Deferred Background

B.1 In-Context Learning

In-context learning is an emergent behavior of LLMs (Brown et al., 2020), which means that these

models can adapt and learn from a limited number of examples provided within their immediate input context. In in-context learning, the prompt is usually constituted by a length of T in-context (independent) examples $(x_t, y_t)_{t \in [T]}$ and $(T + 1)$ -th input x_{T+1} , so the LLM($(z_t)_{t \in [T]}, x_{T+1}$) provides the inference of y_{T+1} , where $z_t = (x_t, y_t)$.

B.2 Online Learning Algorithms

Follow-the-regularized-leader (FTRL). The *follow-the-regularized-leader* algorithm (Shalev-Shwartz, 2007) is an iterative method that updates policy based on the observed data and a regularization term. The idea is to choose the next policy that minimizes the sum of the past losses and a regularization term.

Mathematically, given a sequence of loss vectors $\ell_1, \ell_2, \dots, \ell_t$, the FTRL algorithm updates the policy π at each time step t as follows:

$$\pi_{t+1} = \arg \min_{\pi \in \Pi} \left(\sum_{i=1}^t \langle \ell_i, \pi \rangle + R(\pi) \right),$$

where $R(\pi)$ is a regularization term. The regularization term $R(\pi)$ is introduced to prevent overfitting and can be any function that penalizes the complexity of the model. A function $R(\pi)$ is said to be λ -strongly convex with respect to a norm $\|\cdot\|$ if for all $\pi, \pi' \in \Pi$:

$$R(\pi) \geq R(\pi') + \langle \nabla R(\pi'), \pi - \pi' \rangle + \frac{\lambda}{2} \|\pi - \pi'\|_2^2.$$

A key property that ensures the convergence and stability of the FTRL algorithm is the strong convexity of the regularization term $R(\pi)$. Strong convexity of $R(\pi)$ ensures that the optimization problem in FTRL has a unique solution. The FTRL algorithm's flexibility allows it to encompass a wide range of online learning algorithms, from gradient-based methods like online gradient descent to decision-making algorithms like Hedge (Freund & Schapire, 1997).

Connection to online gradient descent (OGD). The Online Gradient Descent (OGD) (Cesa-Bianchi et al., 1996) algorithm is a special case of the FTRL algorithm when the regularization term is the L_2 -norm square, i.e., $R(\pi) = \frac{1}{2} \|\pi\|_2^2$ and $\Pi = \mathbb{R}^d$. In OGD, at each time step t , the policy π is updated using the gradient of the loss function:

$$\pi_{t+1} = \pi_t - \ell_t.$$

Therefore, the connection between FTRL and OGD can be seen by observing that the update rule for FTRL with L_2 -regularization can be derived from the OGD update rule.

Connection to the Hedge algorithm. The Hedge algorithm (Freund & Schapire, 1997) (also referred to as the Multiplicative Weight Update algorithm (Arora et al., 2012b)) is an online learning algorithm designed for problems where the learner has to choose from a set of actions (denoted as \mathcal{A}) at each time step and suffers a loss based on the chosen action. The FTRL framework can be used to derive the Hedge algorithm by considering an entropy regularization term. Specifically, the regularization term is the negative entropy $R(\pi) = -\sum_{j \in [d]} \pi_j \log \pi_j$ (where d is the dimension of policy π), then the FTRL update rule yields the Hedge algorithm as

$$\pi_{(t+1)j} = \pi_{tj} \frac{\exp(-\ell_{tj} \pi_{tj})}{\sum_{i \in [d]} \exp(-\ell_{ti} \pi_{ti})}$$

for $j \in [d]$.

Follow-the-perturbed-leader (FTPL). Given a sequence of loss vectors $\ell_1, \ell_2, \dots, \ell_{t-1}$, the follow-the-perturbed-leader algorithm (Kalai & Vempala, 2005) at each time step t adds a random perturbation vector ϵ_t to the original loss vectors and then selects the best-response action a_t (that is potentially randomized due to ϵ_t) by solving:

$$a_t \in \arg \min_{a \in \mathcal{A}} \epsilon_{ta} + \sum_{i=1}^{t-1} \ell_{ia},$$

where the perturbation ϵ_t is *sampled* from a pre-defined distribution. Correspondingly, the *policy* π_t is chosen by following equation:

$$\pi_t = \mathbb{E} \left[\arg \min_{\pi \in \Pi} \langle \epsilon_t, \pi \rangle + \sum_{i=1}^{t-1} \langle \ell_i, \pi \rangle \right]. \quad (\text{B.1})$$

Relationship between FTRL and FTPL. The FTRL and FTPL algorithms are deeply related. For example, FTPL with perturbations of Gumbel distribution and FTRL with Entropy Regularization (i.e., Hedge) are equivalent. In general, for the FTPL algorithm with any perturbation distribution, one can always find an FTRL algorithm with a particular regularization such that their update rule is equivalent. However, this relationship does not hold vice versa. For example, Hofbauer & Sandholm (2002) showed that for FTRL with log barrier regularization, there does not exist an equivalent perturbation distribution for FTPL.

Restarting techniques for non-stationary online learning. For non-stationary online learning problems, one common technique is *restarting*: one restarts the standard online learning algorithm periodically (Besbes et al., 2014) (see also e.g., Wei & Luo (2021); Mao et al. (2020)). After each restarting operation, the algorithm will ignore the previous history and execute as if it is the beginning of the interaction with the environment. Since the variation of the loss sequences is bounded, loss sequences between two consecutive restarting operations can be regarded as being *almost stationary*, which makes achieving an overall sublinear dynamic regret guarantee possible.

B.3 Why Focusing on Linear Loss Function?

We note that focusing on the linear loss function $f_t(\pi) := \langle \ell_t, \pi \rangle$ does not lose much of generality. Specifically, for the general convex loss function $(f_t)_{t \in [T]}$, we have $f_t(\pi_{\mathcal{A},t}) - f_t(\pi) \leq \langle \nabla f_t(\pi_{\mathcal{A},t}), \pi_{\mathcal{A},t} - \pi \rangle$ for any $\pi \in \Pi$, which indicates

$$\text{Regret}_{\mathcal{A}}((f_t)_{t \in [T]}) \leq \sum_{t=1}^T \mathbb{E}[\langle \nabla f_t(\pi_{\mathcal{A},t}), \pi_{\mathcal{A},t} \rangle] - \inf_{\pi \in \Pi} \sum_{t=1}^T \mathbb{E}[\langle \nabla f_t(\pi_{\mathcal{A},t}), \pi \rangle].$$

Therefore, one can regard the loss vector $(\ell_t)_{t \in [T]}$ as $\ell_t := \nabla f_t(\pi_{\mathcal{A},t})$ for $t \in [T]$, and control the actual regret by studying the linear loss function (Hazan, 2016). The same argument on the general convex f_t can be applied to the dynamic-regret metric as well. In sum, an algorithm designed for online *linear* optimization can be adapted to solve online *convex* optimization, with the understanding that the instance received at round t corresponds to the gradient of the convex function evaluated at the policy in that round.

B.4 Six Representative General-Sum Games

In game theory, there are six representative two-player general-sum games (Robinson & Goforth, 2005). Firstly, consider the **win-win game** represented by matrices $A = \begin{pmatrix} 1 & 4 \\ 1 & 2 \end{pmatrix}$ and $B = \begin{pmatrix} 1 & 4 \\ 1 & 2 \end{pmatrix}$ for players A and B, respectively. This setup fosters a cooperative dynamic, as both players receive identical payoffs, encouraging strategies that benefit both parties equally.

In contrast, the **prisoner’s dilemma**, depicted by payoff matrices $A = \begin{pmatrix} 1 & 3 \\ 2 & 4 \end{pmatrix}$ and $B = \begin{pmatrix} 4 & 3 \\ 2 & 1 \end{pmatrix}$, illustrates the conflict between individual and collective rationality, where players are tempted to pursue individual gain at the collective’s expense, often resulting in suboptimal outcomes for both.

In the **unfair game**, represented by $A = \begin{pmatrix} 2 & 1 \\ 3 & 4 \end{pmatrix}$ and $B = \begin{pmatrix} 4 & 3 \\ 1 & 2 \end{pmatrix}$, the asymmetry in the payoff structure places one player at a disadvantage, regardless of the chosen strategy. This imbalance often reflects real-world scenarios where power or information asymmetry affects decision-making.

The **cyclic game**, with matrices $A = \begin{pmatrix} 3 & 1 \\ 2 & 4 \end{pmatrix}$ and $B = \begin{pmatrix} 3 & 4 \\ 2 & 1 \end{pmatrix}$, presents a scenario where no stable equilibrium exists. The best strategy for each player changes in response to the other’s actions, leading to a continuous cycle of strategy adaptation without a clear resolution.

The **biased game**, depicted by $A = \begin{pmatrix} 3 & 2 \\ 1 & 4 \end{pmatrix}$ and $B = \begin{pmatrix} 4 & 2 \\ 1 & 3 \end{pmatrix}$, inherently favors one player, often reflecting situations where external factors or inherent advantages influence outcomes, leading to consistently unequal payoffs.

Finally, the **second-best game**, with payoff matrices $A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$ and $B = \begin{pmatrix} 1 & 4 \\ 3 & 2 \end{pmatrix}$, encapsulates scenarios where players settle for less-than-optimal outcomes due to constraints like risk aversion or limited options. This often results in players choosing safer, albeit less rewarding, strategies.

Each of these games exemplifies distinct aspects of strategic decision-making and interactions. From cooperative to competitive and fair to biased scenarios, these matrices provide a rich landscape for exploring the nuances of decision-making behavior in game theory.

C Deferred Results and Proofs in Section 3

C.1 Intuition Why Pre-Trained Language Models Might Exhibit No-Regret Behavior

Intuition why pre-trained language models might exhibit no-regret behavior. Transformer-based LLMs have demonstrated impressive *in-context-learning* and few-/zero-shot learning capabilities (Brown et al., 2020; Garg et al., 2022; Min et al., 2022). One theoretical explanation is that, trained Transformers can implement the *gradient descent algorithm* on the testing loss in certain supervised learning problems (Akyürek et al., 2023; Von Oswald et al., 2023; Dai et al., 2023; Ahn et al., 2023; Zhang et al., 2023a; Mahankali et al., 2023), which is inherently *adaptive* to the loss function used at test time. On the other hand, it is known in online learning that the simple algorithm of *online gradient descent* (Zinkevich, 2003) can achieve no-regret. Hence, it seems reasonable to envision the no-regret behavior of such meta-learners in online learning, due to their fast adaptability. However, it is not straightforward due to the fundamental difference between multi-task/meta-learning and online learning settings, as well as the difference between *stationary* and *non-stationary/adversarial* environments in decision-making. Next, we provide both experimental and theoretical studies to validate this intuition.

C.2 Frameworks for No-Regret Behavior Validation

Trend-checking framework. We propose the following hypothesis test:

H_0 : The sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t=1}^{\infty}$ either diverges or converges to a positive constant

H_1 : The sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t=1}^{\infty}$ converges to 0 or a negative constant

with H_0 and H_1 denoting the null and alternative hypotheses, respectively. The notion of convergence is related to $T \rightarrow \infty$ by definition, making it challenging to verify directly with a finite T . As an alternative, we propose a more tractable hypothesis test, albeit a weaker one, that still captures the essence of our objective:

H_0 : The sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t \in [T]}$ does not exhibit a decreasing pattern

H_1 : The sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t \in [T]}$ shows a decreasing pattern

where the “decreasing pattern” here refers to the case when *more than* 1/2 of the elements in the sequence satisfies that $\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t > \text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t+1]}\right)/(t+1)$. Note that we will only apply the framework when the sequence $\left(\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t\right)_{t \in [T]}$ is non-negative, since a negative regret is even more favorable and directly implies no-regret behaviors.

Ideally, one should check if $\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t$ approaches zero or some negative constant as t goes to infinity. With a finite T value, testing these hypotheses provides a method to quantify this – whether we reject H_0 offers a way to measure it. To this end, one needs to count the number of $\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t - \text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t+1]}\right)/(t+1) > 0$, for which we use Proposition 1 below to provide some understanding of (how small) the probability it happens under various counts. For example, with the default choice of $T = 25$ in our experiments later, one can see from Proposition 1 that: $\mathbb{P}_{H_0}(\mathcal{E}(17, 25)) < 0.032$, $\mathbb{P}_{H_0}(\mathcal{E}(19, 25)) < 0.0035$, $\mathbb{P}_{H_0}(\mathcal{E}(21, 25)) < 0.00014$, i.e., one can easily reject H_0 with high probability. We will report the p -value of H_0 , denoted as p_{trend} , as the output of this framework.

Proposition 1. (p -value of the null hypothesis). *Define the event*

$$\mathcal{E}(s, T) := \left\{ \text{The number of } \frac{\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)}{t} - \frac{\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t+1]}\right)}{t+1} > 0 \text{ for } t = 1, \dots, T \text{ is at least } s \geq \frac{T-1}{2} \right\}.$$

Under the assumption that the null hypothesis H_0 holds, the probability of this event happening is bounded as $\mathbb{P}_{H_0}(\mathcal{E}(s, T)) \leq \frac{1}{2^{T-1}} \sum_{t=s}^{T-1} \binom{T-1}{t}$.

Proof. Under the null hypothesis H_0 , the probability p that $\text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t]}\right)/t - \text{Regret}_{\mathcal{A}}\left((f_{\tau})_{\tau \in [t+1]}\right)/(t+1) > 0$ is less than $\frac{1}{2}$. Therefore, if we consider the event $\mathcal{E}(s, T)$, we have

$$\mathbb{P}_{H_0}(\mathcal{E}(s, T)) = \sum_{k=s}^{T-1} p^k (1-p)^{T-1-k} \binom{T-1}{k} \leq \frac{1}{2^{T-1}} \sum_{k=s}^{T-1} \binom{T-1}{k} \quad (\text{C.1})$$

since $s \geq \frac{T-1}{2}$. □

On the underlying assumption for Equation (C.1). Our *trend-checking* framework was meant to be designed for general sequences $\{a_t\}_{t=1}^T$ for which we *do not know beforehand* how they were generated, since in the online learning setting, by definition, there should be *no* prior assumption on how $\{\text{Regret}_t/t\}_{t=1}^T$ is generated, which very much depends on *both* how the loss sequences and how the policies are generated (by the algorithms).

Our approach implicitly assumes that $(a_{t+1} - a_t)_{t=1}^T$ is mutually independent. We used this assumption since without knowing how $\{\text{Regret}_t/t\}_{t=1}^T$ were generated, one possible (statistical) assumption to model arbitrarily changing sequences is that at each t , some new element is generated randomly and independently, without being affected/biased by any previous elements in the sequence (since we do not know a priori how to model it). Meanwhile, it is possible that the assumption might not hold since it depends on how loss sequences are generated or how LLM behaves. However, it is possible that Equation (C.1) still holds approximately. Specifically, we define

$$\Delta_t = \frac{\text{Regret}_t}{t} - \frac{\text{Regret}_{t+1}}{t+1},$$

and treat $(\Delta_t)_{t=1}^T$ as random variables. We first compute the correlations among those random variables in Figure C.1 using data from Section 3.2, where we can see that the correlations among those random variables are indeed quite small. Meanwhile, this further implies that

$$\begin{aligned} \mathbb{E} \left[\sum_{t=1}^T \mathbf{1}[\Delta_t > 0] \right] &= \sum_{t=1}^T \mathbb{E}[\mathbf{1}[\Delta_t > 0]], \\ \text{Var} \left(\sum_{t=1}^T \mathbf{1}[\Delta_t > 0] \right) &\approx \sum_{t=1}^T \text{Var}(\mathbf{1}[\Delta_t > 0]), \end{aligned}$$

i.e., the random variable $\sum_{t=1}^T \mathbf{1}[\Delta_t > 0]$ indeed has the same first-order and second-order moment as in the case where those random variables $\{\mathbf{1}[\Delta_t > 0]\}_{t \in [T]}$ are independent. Therefore, we regard a Binomial distribution (i.e., assuming $\{\mathbf{1}[\Delta_t > 0]\}_{t \in [T]}$ to be independent) to be an approximation for the actual behaviors of $\sum_{t=1}^T \mathbf{1}[\Delta_t > 0]$, which finally gives Equation (C.1). In fact, when binary random variables have weak correlations (but are not necessarily independent), using the Binomial distribution as an approximation for their summation is also used in the Systems Engineering literature (Hoyland & Rausand, 2009).

C.3 Deferred Experiments for Non-stationary Environments in Section 3.2

We experiment on the setting where the losses are still changing over time, but their total variations across time are bounded, more concretely, sublinear in T . Correspondingly, we consider the stronger metric of *dynamic regret* here to measure the performance. Note that without constraining the variation of the loss vectors, dynamic regret can be linear w.r.t. T in the worst case. Hence, we generate the loss vectors in two different ways: 1) *Gradual variation*. We firstly sample $\ell_1 \sim \text{Unif}([0, 10]^d)$. Then for each $t \geq 2$, we uniformly and randomly generate ℓ_{t+1} under the constraint $\|\ell_{t+1} - \ell_t\|_\infty \leq \frac{1}{\sqrt{t}}$, such that the variations over time are guaranteed to satisfy $\sum_{t=1}^{T-1} \|\ell_{t+1} - \ell_t\|_\infty = o(T)$; 2) *Abrupt variation*. We randomly generate $\ell_1 \sim \text{Unif}([0, 10]^d)$ and m time indices $\{t_i\}_{i \in [m]}$ from $\{1, 2, \dots, T\}$. At each time step t_i for $i \in [m]$, the sign of the loss vector ℓ_{t_i} is flipped, i.e., we let $\ell_{t_i} \leftarrow 10 \cdot \mathbf{1}_d - \ell_{t_i}$. For the specific choice of $T = 25$ in our experiments, we choose $m = 3$. For both cases, the average dynamic regret results are presented in Table 3. GPT-4 achieves sublinear dynamic regret and outperforms *FTRL/FTPL with Restart*, a standard variant of FTRL/FTPL for non-stationary online learning (see e.g., Besbes et al. (2014)). We refer to Appendix B.2 for a detailed introduction of FTRL/FTPL with Restart.

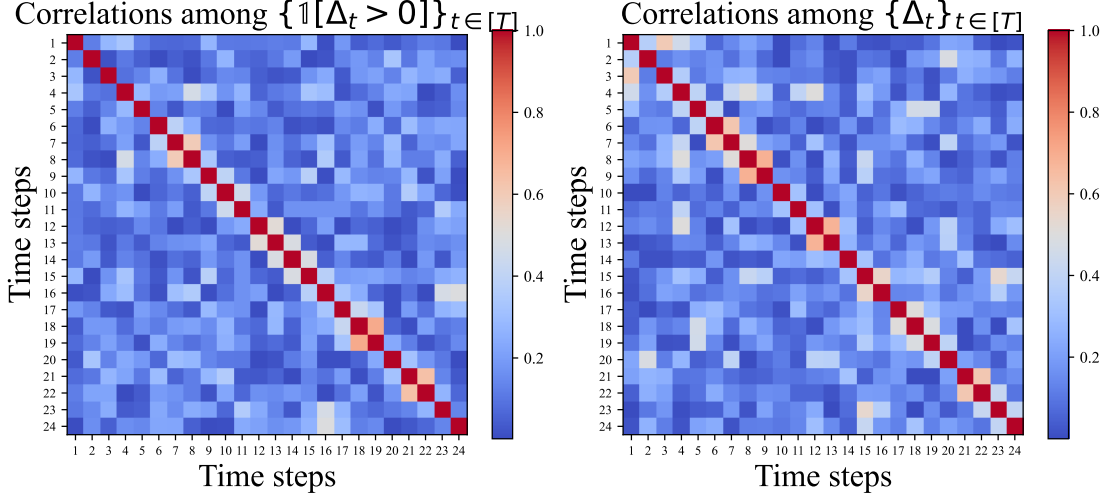


Figure C.1: The absolute value of Pearson correlation coefficient for the random variables $\{1[\Delta_t > 0]\}_{t \in [T]}$ and $\{\Delta_t\}_{t \in [T]}$ using data obtained in Section 3.2.

Dynamic regret		GPT-4	GPT-3.5 Turbo	FTRL	FTPL
Full information	Gradual variation	12.61 ± 7.01 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.0, 0.58, 0.0)$	19.09 ± 11.33 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.0, 0.83, 0.0)$	36.58 ± 24.51	35.19 ± 22.51
	Abrupt variation	30.0 ± 19.91 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.01, 0.87, 0.0)$	33.65 ± 22.51 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.08, 0.96, 0.0)$	36.52 ± 27.68	36.24 ± 28.22
Bandit	Gradual variation	21.39 ± 10.86 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.0, 0.78, 0.0)$	28.42 ± 21.6 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.0, 0.83, 0.0)$	37.64 ± 21.97	36.37 ± 20.7
	Abrupt variation	35.94 ± 28.93 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.42, 0.95, 0.0)$	30.76 ± 25.48 $(p_{trend}, \hat{\beta}_0, p_{reg}) = (0.92, 1.01, 0.0)$	36.52 ± 27.68	38.82 ± 26.17

Table 3: Dynamic regret of GPT-3.5 Turbo/GPT-4 in a non-stationary environment with either full-information or bandit feedback. Every experiment is conducted with 25 rounds. No-regret behaviors of GPT-3.5 Turbo/GPT-4 are validated by both of our frameworks (low p -values and $\hat{\beta}_0 < 1$). The only exception is GPT-3.5 Turbo on loss sequence with abrupt variations under bandit feedback. This indicates that GPT-3.5 Turbo may not be capable of dealing with an abruptly changing environment with limited feedback, although the average regret achieved eventually is still lower than that of other baselines.

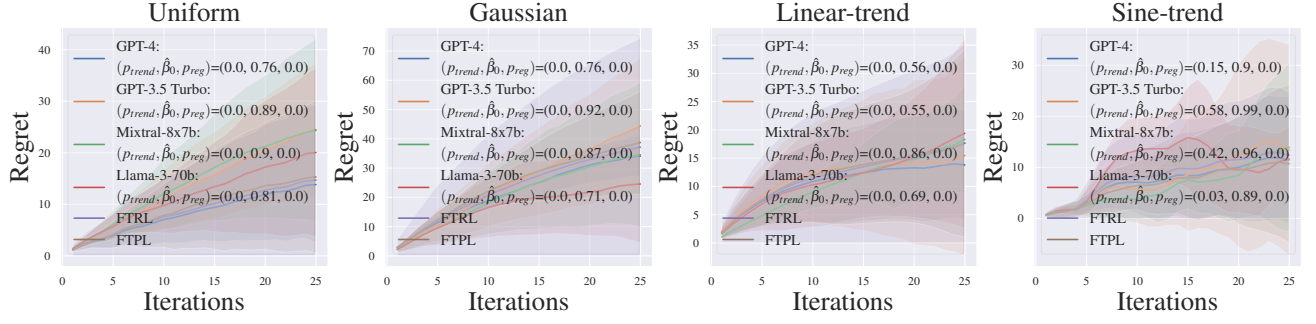


Figure C.2: Regret of pre-trained LLMs for online learning with bandit feedback in 4 different settings. It performs comparably and sometimes even better than well-known no-regret learning algorithms, variants of FTRL and FTPL with bandit-feedback.

C.4 Deferred Experiments for Bandit-feedback Environments in Section 3.2

Although pre-trained LLMs have achieved good performance in online learning with full-information feedback, it is unclear whether they can still maintain no-regret with only bandit feedback. For such problems, we modify the prompt and protocol of interactions slightly, where we still ask the LLM agent to provide a policy π_t at time step t , then sample one $a_t \sim \pi_t(\cdot)$. In the bandit setting, the LLM agent can only access $(a_t, \ell_{t a_t})$. Instead of directly feeding it to the agent, we feed an estimate of the loss vector $\widehat{\ell}_t \in \mathbb{R}^d$, where $\widehat{\ell}_t(a) \leftarrow \frac{\ell_t(a)}{\pi_t(a)} \mathbb{1}(a_t = a)$ for all $j \in [d]$. Note that such an operation of *re-weighting* the loss (when the loss is non-negative) by the inverse of the probability is standard in online learning when adapting full-information-feedback no-regret algorithms to the bandit-feedback ones (Auer et al., 2002). Later, we will also show the benefits of such operations (c.f. Section 4). We compare the performance of pre-trained LLMs with that of the counterparts of FTRL with bandit feedback, e.g., EXP3 (Auer et al., 2002) and the bandit-version of FTPL (Abernethy et al., 2015), in both Figure C.2 and Table 3, where GPT-4 consistently achieves lower regret.

C.5 Additional Figures for Section 3.3

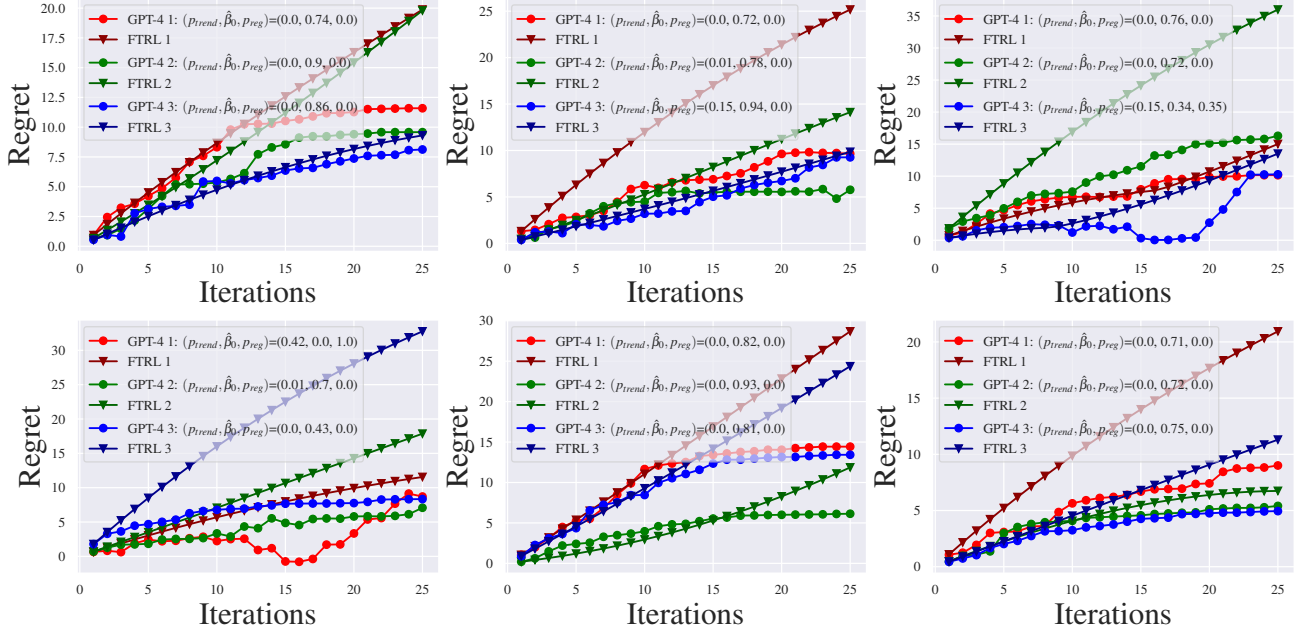


Figure C.3: Regret of GPT-4 and the FTRL algorithm in 6 randomly generated three-player general-sum games. GPT-4 has comparable (even better) no-regret properties when compared with the FTRL algorithm, according to the frameworks in Section 3.1 and the graphic trends..

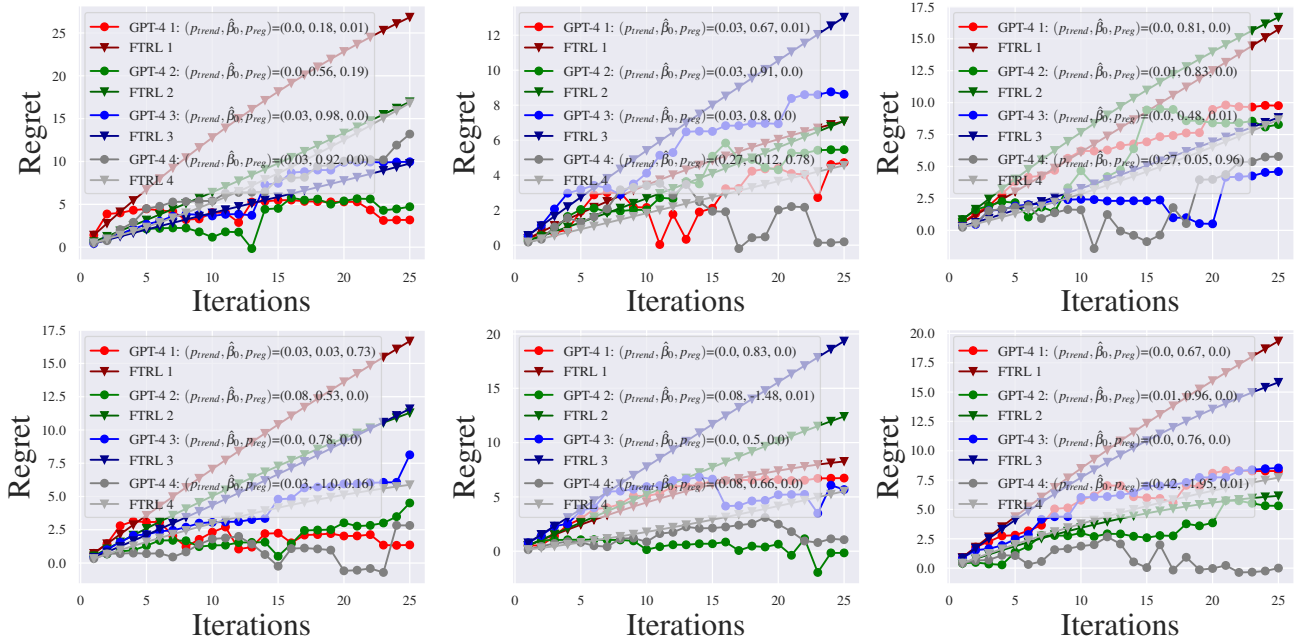


Figure C.4: Regret of GPT-4 and the FTRL algorithm in 6 randomly generated four-player general-sum games. GPT-4 has comparable (even better) no-regret properties when compared with the FTRL algorithm, according to the frameworks in Section 3.1 and the graphic trends.

C.6 Additional Results for Section 3.4

For *Example 2*, we evaluate LLMs on both the $c = 100$ and $c = 200$ cases. The results and comparisons are presented in Figure C.5 using a temperature of 0 to minimize the randomness for such fixed problem instances, where we can confirm that GPT-4 with raw history identifies the pattern and is able to achieve decreasing, negative regret during the first $c = 100$ or $c = 200$ rounds), whereas FTRL, FTPL, and GPT-4 with only summarized history cannot detect the trend and then make adaptive decisions. Meanwhile, after first c rounds, the LLM with raw history can identify that the pattern for the loss vectors has changed to adjust its policy, and its regret grows more slowly than the LLM with only summarized history.

Such observations further demonstrate the fundamental differences in the stochastic settings considered in Krishnamurthy et al. (2024) and our non-stochastic settings: the summarized history, an essential factor for the successful configuration in Krishnamurthy et al. (2024), can be good statistics in the i.i.d. setting (as a good estimate of the *mean* of the losses), while it loses information and can be highly ineffective in the non-stochastic settings that are highly adversarial (Feder et al., 1992). In contrast, with raw history, GPT-4 was able to better identify the pattern of the sequence and make good predictions to achieve even negative regret values.




Figure C.5: Comparing LLMs on *Example 2* in Section 3.4 with raw history as the input and summarized history as the input.

Explaining the better performance of LLMs on losses with trends via in-context learning. LLMs’ in-context-learning capability of being able to *infer* the underlying *trend* in the above case might offer one explanation for the observations above. Specifically, the task of predicting ℓ_{T+1} given past loss sequences $\ell_{1:T}$ could be understood as an in-context learning problem as follows: the demonstration/in-context dataset is given by the following input and label pairs $D = \{x_t, y_t\}_{t \in [T-1]}$, where $x_t = \ell_{1:t}$ and $y_t = \ell_{t+1}$ for each $t \in [T-1]$. Then, LLMs given such demonstration/context D will make prediction based on $x_T = \ell_{1:T}$ (to predict y_T , i.e., the next loss vector ℓ_{T+1}). In other words, in-context learning, in this case, is firstly learning the *trend* from the $T-1$ pairs of inputs and labels, and then making a prediction of the next loss. Hence, when there exists an underlying pattern, in-context-learning can accurately predict the next loss (when raw history is given), and thus achieves good no-regret performance. This perspective may offer an explanation of why LLMs can achieve better performance than FTRL/FTPL when the loss sequences have an obvious trend. Note that, this may also be used to explain why raw-history-based input outperforms the summarized-history-based input in the experiments above – the latter loses such a “context” information, as the mean of

the history losses is not sufficient to predict/infer the underlying trend (even when there exists one). Finally, note that, this “trend prediction” explanation does not apply to general loss sequences, for which our explanation in Section 4 that connects LLMs’ behaviors to FTPL still applies.

C.7 Ablation Study on the Prompt

Ablation study on online learning. To systematically understand the effects of our prompt on the final performance of the LLM agent, we create three different variants of our prompt and report the regret by using different prompts in Figure C.6. Specifically, for **Ablation1**, we remove examples to illustrate the game rules. For **Ablation2**, we remove the number of iterations. For **Ablation3**, we incorporate some *hints* for the LLM for decision-making, including the hints to suggest it to pay attention to the loss history, to behave more greedily at the end of an episode, and also to explain the reason of its decision step-by-step. The latter hint is a popular technique in prompt engineering known as the *Chain-of-Thought* prompting (Wei et al., 2022b). Finally, we recall that d is the number of actions in all prompts.




Figure C.6: Ablation study on our prompt design.

Original prompt
<p>You are solving a decision-making problem for 25 rounds.</p> <p>There are d number of action (which is 0 to $d-1$).</p> <p>At each round, you need to choose a policy, it specifies your probability to choose each action.</p> <p>This policy should be d-dimensional, and the sum of its components should equal 1. After that, you will be shown the reward vector for choosing each action.</p> <p>Remember that this reward vector is decided by the external system and can be potentially different for different rounds.</p> <p>It is not decided by what policies you have chosen. The reward vector is also d-dimensional.</p> <p>It represents the reward of choosing action from 0 to $d-1$.</p> <p>For example, a reward vector of $[0.8, 3.2]$ means reward for action_0 is 0.8 and the reward for action_1 is 3.2.</p> <p>Then your reward for this round will be calculated according to the reward of each action and your probability of choosing each action.</p>

For example, if you choose the policy [0.2, 0.8] and get the reward vector [1, 2], then your expected reward is $0.2 \times 1 + 0.8 \times 2 = 1.8$

Your goal is to maximize your accumulative expected reward.

You can adjust your policy based on the reward vectors for all previous rounds.

You're required to provide your policy in numeric format.

Your response's last line should be formatted as
'Policy: [your d -dimensional policy]'

Ablation1: no examples

You are solving a decision-making problem for 25 rounds.

There are d number of action (which is 0 to $d-1$).

At each round, you need to choose a policy,
it specifies your probability to choose each action.

This policy should be d -dimensional, and the sum of its components should equal 1. After that, you will be shown the reward vector for choosing each action.

Remember that this reward vector is decided by the external system and can be potentially different for different rounds.

It is not decided by what policies you have chosen.
The reward vector is also d -dimensional.

It represents the reward of choosing action from 0 to $d-1$.

Then your reward for this round will be calculated according to the reward of each action and your probability of choosing each action.

Your goal is to maximize your accumulative expected reward.

You can adjust your policy based on the reward vectors for all previous rounds.

You're required to provide your policy in numeric format.

Your response's last line should be formatted as
'Policy: [your d -dimensional policy]'

Ablation2: no round information

You are solving a decision-making problem.

There are d number of action (which is 0 to $d-1$).

At each round, you need to choose a policy,
it specifies your probability to choose each action.

This policy should be d -dimensional, and the sum of its components should equal 1. After that, you will be shown the reward vector for choosing each action.

Remember that this reward vector is decided by the external system and can be potentially different for different rounds.

It is not decided by what policies you have chosen.
The reward vector is also d -dimensional.

It represents the reward of choosing action from 0 to $d-1$.

For example, a reward vector of $[0.8, 3.2]$ means reward for action_0 is 0.8 and the reward for action_1 is 3.2.

Then your reward for this round will be calculated according to the reward of each action and your probability of choosing each action.

For example, if you choose the policy $[0.2, 0.8]$ and get the reward vector $[1, 2]$, then your expected reward is $0.2 \times 1 + 0.8 \times 2 = 1.8$

Your goal is to maximize your accumulative expected reward.

You can adjust your policy based on the reward vectors for all previous rounds.

You're required to provide your policy in numeric format.

Your response's last line should be formatted as
'Policy: [your d -dimensional policy]'.

Ablation3: adding hints

You are solving a decision-making problem for 25 rounds.

There are d number of action (which is 0 to $d-1$).

At each round, you need to choose a policy,
it specifies your probability to choose each action.

This policy should be d -dimensional, and the sum of its components should equal 1. After that, you will be shown the reward vector for choosing each action.

Remember that this reward vector is decided by the external system and can be potentially different for different rounds.

It is not decided by what policies you have chosen.
The reward vector is also d -dimensional.

It represents the reward of choosing action from 0 to $d-1$.

For example, a reward vector of $[0.8, 3.2]$ means reward for action_0 is 0.8 and the reward for action_1 is 3.2.

Then your reward for this round will be calculated according to the reward of each action and your probability of choosing each action.

For example, if you choose the policy $[0.2, 0.8]$ and get the reward vector $[1, 2]$, then your expected reward is $0.2 \times 1 + 0.8 \times 2 = 1.8$

Your goal is to maximize your accumulative expected reward.




Figure C.7: Regret of GPT-4 for repeated games under 3 different prompt ablations. Its performance is consistent among three different prompts.

You can adjust your policy based on the reward vectors for all previous rounds.

You're required to provide your policy in numeric format.

Your response's last line should be formatted as
'Policy: [your d -dimensional policy]'.

Let's think step by step. Explicitly examining history is important.

Please explain how you chose the policy by guessing
what reward you might receive for each action according to the history.

You should explore for first several rounds and behave greedily for
later rounds, for example, choosing one action with probability more
than 0.99.

Please also explain whether you are behaving more greedily and less
greedily by explicitly considering the policy you just used for
last round.

We can see in Figure C.6 that the performances of LLM agents are consistent under different variants of the prompts.

Ablation study on repeated games. For the game setting, we also investigate whether explicitly informing LLM agents that they are ``playing a repeated matrix game with some other opponents'' would affect the performance. Therefore, we evaluate three different prompts by informing LLM agents that they are playing a matrix game, solving multi-arm bandit, or solving general decision-making problems, in the first line of the prompt. We show the performance of such three prompts in Figure C.7, where it is seen that LLM agents' performance on repeated games is consistent among these variants of the prompts.

C.8 Results for GPT-4 Turbo

When preparing our paper, we are aware of a more advanced version of GPT-4, GPT-4 Turbo. Although GPT-4 has successfully exhibited no-regret behaviors in most problems, we also evaluate and report the performance of its more advanced version here. Specifically, we compare GPT-4 and GPT-

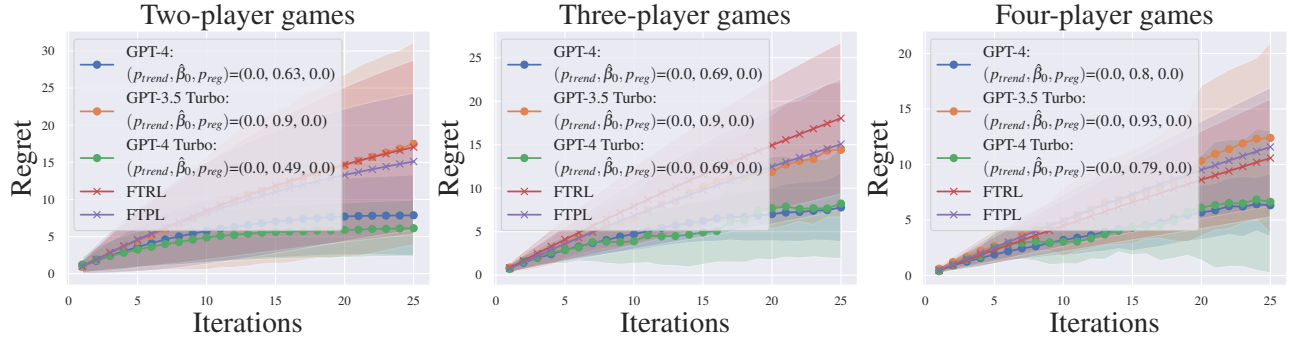


Figure C.8: Comparisons of GPT-4 Turbo with GPT-4 and GPT-3.5 Turbo. Due to the symmetry of agents in randomly generated games, we report the regret of one agent for ease of presentation. We can see that GPT-4 Turbo performs on par with or better than GPT-4.

4 Turbo for the repeated game scenario in Figure C.8, where we validate that GPT-4 Turbo performs either comparably with or better than GPT-4.

C.9 LLM Agents' Explanation on Their Output Policies

We here present two examples of LLMs' explanations of their output policies, showing how they *reason* in their decision-making.

Output Example 1

Looking at the history of the reward vectors, it is clear to see that the reward for choosing action 0 is consistently higher than that for choosing action 1. This pattern is evident across all 5 rounds of the game.

In the first round, the difference in rewards is 3.52, this remained relatively stable with the difference in the last round being 3.52 as well. It's indicative that action 0 consistently brings a higher return than action 1.

Therefore, I'll lean more towards selecting action 0 in my policy.

However, acting deterministically, that is always choosing action 0, would introduce the risk of becoming too predictable and letting the opponent exploit this predictability

So, in order to maintain some unpredictability while still capitalizing on the higher average return of action 0, I'll set my policy as follows: select action 0 with a probability of 0.8 and action 1 with a probability of 0.2.

This way, the majority of the time I'm choosing the action that gives a higher reward, but I'm also occasionally choosing the other action to keep the opponent guessing, and avoid any predictable patterns that can be exploited

So, my policy is as follows: Policy: [0.8, 0.2]

Output Example 2

The total reward from action_1 is 130.92 and the average reward is around 10.07.

On the other hand, the total reward from action_0 is 33.418 and the average reward is around 8.354.

Although action_0 has a less average reward, its highest reward is 26.977, which is considerably high and shows its potential.

For this round, I want to continue to explore action_0 but also significantly exploit action 1 as it reliably gives higher reward.

So, I'd like to decrease the probability of selecting action_0 compared to last time but still keep it in the policy as a possible option.

The policy for the previous round was not provided but assuming it was more balanced, I'm now going to lean more towards action_1.

So, the policy will be: [0.3, 0.7]. Here, I'm behaving more greedily compared to earlier rounds, as we are choosing the action which has been giving us higher rewards with more probability.

As observed from the output, the LLM tends to use the history of the reward vectors by looking at the “sum/average of the accumulated rewards”, and tends to introduce randomization in decision-making, in order to be “less predictable and exploitable”. These are several key components in achieving no-regret in online learning and games (Cesa-Bianchi & Lugosi, 2006), explaining the empirical evidence we had in Section 3.

C.10 Case Studies on Real-world Applications

In this subsection, we evaluate the sequential decision-making abilities of LLMs in realistic scenarios from the perspective of regret and dynamic regret. While several studies have explored sequential decision-making using synthetic scenarios (Krishnamurthy et al., 2024; Wu et al., 2024b; Xia et al., 2024; Akata et al., 2023) or real-world data scenarios (Liu et al., 2023b; Wang et al., 2023c; Wu et al., 2024a), none have explicitly analyzed regret or dynamic regret. As a result, the *theoretical optimality* of such a sequential decision-making process remains unclear.

Our first case study investigates single-agent sequential decision-making using real-world data, leveraging the same dataset and experimental setup as (Liu et al., 2023b). The second case study explores a two-player negotiation scenario, providing insights into dynamic interactions and their impact on decision-making performance.

C.10.1 Sequential Recommendation

We consider the task of sequential recommendation, a task that people have been employing LLMs to solve with success (Liu et al., 2023b; Wang et al., 2023c; Wu et al., 2024a). Note that how existing literature (Liu et al., 2023b) uses LLMs to solve this task fits exactly into our online learning framework, where humans feed a history of items the user have interacted with to the LLM and then ask the LLM to recommend the item (or several items) the user may want to interact next. The entire process carries on repeatedly.

Formally, the problem is as follows. Given a sequence of history items the user has interacted with $(x_1, x_2, \dots, x_{t-1})$, where each $x_i \in D$ for $i \in [t-1]$ and D is the collection of all items, the LLM needs to recommend n items that the user might interact with in the next step t . Typically, the LLM should also give a priority on the n items it recommends. For simplicity here, we here assume they are of equal priority. In other words, at step t , the LLM will take an action $a_t \subseteq D$ with $|a_t| = n$, hoping

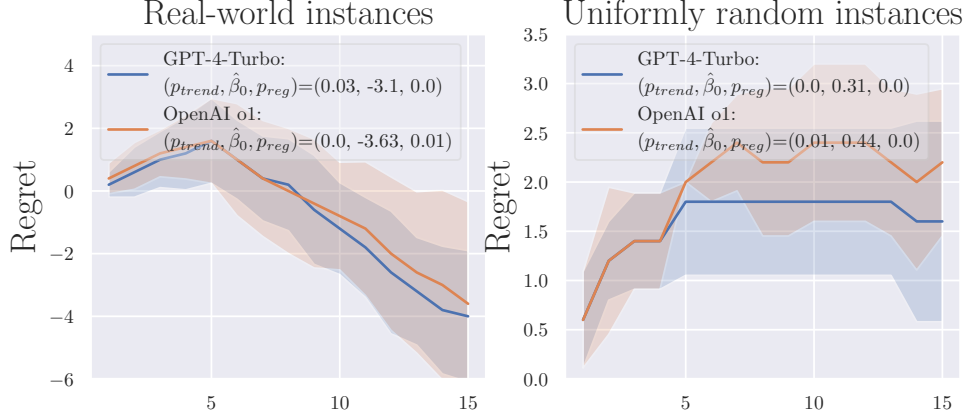


Figure C.9: We evaluate GPT-4-Turbo and OpenAI o1 on both real-world data and uniformly random synthetic data, where we can see both models can still achieve sublinear regret.

what the user will interact at step t belongs to a_t . Hence, the loss is given by $\ell_t(a_t, x_t) := \mathbf{1}[x_t \notin a_t]$. Correspondingly, the regret by our definition is given by

$$\text{Regret}(x_{1:T}) = \sum_{t=1}^T \ell_t(x_t, a_t) - \min_a \sum_{t=1}^T \ell_t(x_t, a).$$

We refer to (Liu et al., 2023b) for a more detailed introduction. Meanwhile, we use the real-world data and follow the experimental setup of (Liu et al., 2023b).

In the left one of Figure C.9, we can observe that LLMs can achieve expressively low and sublinear regret on such a real-world application with real-world data. As a comparison, in the right one of Figure C.9, we replace the real-world data with synthetic data generated in a uniformly random way (it is worth mentioning that the prompt setting still follows the setup of sequential recommendation of Liu et al. (2023b)), where we can see that LLMs can still be no-regret. However, interestingly, LLMs perform better on real-world data, which validates that real-world applications can exhibit certain trends/structures, for which LLMs can exploit and achieve superior performance as we have shown in our paper through synthetic problems with trends.

C.10.2 Interactive Negotiation

The experiment was designed to simulate negotiation scenarios between two LLMs, designated as LLM A and LLM B, across multiple turns. The primary objectives were to analyze multi-agent sequential decision-making processes and quantify regret. For each repetition, an LLM generated unique negotiation topics. Based on these topics, the LLM also created the context, objectives, and relevant background information to design engaging and interactive negotiation scenarios.

Negotiation Process. The negotiation process was executed in a turn-based manner, with each turn comprising three steps:

1. **Intention Generation:** Each LLM defined its goal for the turn, specifying what it aimed to achieve with its response.
2. **Response Generation:** Based on the defined intention and the dialogue history, each LLM generated a response.

3. **Alternative Response Generation:** Three distinct alternative replies were produced for each original response. These alternatives represented diverse negotiation strategies or perspectives while preserving the original intention.

Response Evaluation. After the dialogue concluded, all responses—both original and alternatives—were evaluated using a scoring scale from 1 to 10 based on the following criteria for each turn:

- **Clarity:** How clear and understandable the reply is.
- **Relevance:** How pertinent the reply is to the negotiation topic and the defined intention.
- **Engagement:** How engaging or persuasive the reply is in fostering further dialogue.
- **Alignment with the Stated Intention:** How well the conversation aligns with the turn’s stated intention following the reply. For alternative replies, this was assessed by hypothetically replacing the original reply with an alternative and evaluating the alignment based on the entire conversation.

Each response was scored using an LLM as the evaluator. Although human evaluation would be preferable, the use of an LLM as a scorer was chosen for scalability. This approach is common in the LLM domain and is sometimes referred to as G-eval (where “G” stands for GPT) (Liu et al., 2023c).

Dynamic Regret Analysis. Finally, dynamic regret was calculated to measure suboptimality by comparing the scores of the original replies against the highest-scoring alternative responses. Since calculating regret typically requires hindsight knowledge of the best possible responses, which requires rollout of every possible dialogues, we decide to analyze on dynamic regret. Dynamic regret analysis provided a quantitative measure of decision-making effectiveness across turns. This analysis offered insights into how regret dynamics can inform improved decision-making strategies in real-world negotiation contexts.

Example. Here is an example from our simulation:

Step 1: Generate Topics and Backgrounds. Topics and backgrounds were generated using a language model.

Topic: The Trade-Off Negotiation Between Eco-Tech Innovator and Traditional Manufacturing Tycoon

Background of Player A: Eco-Tech Innovator (Jordan Green). Jordan Green is the CEO of a rapidly growing startup, EcoWave Technologies, which specializes in developing sustainable energy solutions and eco-friendly manufacturing processes. With a background in environmental science and engineering, Jordan is passionate about reducing carbon footprints and promoting renewable energy sources. Their innovative products, such as biodegradable materials and energy-efficient machinery, have garnered attention and accolades within the green tech community. However, despite the startup’s promise, EcoWave faces challenges in scaling production and reaching wider markets due to limited financial resources and manufacturing capabilities.

Background of Player B: Traditional Manufacturing Tycoon (Robert Steele). Robert Steele is the owner of Steele Industries, a well-established manufacturing company known for its mass production of consumer goods. With decades of experience in the industry, Robert has built a reputation for efficiency and profitability, but his company has been criticized for its environmental impact and reliance on outdated practices. Facing increasing pressure from regula-

tors and consumers to adopt sustainable practices, Robert is aware that failing to adapt could threaten the future of his business. However, he is also concerned about the costs and risks associated with transitioning to new technologies and processes, which could disrupt his current operations.

Detailed Situation: Jordan and Robert have agreed to meet to discuss a potential partnership. Robert is interested in integrating EcoWave’s sustainable technologies into his manufacturing processes, which could help Steele Industries improve its environmental footprint and public image. However, he is also hesitant about the costs associated with adopting these new technologies and is particularly concerned about how this might affect his established workforce and production timelines. On the other hand, Jordan sees this as a golden opportunity to scale EcoWave’s innovative solutions and gain access to Robert’s extensive distribution networks. However, Jordan is wary of compromising on key values regarding environmental integrity and sustainable practices, fearing that a partnership with a traditional manufacturer could dilute EcoWave’s mission. The negotiation aims to balance Robert’s need for cost-effectiveness and operational continuity with Jordan’s commitment to sustainability and innovation. Both players must navigate their differing priorities, explore potential compromises, and consider the long-term implications of their decisions on their businesses and the environment.

Step 2: Generate Conversation. We iterated between **Step 2-1** and **Step 2-2** to create a conversation with a length of 20 turns. In this setup, each LLM (A and B) contributed 10 responses based on the evolving conversation history.

Step 2-1: Intention Generation and Reply Generation. We provided the initial turn’s intention for LLM A:

Intention: In my next response, I aim to outline potential trade-offs that both I and Robert might consider in order to reach a mutually beneficial agreement. I will also suggest ways we can ensure that the partnership aligns with their core values of sustainability and operational efficiency. This will help set the stage for a constructive dialogue that addresses our respective concerns and priorities.

Reply of Player A based on the intention: Thank you for meeting with me, Robert. I believe we have a unique opportunity to collaborate that could benefit both EcoWave and Steele Industries. I understand your concerns about the costs and risks of integrating new technologies, but I think we can explore some trade-offs. For instance, if we can demonstrate the long-term cost savings and efficiency gains from adopting sustainable practices, it might help alleviate some of your concerns. Additionally, I’m open to discussing a phased implementation approach that allows for gradual integration without disrupting your current operations. How do you feel about that, and what trade-offs might you be willing to consider to make this partnership work?

Step 2-2: Alternative Generation. The LLM also generated three alternative responses based on the aforementioned intention.

Alternative 1: I appreciate your concerns, Robert, and it’s great to see your commitment to a smooth transition. Instead of jumping straight into a pilot project, how about we organize a joint workshop with your key team members? This would allow us to collaboratively explore areas where EcoWave’s technologies can fit into your existing operations. We can analyze potential impacts, address any workforce concerns, and identify specific processes that could benefit from our innovations. This way, we ensure that everyone is on board and feels included in the transition. What are your thoughts on setting up a workshop to discuss this further?

Alternative 2: I appreciate your concerns, Robert, and I think it’s important for us to find a

starting point that feels comfortable for you. Instead of a full-scale integration right away, how about we explore a collaborative approach where we can jointly assess some specific production lines that may benefit from our technologies? We could conduct a thorough analysis of these areas together, identifying potential efficiency improvements and cost benefits through a detailed study. This way, we can base our decisions on solid data while ensuring that your existing operations remain stable. Additionally, I'd be more than willing to work closely with your team to address any training needs that may arise. Would this kind of collaborative assessment resonate with you?

Alternative 3: I appreciate your insights, Robert, and I recognize the importance of ensuring a smooth transition. Instead of jumping straight into a pilot project, how about we start by conducting a joint assessment of your current operations? This would allow us to pinpoint specific areas where EcoWave's technologies could fit seamlessly. By collaborating on this assessment, we can address your concerns and tailor our solutions to your unique processes. Additionally, I'd be happy to share case studies from other manufacturers who have successfully integrated our technologies, demonstrating the benefits they achieved without significant disruptions. What do you think about this collaborative approach?

Step 3: Evaluation. Each reply and its alternatives were evaluated using G-eval.

Experiment Setting. We set the maximum number of turns to 10 and evaluated two language models: GPT-4-Turbo and OpenAI's o1 model. For each model, we tested performance across 10 different scenarios.

Results. Interestingly, both models demonstrated no-dynamic-regret behavior within the regression framework. Both models also exhibited low p -values, with GPT-4-Turbo achieving 0.09 and OpenAI o1 achieving 0.02. The detailed results are shown in Figure C.10.




Figure C.10: We evaluate GPT-4-Turbo and OpenAI o1 on two-player Negotiation, where we can see both models can achieve no-regret in the regression framework.

D Deferred Results and Proofs in Section 4

D.1 Pre-Trained LLMs Have Similar Regret as Humans (Who Generate Data)

We first provide a direct observation based on some existing speculation on the capability of Transformer-based LLMs. Recently, a growing literature has evidenced that the intelligence level of LLM agents are determined by, and in fact mimic, those of human beings who generate the data for pre-training the models (Park et al., 2022; Argyle et al., 2023; Horton, 2023). The key rationale was that, LLMs (with Transformer parameterization) can approximate the *pre-training data distribution* very well (Xie et al., 2022; Zhang et al., 2023b; Lee et al., 2023). In such a context, one can expect that LLM agents can achieve similar regret as human decision-makers who generate the pre-training data, as we formally state below.

Observation 1. An LLM agent is said to be pre-trained with an ϵ -decision error if, for any arbitrary t and loss sequences $(\ell_i)_{i \in [t]}$, the following condition holds:

$$\sup_{\pi \in \Pi} |P_{data}(\pi | (\ell_i)_{i \in [t]}) - P_{LLM}(\pi | (\ell_i)_{i \in [t]})| \leq \epsilon,$$

where P_{data} and P_{LLM} are the pre-training data distribution and the decision policy distribution of the pre-trained LLM, respectively. Then, the regret of an LLM agent with ϵ -decision error is bounded as:

$$(D-)Regret_{LLM}((\ell_t)_{t \in [T]}) \in \left[(D-)Regret_{data}((\ell_t)_{t \in [T]}) \pm \epsilon \|\ell_t\| \sup_{\pi \in \Pi} \|\pi\| \right],$$

where $[a \pm b] := [a - b, a + b]$.

Observation 1 shows that the pre-trained LLM-agent's regret can be controlled by that of the pre-training dataset and the decision error ϵ . A small ϵ can be achieved if LLM is constructed by a rich function class, e.g., the Transformer architecture (Zhang et al., 2023b; Lin et al., 2024).

Proof of Observation 1. For given $(\ell_t)_{t \in [T]}$,

$$\sum_{t=1}^T \int_{\pi_t \in \Pi} P_{LLM}(\pi_t | (\ell_i)_{i \in [t-1]}) \langle \ell_t, \pi_t \rangle d\pi_t \leq \sum_{t=1}^T \int_{\pi_t \in \Pi} (P_{data}(\pi_t | (\ell_i)_{i \in [t-1]}) + \epsilon) \langle \ell_t, \pi_t \rangle d\pi_t$$

holds, where we use the convention of $P_{LLM}(\pi_t | (\ell_0)) := P_{LLM}(\pi_t)$ and $P_{data}(\pi_t | (\ell_0)) := P_{data}(\pi_t)$. Hence,

$$\begin{aligned} \text{Regret}_{LLM}((\ell_t)_{t \in [T]}) &= \sum_{t=1}^T \int_{\pi_t \in \Pi} P_{LLM}(\pi_t | (\ell_i)_{i \in [t-1]}) \langle \ell_t, \pi_t \rangle d\pi_t - \inf_{\pi \in \Pi} \sum_{t=1}^T \langle \ell_t, \pi \rangle \\ &\leq \sum_{t=1}^T \int_{\pi_t \in \Pi} (P_{data}(\pi_t | (\ell_i)_{i \in [t-1]}) + \epsilon) \langle \ell_t, \pi_t \rangle d\pi_t - \inf_{\pi \in \Pi} \sum_{t=1}^T \langle \ell_t, \pi \rangle \\ &= \sum_{t=1}^T \int_{\pi_t \in \Pi} (P_{data}(\pi_t | (\ell_i)_{i \in [t-1]})) \langle \ell_t, \pi_t \rangle d\pi_t - \inf_{\pi \in \Pi} \sum_{t=1}^T \langle \ell_t, \pi \rangle + \sum_{t=1}^T \int_{\pi_t \in \Pi} \langle \ell_t, \epsilon \pi_t \rangle d\pi_t \\ &\leq \text{Regret}_{data}((\ell_t)_{t \in [T]}) + \epsilon \|\ell\|_p \|\pi\|_q T \end{aligned}$$

where $\frac{1}{p} + \frac{1}{q} = 1$ and $p, q \geq 1$. Similarly, we can establish the lower bound for $\text{Regret}_{LLM}((\ell_t)_{t \in [T]})$. To prove the result for the dynamic-regret case, we can simply change the term $\inf_{\pi \in \Pi} \sum_{t=1}^T \langle \ell_t, \pi \rangle$ in the above derivation to $\sum_{t=1}^T \inf_{\pi \in \Pi} \langle \ell_t, \pi \rangle$. \square

D.2 Background and Motivations for (Generalized) Quantal Response

Formally, the quantal response is defined as follows:

Definition D.1 (Quantal response). *Given a loss vector $\ell \in \mathbb{R}^d$, a noise distribution $\epsilon \sim P_{noise}$, and $\eta > 0$, the quantal response is defined as*

$$P_{quantal}^\eta(a | \ell) = \mathbb{P} \left(a \in \arg \min_{a' \in A} z(a') \right), \quad \text{where } z = \ell + \eta \epsilon.$$

In essence, this implies that humans are rational but with respect to (w.r.t.) the latent variable z , a perturbed version of ℓ , instead of ℓ per se. This addition of noise to the actual loss vector characterizes the bounded rationality of humans in decision-making.

Further motivations for generalized quantal response. Note that a *dynamic* version of quantal response in Definition 4.1 also has implications from behavior economics, and has been recently used to model human behaviors in sequential decision-making (Ding et al., 2022) (in stochastic and stationary environments). Indeed, such a response against multiple loss vectors is believed to be natural, and has also been widely adopted in well-known no-regret learning algorithms of *smooth/stochastic fictitious play* (Fudenberg & Kreps, 1993) and *follow-the-perturbed-leader* (Kalai & Vempala, 2005), whose formal definitions can be found in Appendix B.2. Finally, note that the response model in Definition 4.1 does not necessarily involve a *sequential* decision-making process, i.e., the set of losses may not come from the history of an online learning process.

D.3 The Example Instantiating Assumption 1

Example 1 (An example instantiating Assumption 1). We consider a common decision-making task that may generate the training data, recommender systems. An instance of the text data could be: “On September 29, 2023, user X clicked movie A three times, movie B eight times, and movie C five times”. This sentence corresponds to $x_{N_{i-1}+1:N_i}$ for some $i \in [t]$ and serves as a natural language depiction of the numerical ℓ_i . The corresponding label $x_{N_t+1:N_{t+1}}$ can be obtained by some user survey: “User X’s favorite movie is movie B”. Meanwhile, z represents user X’s latent, genuine preference for each movie – information that is private to the user, and cannot be observed or collected in the pre-training dataset. In this example, Assumption 1 suggests that $x_{1:N_t}$, which records the frequency of interactions with each movie, serves as an imperfect estimate of the user’s latent, genuine preference for the movies, while the text $x_{N_t+1:N_{t+1}}$ depicts the user’s favorite movie only based on her latent z .

D.4 Alignment of Assumption 1 with Quantal Response

Before presenting the technical lemma, based on Assumption 1, we denote the (potentially unknown) mappings that decode semantic information in Assumption 1 into numeric values as f , g , such that $f(x_{N_{i-1}+1:N_i}) = \ell_i \in \mathbb{R}^d$ for each $i \in [t]$ and $g(x_{N_t+1:N_{t+1}}) = a \in \mathcal{A}$.

Lemma 1. Fix $t \in [T]$, $\sigma > 0$. If we model the noise of data collection to be i.i.d. Gaussian distribution in the numeric value space, i.e.,

$$\mathbb{P}\left(\left\{f(x_{N_{i-1}+1:N_i})\right\}_{i \in [t]} \mid z\right) = \prod_{i=1}^t \mathbb{P}\left(f(x_{N_{i-1}+1:N_i}) \mid z\right) \propto \prod_{i=1}^t \exp\left(-\frac{\|f(x_{N_{i-1}+1:N_i}) - z\|_2^2}{2\sigma^2}\right),$$

the prior distribution of the latent variable z is also Gaussian, i.e., $z \sim \mathcal{N}(\mathbf{0}_d, \sigma^2 I)$, and the text labels satisfy that $\mathbb{P}(g(x_{N_t+1:N_{t+1}}) \mid z) = \mathbb{1}\left(g(x_{N_t+1:N_{t+1}}) \in \arg \min_{a \in \mathcal{A}} z_a\right)$, then we have

$$\mathbb{P}\left(g(x_{N_t+1:N_{t+1}}) \mid x_{1:N_t}\right) = P_{\text{quantal}}^{\sigma\sqrt{t+1}}\left(g(x_{N_t+1:N_{t+1}}) \mid \left\{f(x_{N_{i-1}+1:N_i})\right\}_{i \in [t]}\right),$$

with $P_{\text{noise}} = \mathcal{N}(\mathbf{0}_d, I)$ in Definition 4.1, i.e., the action $a = g(x_{N_t+1:N_{t+1}})$ extracted from the text $x_{N_t+1:N_{t+1}}$ is a quantal response w.r.t. the loss vectors $\left(f(x_{N_{i-1}+1:N_i})\right)_{i \in [t]}$.

Proof. Note that

$$\mathbb{P}(z \mid x_{1:N_t}) = \int_{\ell_{1:t}} \mathbb{P}(z, \ell_{1:t} \mid x_{1:N_t}) d\ell_{1:t} = \int_{\ell_{1:t}} \mathbb{P}(\ell_{1:t} \mid x_{1:N_t}) \mathbb{P}(z \mid x_{1:N_t}, \ell_{1:t}) d\ell_{1:t}.$$

For $\mathbb{P}(\ell_{1:t} | x_{1:N_t})$, since we have assumed the existence of function f to decode $\ell_{1:t}$ from $x_{1:N_t}$, it holds that

$$\mathbb{P}(\ell_{1:t} | x_{1:N_t}) = \prod_{i=1}^t \delta(\ell_i - f(x_{N_{i-1}+1:N_i})),$$

where we use δ to denote the d -dimensional Dirac-delta function. For $\mathbb{P}(z | x_{1:N_t}, \ell_{1:t})$, by Assumption 1, it holds that

$$\mathbb{P}(z, x_{1:N_t}, \ell_{1:t}) = \mathbb{P}(z, \ell_{1:t}) \mathbb{P}(x_{1:N_t} | \ell_{1:t}),$$

which leads to $\mathbb{P}(x_{1:N_t} | \ell_{1:t}) = \mathbb{P}(x_{1:N_t} | \ell_{1:t}, z)$ by Bayes rule. This implies that the random variable $x_{1:N_t}$ and z are independent conditioned on $\ell_{1:t}$. Therefore, it holds that $\mathbb{P}(z | x_{1:N_t}, \ell_{1:t}) = \mathbb{P}(z | \ell_{1:t})$. Finally, we can compute

$$\begin{aligned} \mathbb{P}(z | x_{1:N_t}) &= \int_{\ell_{1:t}} \mathbb{P}(z, \ell_{1:t} | x_{1:N_t}) d\ell_{1:t} = \int_{\ell_{1:t}} \prod_{i=1}^t \delta(\ell_i - f(x_{N_{i-1}+1:N_i})) \mathbb{P}(z | \ell_{1:t}) d\ell_{1:t} \\ &= \mathbb{P}\left(z \mid \left(\ell_i = f(x_{N_{i-1}+1:N_i})\right)_{i \in [t]}\right). \end{aligned}$$

Based on this, we conclude that

$$\begin{aligned} \mathbb{P}(g(x_{N_t+1:N_{t+1}}) | x_{1:N_t}) &= \int_z \mathbb{P}(g(x_{N_t+1:N_{t+1}}) | z, x_{1:N_t}) \mathbb{P}(z | x_{1:N_t}) dz \\ &= \int_z \mathbb{P}(g(x_{N_t+1:N_{t+1}}) | z) \mathbb{P}(z | \{\ell_i = f(x_{N_{i-1}+1:N_i})\}_{i \in [t]}) dz \\ &= \mathbb{P}\left(g(x_{N_t+1:N_{t+1}}) \mid \left(\ell_i = f(x_{N_{i-1}+1:N_i})\right)_{i \in [t]}\right) \end{aligned}$$

where the first equality is by the independence between $x_{N_t+1:N_{t+1}}$ and $x_{1:N_t}$ conditioned on z , due to Assumption 1. Therefore, it suffices to consider the probability of $\mathbb{P}(a | \ell_{1:t})$ only, in order to analyze $\mathbb{P}(g(x_{N_t+1:N_{t+1}}) | x_{1:N_t})$, where we recall the definition that $a = g(x_{N_t+1:N_{t+1}})$. Since $z \sim \mathcal{N}(\mathbf{0}_d, \sigma^2 I)$, and $\ell_i | z \sim \mathcal{N}(z, \sigma^2 I)$, we have

$$z | \ell_{1:t} \sim \mathcal{N}\left(\frac{1}{t+1} \sum_{i \in [t]} \ell_i, \frac{\sigma^2}{t+1} I\right), \quad (\text{D.1})$$

by the posterior distribution of Gaussian distribution. Now we conclude that

$$\begin{aligned} \mathbb{P}(a | \ell_{1:t}) &= \int_z \mathbb{P}(a | z, \ell_{1:t}) \mathbb{P}(z | \ell_{1:t}) dz = \int_z \mathbb{P}(a | z) \mathbb{P}(z | \ell_{1:t}) dz \\ &= \int_z \mathbb{1}(a \in \arg \min_{a' \in \mathcal{A}} z_{a'}) \mathbb{P}(z | \ell_{1:t}) dz = \int_z \mathbb{1}\left(a \in \arg \min_{a' \in \mathcal{A}} \left(\frac{\sigma}{\sqrt{t+1}} \epsilon + \frac{1}{t+1} \sum_{i \in [t]} \ell_i\right)_{a'}\right) \mathbb{P}(\epsilon) d\epsilon \\ &= \int_z \mathbb{1}\left(a \in \arg \min_{a' \in \mathcal{A}} \left(\sigma \sqrt{t+1} \epsilon + \sum_{i \in [t]} \ell_i\right)_{a'}\right) \mathbb{P}(\epsilon) d\epsilon = \mathbb{P}\left(a \in \arg \min_{a' \in \mathcal{A}} \left(\sigma \sqrt{t+1} \epsilon + \sum_{i \in [t]} \ell_i\right)_{a'}\right) \\ &= P_{\text{quantal}}^{\sigma \sqrt{t+1}}(a | \ell_{1:t}), \end{aligned}$$

where $\mathbb{P}(\epsilon) = \mathcal{N}(\mathbf{0}_d, I)$. This completes the proof. \square

D.5 Relationship between FTPL and Definition 4.1

Fact 1. Performing generalized quantal response of Definition 4.1 at every iteration $t \in [T]$ w.r.t. history loss vectors $\ell_{1:t-1}$ is essentially executing an FTPL algorithm.

Proof. Before we move to the proof, we will define the random variable which has distribution P_{noise} as Z_{noise} . Note that at round $t \geq 2$ (as the policy at round $t = 1$ is fixed), we have

$$P_{\text{quantal}}^{\eta_{t-1}}(a | \ell_{1:t-1}) := \mathbb{P} \left(a \in \arg \min_{a' \in \mathcal{A}} \left(\sum_{i=1}^{t-1} \ell_i + \eta_{t-1} \epsilon \right) (a') \right) \quad (\text{D.2})$$

which is exactly the case when ϵ_t in Equation (B.1) satisfies $\epsilon_t \stackrel{d}{=} \eta_{t-1} \epsilon$. \square

D.6 Formal Statement and Proof of Theorem 4.1

Theorem D.1. (Emergence of no-regret behavior). *Under the assumptions of Lemma 1, suppose the function class of LLM_θ is expressive enough such that for all $t \in [T]$, $\max_{\theta \in \Theta} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} \sum_{j=1}^{N_{t+1}} \log \text{LLM}_\theta(x_j | x_{1:j-1}) = \max_{\{q_j \in \{\mathcal{V}^{j-1} \rightarrow \Delta(\mathcal{V})\}\}_{j \in [N_{t+1}]}} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} \sum_{j=1}^{N_{t+1}} \log q_j(x_j | x_{1:j-1})$, where we define $q_1(x_1 | x_{1:0}) := q_1(x_1)$, and θ^\star maximizes Equation (4.1). Then, there exist (simple) algorithms using $\text{LLM}_{\theta^\star}$ to achieve no (dynamic) regret for (non-stationary) online learning with full-information/bandit feedback. To be specific, for (2) and (4), by defining the variation bound $\sum_{t=1}^{T-1} \|\ell_{t+1} - \ell_t\|_\infty \leq V_T$ such that $V_T \leq T$ and $V_T = \Theta(T^\rho)$ for some $\rho \in (0, 1)$, it holds that for large enough T, d :*

- (1) For online learning with full-information feedback, $\text{Regret}_{\text{LLM}_{\theta^\star}}((\ell_t)_{t \in [T]}) \leq \mathcal{O}(\sqrt{T \log d})$;
- (2) For non-stationary online learning with full-information feedback, $D\text{-Regret}_{\text{LLM}_{\theta^\star}}((\ell_t)_{t \in [T]}) \leq \mathcal{O}((\log d V_T)^{1/3} T^{2/3})$;
- (3) For online learning with bandit feedback, $\mathbb{E}[\text{Regret}_{\text{LLM}_{\theta^\star}}((\ell_t)_{t \in [T]})] \leq \mathcal{O}((\log d)^{1/2} d T^{1/2+1/\log T} \log T)$;
- (4) For non-stationary online learning with bandit feedback, $\mathbb{E}[D\text{-Regret}_{\text{LLM}_{\theta^\star}}((\ell_t)_{t \in [T]})] \leq \mathcal{O}((T^2 d^2 V_T)^{1/3} (\log d)^{1/2} T^{1/\log T} \log T)$.

Proof. Note that

$$\begin{aligned} & \max_{\{q_j \in \{\mathcal{V}^{j-1} \rightarrow \Delta(\mathcal{V})\}\}_{j \in [N_{t+1}]}} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} \sum_{j=1}^{N_{t+1}} \log q_j(x_j | x_{1:j-1}) \\ &= \max_{q \in \Delta(\mathcal{V}^{N_{t+1}})} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} \log q(x_{1:N_{t+1}}) \\ &= \max_{q \in \Delta(\mathcal{V}^{N_{t+1}})} -\text{KL}(P_t^{\text{text}} || q) + \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} [P_t^{\text{text}}(x_{1:N_{t+1}})], \end{aligned}$$

where $\text{KL}(q || p)$ denotes the Kullback–Leibler divergence between two distributions p, q . Now we define $\text{LLM}_\theta(x_{1:N_{t+1}}) = \prod_{t=1}^{N_{t+1}} \text{LLM}_\theta(x_j | x_{1:j-1})$. It is easy to verify that $\text{LLM}_\theta(x_{1:N_{t+1}}) \in \Delta(\mathcal{V}^{N_{t+1}})$, i.e., it also defines a valid joint distribution over tokens. Therefore, we have

$$\max_{\theta \in \Theta} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} \sum_{j=1}^{N_{t+1}} \log \text{LLM}_\theta(x_j | x_{1:j-1}) = \max_{\theta \in \Theta} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{\text{text}}} \log \text{LLM}_\theta(x_{1:N_{t+1}}).$$

Now, due to our assumption that

$$\begin{aligned} & \max_{\theta \in \Theta} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{text}} \sum_{j=1}^{N_{t+1}} \log \text{LLM}_{\theta}(x_j | x_{1:j-1}) \\ &= \max_{\{q_j \in \{\mathcal{V}^{j-1} \rightarrow \Delta(\mathcal{V})\}\}_{j \in [N_{t+1}]}} \mathbb{E}_{x_{1:N_{t+1}} \sim P_t^{text}} \sum_{j=1}^{N_{t+1}} \log q_j(x_j | x_{1:j-1}), \end{aligned}$$

we conclude that

$$\min_{\theta \in \Theta} \text{KL}(P_t^{text} || \text{LLM}_{\theta}) = \min_{q \in \Delta(\mathcal{V}^{N_{t+1}})} \text{KL}(P_t^{text} || q) = 0,$$

which implies that $\text{LLM}_{\theta^*} = P_t^{text}$. Correspondingly, if we define $\text{LLM}_{\theta^*}(x_{N_t+1:N_{t+1}} | x_{1:N_t})$ to be the distribution induced by the joint distribution $\text{LLM}_{\theta^*}(x_{1:N_{t+1}})$, it holds that

$$\text{LLM}_{\theta^*}(x_{N_t+1:N_{t+1}} | x_{1:N_t}) = \mathbb{P}(x_{N_t+1:N_{t+1}} | x_{1:N_t}).$$

In other words, intuitively, LLM_{θ^*} has learned the corresponding *pre-training* distribution perfectly. Note that this has been a common assumption in the Bayesian perspective of ICL (Xie et al., 2022; Lee et al., 2023; Zhang et al., 2023b). Therefore, to analyze the actions taken by LLM_{θ^*} , it suffices to consider $\mathbb{P}(g(x_{N_t+1:N_{t+1}}) | x_{1:N_t})$, which is equal to $P_{quantal}^{\sigma \sqrt{t+1}}(g(x_{N_t+1:N_{t+1}}) | \{f(x_{N_{i-1}+1:N_i})\}_{i \in [t]})$ by Lemma 1. Therefore, we proved that LLM_{θ^*} is essentially mimicking the well-known no-regret algorithm, FTPL with perturbation distribution as $\mathcal{N}(\mathbf{0}_d, \sigma^2 t I)$ for round $t \in [T]$, according to Equation (D.2) of Fact 1, for which we can establish the corresponding regret guarantee for each case:

(1) Combining the above result with Lemma 2, we can derive the regret bound for online learning with full-information feedback.

(2) Combining the above result with Lemma 2 and Lemma 4, we get that

$$\text{D-Regret}_{\text{LLM}_{\theta^*}}((\ell_i)_{i \in [T]}) \leq \min_{\Delta_T \in [T]} \frac{2T}{\Delta_T} C \sqrt{\Delta_T \log d} + 2\Delta_T V_T,$$

for some constant C . We firstly consider the following problem

$$\min_{u > 0} \frac{2T}{u} C \sqrt{u \log d} + 2u V_T,$$

where the optimal solution is $u^* = \left(\frac{C^2 T^2 \log d}{4V_T^2} \right)^{1/3}$. Therefore, if we have $u^* \in [1, T]$, we can choose $\Delta_T = \lceil u^* \rceil$, which results in a regret bound of

$$\text{D-Regret}_{\text{LLM}_{\theta^*}}((\ell_i)_{i \in [T]}) \leq \frac{2T}{\sqrt{u^*}} C \sqrt{\log d} + 4u^* V_T = \mathcal{O}\left((\log d V_T)^{1/3} T^{2/3}\right).$$

Now we check the conditions for $u^* \in [1, T]$. It is direct to see that since $V_T \leq T$, $u^* \geq 1$ holds as long as d is sufficiently large. To ensure $u^* \leq T$, we get the condition $V_T \geq C \sqrt{\frac{\log d}{4T}}$, which holds as long as T is large enough.

(3) Combining the above result with Lemma 3, we can prove a regret guarantee for online learning with bandit feedback.

(4) Combining this result with Lemma 3 and Lemma 4, it holds that

$$\mathbb{E}[\text{D-Regret}_{\text{LLM}_{\theta^*}}((\ell_i)_{i \in [T]})] \leq \min_{\Delta_T \in [T]} \frac{2T}{\Delta_T} C (\log d)^{\frac{1}{2}} d \Delta_T^{\frac{1}{2} + \frac{1}{\log T}} \log \Delta_T + 2\Delta_T V_T,$$

for some constant C . By adopting a similar analysis as that of (2), we choose $u^* = \left(\frac{C'T^2 d^2}{V_T^2}\right)^{1/3}$ for some constant C' . If $u^* \in [1, T]$, we choose $\Delta_T = \lceil u^* \rceil$ and derive the following regret:

$$\mathbb{E}[\text{D-Regret}_{\text{LLM}_{\theta^*}}((\ell_i)_{i \in [T]})] \leq \mathcal{O}\left((T^2 d^2 V_T)^{1/3} (\log d)^{1/2} T^{1/\log T} \log T\right).$$

Now we check the condition of $u^* \in [1, T]$. Note that since $V_T \leq T$, $u^* \geq 1$ holds as long as d is sufficiently large. For $u^* \leq T$, we have $V_T \geq \sqrt{\frac{C'd^2}{T}}$, which holds as long as T is large enough.

Now, we present Lemma 2 - Lemma 4. Before proceeding, we assume $\|\ell_t\|_\infty \leq B = 1$ for simplicity of presentations hereafter. The results and proof are not affected by the constant bound B .

Lemma 2 (Regret guarantee of FTPL with full-information feedback). *Suppose the noise distribution of FTPL satisfies that $\epsilon_t \sim \mathcal{N}(\mathbf{0}_d, \zeta_t^2 I)$ in Equation (B.1) and $\zeta_t = \sigma\sqrt{t}$, then for online learning with full-information feedback,*

$$\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]}) \leq 4\left(\sigma + \frac{1}{\sigma}\right)\sqrt{T \log d} = \mathcal{O}(\sqrt{T \log d}).$$

Proof. By Theorem 8 of Abernethy et al. (2014), we have

$$\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]}) \leq \sqrt{2 \log d} \left(\eta_T + \sum_{t=1}^T \frac{1}{\eta_t} \|\ell_t\|_\infty^2 \right).$$

Therefore, plugging $\zeta_t = \sigma\sqrt{t}$ and $\|\ell_t\|_\infty^2 \leq 1$ provides

$$\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]}) \leq \sqrt{2 \log d} \left(\sigma\sqrt{T} + \sum_{t=1}^T \frac{1}{\sigma\sqrt{t}} \right) \leq 4\left(\sigma + \frac{1}{\sigma}\right)\sqrt{T \log d},$$

completing the proof. \square

Lemma 3 (Regret guarantee of FTPL with bandit feedback). *Suppose the noise distribution of FTPL satisfies that $\epsilon_t \sim \mathcal{N}(\mathbf{0}_d, \zeta_t^2 I)$ in Equation (B.1) and $\zeta_t = \sigma\sqrt{t}$, then for online learning with bandit feedback,*

$$\mathbb{E}[\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]})] \leq \mathcal{O}((\log d)^{\frac{1}{2}} d T^{\frac{1}{2} + \frac{1}{\log T}} \log T).$$

Proof. The proof of the bandit problem is more complex. We first define the following notation. We denote $G_t = \sum_{t'=1}^t -\ell_{t'}$, $\widehat{G}_t = \sum_{t'=1}^t -\widehat{\ell}_{t'}$, $\Phi(G) = \max_{\pi} \langle \pi, G \rangle$, $\Phi_t(G) = \mathbb{E}_{\epsilon \sim \mathcal{N}(\mathbf{0}_d, I)} \Phi(G + \zeta_t \epsilon)$, and D_{Φ_t} to be the Bregman divergence with respect to Φ_t , where we recall the construction of the empirical estimator $\widehat{\ell}_{t'}$ of $\ell_{t'}$ in Section 3.2. By Li & Tewari (2017), $\pi_t = \nabla \Phi_t(\widehat{G}_t)$. Now due to the convexity of Φ ,

$$\Phi(G_T) = \Phi(\mathbb{E}[\widehat{G}_T]) \leq \mathbb{E}[\Phi(\widehat{G}_T)].$$

Therefore,

$$\mathbb{E}[\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]})] = \Phi(G_T) - \mathbb{E}\left[\sum_{t=1}^T \langle \pi_t, -\widehat{\ell}_t \rangle\right] \leq \mathbb{E}\left[\Phi(\widehat{G}_T) - \sum_{t=1}^T \langle \pi_t, -\widehat{\ell}_t \rangle\right].$$

By recalling the definition of the Bregman divergence, we have

$$\begin{aligned} -\sum_{t=1}^T \langle \pi_t, -\widehat{\ell}_t \rangle &= -\sum_{t=1}^T \langle \nabla \Phi_t(\widehat{G}_t), -\widehat{\ell}_t \rangle = -\sum_{t=1}^T \langle \nabla \Phi_t(\widehat{G}_t), \widehat{G}_t - \widehat{G}_{t-1} \rangle \\ &= \sum_{t=1}^T D_{\Phi_t}(\widehat{G}_t, \widehat{G}_{t-1}) + \Phi_t(\widehat{G}_{t-1}) - \Phi_t(\widehat{G}_t). \end{aligned}$$

Therefore,

$$\begin{aligned} \mathbb{E}[\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]})] &\leq \underbrace{\mathbb{E}\left[\sum_{t=1}^T D_{\Phi_t}(\widehat{G}_t, \widehat{G}_{t-1})\right]}_{(i)} + \underbrace{\mathbb{E}\left[\sum_{t=1}^T \Phi_t(\widehat{G}_{t-1}) - \Phi_{t-1}(\widehat{G}_{t-1})\right]}_{(ii)} + \underbrace{\mathbb{E}[\Phi(\widehat{G}_T) - \Phi_T(\widehat{G}_T)]}_{(iii)}, \end{aligned}$$

where $(iii) \leq 0$ due to the convexity of Φ . For (ii) , we use Lemma 10 of [Abernethy et al. \(2014\)](#) to obtain

$$\mathbb{E}\left[\sum_{t=1}^T \Phi_t(\widehat{G}_{t-1}) - \Phi_{t-1}(\widehat{G}_{t-1})\right] \leq \zeta_T \mathbb{E}_\epsilon[\Phi(\epsilon)] \leq \mathcal{O}(\sqrt{2T \log d}).$$

For (i) , by Theorem 8 of [Li & Tewari \(2017\)](#), for any $\alpha \in (0, 1)$, the following holds:

$$\begin{aligned} \mathbb{E}\left[\sum_{t=1}^T D_{\Phi_t}(\widehat{G}_t, \widehat{G}_{t-1})\right] &\leq \sum_{t=1}^T \zeta_t^{\alpha-1} \frac{4d}{\alpha(1-\alpha)} \\ &\leq \frac{4d}{\alpha(1-\alpha)} \mathcal{O}(T^{\frac{1+\alpha}{2}}). \end{aligned}$$

By tuning $\alpha = \frac{2}{\log T}$, we proved that $\mathbb{E}[\text{Regret}_{\text{FTPL}}((\ell_i)_{i \in [T]})] \leq \mathcal{O}((\log d)^{\frac{1}{2}} d T^{\frac{1}{2} + \frac{1}{\log T}} \log T)$. \square

Lemma 4. Denote the variation of loss vectors as $L_T = \sum_{t=1}^{T-1} \|\ell_{t+1} - \ell_t\|_\infty$. Suppose there exists an algorithm \mathcal{A} for online learning with full-information feedback with regret guarantee that $\text{Regret}_{\mathcal{A}}((\ell_i)_{i \in [T]}) \leq f(T, d)$ for some function f , where T denotes the horizon and d denotes the policy dimension. Then, there exists another algorithm \mathcal{A}' that can achieve

$$D\text{-Regret}_{\mathcal{A}'}((\ell_i)_{i \in [T]}) \leq \min_{\Delta_T \in [T]} \left(\frac{T}{\Delta_T} + 1 \right) f(\Delta_T, d) + 2\Delta_T L_T.$$

Similarly, suppose there exists an algorithm \mathcal{B} for online learning with bandit feedback with regret guarantee that $\mathbb{E}[\text{Regret}_{\mathcal{B}}((\ell_i)_{i \in [T]})] \leq g(T, d)$ for some function g ; then there exists another algorithm \mathcal{B}' that can achieve

$$\mathbb{E}[D\text{-Regret}_{\mathcal{B}'}((\ell_i)_{i \in [T]})] \leq \min_{\Delta_T \in [T]} \left(\frac{T}{\Delta_T} + 1 \right) g(\Delta_T, d) + 2\Delta_T L_T.$$

Proof. We denote \mathcal{A}' as the algorithm that restarts \mathcal{A} every Δ_T iterations. We break the time index $[T]$ into m batches $\mathcal{T}_{1:m}$ of size Δ_T (except for, possibly the last batch). Denote $\ell_i^\star := \min_{j \in [d]} \ell_{ij}$. By Equation (6) of [Besbes et al. \(2014\)](#), it holds that for each $k \in [m]$

$$\min_{j \in [d]} \left(\sum_{t \in \mathcal{T}_k} \ell_t \right)_j - \sum_{t \in \mathcal{T}_k} \ell_t^\star \leq 2\Delta_T L_k,$$

where we define $L_k = \sum_{t \in \mathcal{T}_k} \|\ell_{t+1} - \ell_t\|_\infty$. Therefore, we have

$$\begin{aligned} \text{D-Regret}_{\mathcal{A}'}((\ell_i)_{i \in [T]}) &\leq \min_{j \in [d]} \left(\sum_{t \in [T]} \ell_t \right)_j - \sum_{t \in [T]} \ell_t^\star + \sum_{k \in [m]} \text{Regret}_{\mathcal{A}}((\ell_i)_{i \in [\mathcal{T}_k]}) \\ &\leq 2\Delta_T \left(\sum_{k \in [m]} L_k \right) + (T/\Delta_T + 1)g(\Delta_T, d). \end{aligned} \quad (\text{D.3})$$

By Equation (4) of [Besbes et al. \(2014\)](#) that $\sum_{k \in [m]} L_k \leq L_T$ and this inequality holds for any $\Delta_T \in [T]$, we proved $\text{D-Regret}_{\mathcal{A}'}((\ell_i)_{i \in [T]}) \leq \min_{\Delta_T \in [T]} \left(\frac{T}{\Delta_T} + 1 \right) f(\Delta_T, d) + 2\Delta_T L_T$.

Similarly, if we take the expectation for Equation (D.3), it holds that

$$\begin{aligned} \mathbb{E}[\text{D-Regret}_{\mathcal{B}'}((\ell_i)_{i \in [T]})] &\leq \min_{j \in [d]} \left(\sum_{t \in [T]} \ell_t \right)_j - \sum_{t \in [T]} \ell_t^\star + \sum_{k \in [m]} \mathbb{E}[\text{Regret}_{\mathcal{B}}((\ell_i)_{i \in [\mathcal{T}_k]})] \\ &\leq \min_{\Delta_T \in [T]} \left(\frac{T}{\Delta_T} + 1 \right) g(\Delta_T, d) + 2\Delta_T L_T, \end{aligned}$$

thus completing the proof. \square

Combining the results above completes the proof for Theorem 4.1. \square

D.6.1 Implications of Theorem 4.1 for Repeated Games

Remark D.1 (Implication for playing repeated games). *First, we note that the no-regret guarantee in the online setting is stronger than and thus implies that in the game setting, since regret by definition handles arbitrary/adversarial environments, while in playing games the opponents are not necessarily as adversarial. Second, it is folklore that if all players in the repeated game follow no-regret learning algorithms, then the time-average policies of all players during learning constitute an approximate **coarse correlated equilibrium** of the game ([Cesa-Bianchi & Lugosi, 2006](#)). Hence, the results (1) and (2) in Theorem 4.1 imply that a coarse correlated equilibrium will emerge in the long run from the interactions of the LLM agents (under certain assumptions as in the theorem).*

D.7 Extending Theorem 4.1 with Relaxed Assumptions

D.7.1 Relaxation under More General Data Distributions

We first remark on the possibility of relaxing the Gaussian assumptions on the data distributions.

Remark D.2 (Relaxing the Gaussian distribution assumption). *In the proof of Lemma 1, to obtain the result that the action is a quantal response w.r.t. $\ell_{1:T}$, one does not necessarily require both the prior distribution of z and the conditional distribution of ℓ_i given z to be Gaussian. Instead, for any joint distribution $\mathbb{P}(z, \ell_{1:T})$, as long as its posterior distribution satisfies Equation (D.1), it would suffice. It is a combined effect of both the prior and the conditional distributions.*

More formally, we can extend Theorem 4.1 to the case with a much more general prior task distribution than the Gaussian one, where the key is that Equation (D.1) only needs to hold approximately.

Theorem D.2. *In Theorem 4.1, we can relax the assumption on $\mathbb{P}(z)$ to one where we only require $\mathbb{P}(z)$ to be i.i.d for each coordinate of z and $0 < \mathbb{P}(z_j) < \infty$, $|\nabla \mathbb{P}(z_j)| < \infty$ for any $j \in [d]$, $z_j \in \mathbb{R}$, and the bounds for (1) and (2) of Theorem 4.1 still hold, with only a degradation of $\mathcal{O}(d^2 \log T)$.*

The key idea of the proof is that when t is large enough, the prior distribution does not affect the posterior distribution, which is also referred to as the *Bernstein–von Mises theorem* (Van der Vaart, 2000).

Proof. Since we extend Theorem 4.1 to settings with general task prior distribution only requiring the coordinates to be i.i.d, from now on, we consider the j -th coordinate only. To begin with, fix $t \in [T]$, we define the log-likelihood of the posterior as

$$L_t(z_j) := \log \prod_{i=1}^t \frac{1}{\sigma^d (2\pi)^{d/2}} e^{-\frac{1}{2\sigma^2}(\ell_{ij} - z_j)^2} = -t \log \sigma - \frac{t}{2} \log 2\pi - \sum_{i=1}^t \frac{1}{2\sigma^2} (\ell_{ij} - z_j)^2.$$

Then, the MLE estimator $\widehat{z}_{j,t}$ is defined as

$$\widehat{z}_{j,t} := \arg \max_{z_j \in \mathbb{R}} L_t(z_j) = \frac{1}{t} \sum_{i=1}^t \ell_{ij}.$$

We also define $\widehat{J}_t : \mathbb{R} \rightarrow \mathbb{R}$ as:

$$\widehat{J}_t(z_j) := -\frac{\nabla^2 L_t(z_j)}{t} = \frac{1}{\sigma^2}.$$

For Assumption 1 of Kasprzak et al. (2022) to hold, any $\delta > 0$, $M_2 > 0$ suffices.

For Assumption 2 of Kasprzak et al. (2022) to hold, we can choose $\widehat{M}_1 = \max_{z_j \in [-\delta, 1+\delta]} \frac{1}{\mathbb{P}(z_j)}$

For Assumption 7 of Kasprzak et al. (2022) to hold, we choose δ to be σ .

For Assumption 8 of Kasprzak et al. (2022) to hold, one can choose $M_2 = \frac{\sigma}{2}$.

For Assumption 9 of Kasprzak et al. (2022) to hold, we have

$$\kappa \leq - \sup_{(z_j - \widehat{z}_{j,t})^2 \geq \delta} \frac{L_t(z_j) - L_t(\widehat{z}_{j,t})}{t} = -\frac{1}{2\sigma^2 t} \sup_{(z_j - \widehat{z}_{j,t})^2 \geq \delta} \sum_{i=1}^t (\ell_{ij} - \widehat{z}_{j,t})^2 - (\ell_{ij} - z_j)^2 = \frac{1}{4\sigma}.$$

For Assumption 10 of Kasprzak et al. (2022) to hold, we choose $M_1 = \sup_{z_j \in [-\delta, 1+\delta]} \left| \frac{\nabla \mathbb{P}(z_j)}{\mathbb{P}(z_j)} \right|$, $\widetilde{M}_1 = \sup_{z_j \in [-\delta, 1+\delta]} |\mathbb{P}(z_j)|$ since we have assumed that $0 < \mathbb{P}(z_j) < \infty$, $|\nabla \mathbb{P}(z_j)| < \infty$.

By Theorem 6.1 of Kasprzak et al. (2022), we have

$$\begin{aligned} & \int_{z_j} |\mathbb{P}(z_j/\sqrt{t} + \widehat{z}_j) (\ell_{ij})_{i \in [t]} - C e^{-\frac{1}{2\sigma^2} z_j^2}| dz_j \\ &= \sqrt{t} \int_{z_j} |\mathbb{P}(z_j) (\ell_{ij})_{i \in [t]} - \mathcal{N}(\widehat{z}_j, \frac{\sigma^2}{t})| dz_j \leq D_1 t^{-1/2} + D_2 t^{1/2} e^{-t\kappa} + 2\widehat{\mathcal{D}}(t, \delta), \end{aligned}$$

where C is the normalization constant and

$$\begin{aligned} D_1 &= \frac{\sqrt{\widetilde{M}_1 \widehat{M}_1}}{\sigma} \left(\frac{\sqrt{3}\sigma^2}{2 \left(1 - \sqrt{\widehat{\mathcal{D}}(t, \delta)} \right)} M_2 + M_1 \right) \\ D_2 &= \frac{2\widehat{M}_1 \widehat{J}_t^p(\widehat{z}_j, \delta)}{(2\pi)^{1/2} (1 - \widehat{\mathcal{D}}^p(t, \delta))} \\ \widehat{\mathcal{D}}(t, \delta) &= e^{-\frac{1}{2}(\sqrt{t}-1)^2} \\ \widehat{J}_t^p(\widehat{z}_j, \delta) &= \frac{1}{\sigma^2} + \frac{\delta M_2}{3}. \end{aligned}$$

Therefore, we conclude that the TV distance between z (conditioned on $(\ell_i)_{i \in [t]}$) and $\mathcal{N}(\widehat{z}, \frac{\sigma^2}{t})$ satisfies that

$$\int_z \left| \mathbb{P}(z | (\ell_i)_{i \in [t]}) - \mathcal{N}\left(\widehat{z}, \frac{\sigma^2}{t}\right) \right| dz \leq \sum_{j=1}^d \int_{z_j} \left| \mathbb{P}(z_j | (\ell_{ij})_{i \in [t]}) - \mathcal{N}\left(\widehat{z}_j, \frac{\sigma^2}{t}\right) \right| dz_j \leq \mathcal{O}(d/t),$$

due to the independence of $(z_j)_{j \in [d]}$ conditioned on $\ell_{1:t}$. Now we denote algorithm $\widehat{\text{FTPL}}$ to be the FTPL algorithm w.r.t. the noise distribution $\mathbb{P}(z | (\ell_i)_{i \in [t]})$, and FTPL to be the algorithm w.r.t. the noise distribution $\mathcal{N}(\widehat{z}, \frac{\sigma^2}{t})$. Therefore, we have

$$\begin{aligned} \left| \text{Regret}_{\text{FTPL}}((\ell)_{i \in [T]}) - \text{Regret}_{\widehat{\text{FTPL}}}((\ell)_{i \in [T]}) \right| &\leq \sum_{t=1}^T d \|\pi_t - \widehat{\pi}_t\|_\infty \\ &\leq d \sum_{t=1}^T \int_z \left| \mathbb{P}(z | (\ell_i)_{i \in [t]}) - \mathcal{N}\left(\widehat{z}, \frac{\sigma^2}{t}\right) \right| dz = \mathcal{O}(d^2 \log T). \end{aligned}$$

In other words, using $\mathbb{P}(z | (\ell_i)_{i \in [t]})$ as the noise distribution only increases the regret by $\mathcal{O}(d^2 \log T)$. Similarly, it is easy to see that

$$\left| \text{D-Regret}_{\text{FTPL}}((\ell)_{i \in [T]}) - \text{D-Regret}_{\widehat{\text{FTPL}}}((\ell)_{i \in [T]}) \right| \leq \mathcal{O}(d^2 \log T),$$

which completes the proof. \square

D.7.2 Relaxation under Decision-Irrelevant Pre-Training Data

We then remark on the possible relaxation when the training data may not all come from decision-making tasks.

Remark D.3 (Pre-training with relaxed data assumptions). *Note that the pre-training (text) data are so far assumed to be related to decision-making problems (though not necessarily sequential ones), see Assumption 1 and Example 1 for instance. It can also be generalized to the text datasets involving Question-Answering (Q-A), a typical task in natural language processing, where the true/fact answer, sampled answers from different human users (with possibly wrong or biased answers), correspond to the latent z (and associated maximizer a) and $\ell_{1:t}$, respectively. Moreover, in practice, the pre-training data may also involve non-decision-making/Q-A texts, given the diversity of the datasets. For such scenarios, we will make the assumptions on the data distribution conditioned on the prompt for decision-making. Specifically, when interacting with the LLM, human users will provide prompts (see e.g., our Figure 3.1), to induce it to make decisions. This will query the conditional distribution of*

$$\mathbb{P}\left(g(x_{N_t+1:N_{t+1}}) \mid x_{1:N_t}, \text{decision-making prompt}\right)$$

to generate the control action. Correspondingly, Assumption 1 will thus only need to be made on

$$\mathbb{P}\left(z, \ell_{1:t}, x_{1:N_{t+1}}, \text{decision-making prompt}\right),$$

while we do not need to make such assumptions on other prompts, e.g., corpora that are not related to decision-making.

D.8 Comparison with Lee et al. (2023); Lin et al. (2024); Liu et al. (2023e)

Similar assumptions and pre-training objectives have also been considered in the very recent work of Lee et al. (2023); Lin et al. (2024); Liu et al. (2023e) for studying in-context reinforcement learning property of Transformers/LLM-agents under supervised pre-training. Lee et al. (2023) established its equivalence to *posterior sampling* (Osband et al., 2013), an important RL algorithm with provable regret guarantees when the environments are *stationary*, and Lin et al. (2024) generalized the study to the setting of algorithm distillation as in Laskin et al. (2023). Liu et al. (2023e) adopted the similar data generation assumption as Lee et al. (2023) without assuming optimal labels are available in the pre-training datasets, but leverages external oracles for *planning*. Consequently, the resulting LLM agent would still perform the posterior sampling algorithm. However, these results cannot directly imply the no-regret guarantee in our online learning setting, due to the known fact that posterior sampling can perform poorly under potentially *adversarial* or *non-stationary* environments (Zimmert & Seldin, 2021; Liu et al., 2023d). In contrast, we here establish the equivalence of the pre-trained LLM to the FTPL algorithm (under different pre-training data distribution specifications), with the ability to handle arbitrary loss sequences, even though the LLMs are only trained on a fixed/stationary distribution of texts (tasks).

D.9 Details of Estimating the Parameters of Our Hypothetical Model

To further validate our model and data distribution assumptions, we also propose to estimate the parameter $\{\eta_t\}_{t \in [T-1]}$ in Definition 4.1, using data from interacting with LLMs (following the same protocol as before), with P_{noise} being a standard normal distribution (note that we do not need to estimate η_0 by Definition 4.1). Specifically, given n episodes of the LLM agent’s behavior $\{(\ell_t^{(j)}, \pi_t^{(j)})_{t \in [T]}\}_{j \in [n]}$, motivated by our Lemma 1 and Theorem 4.1, we estimate $\{\eta_t\}_{t \in [T-1]}$ by solving the following problem

$$\sigma^* \in \arg \min_{\sigma > 0} \sum_{t \in [T-1]} \sum_{j \in [n]} \left\| \pi_{t+1}^{(j)} - P_{quantal}^{\sigma \sqrt{t+1}} \left(\cdot \mid \ell_{1:t}^{(j)} \right) \right\|_1, \quad \eta_t^* = \sigma^* \sqrt{t+1}, \quad \forall t \in [T-1].$$

We solve this single-variable optimization problem by grid search over $[0, 10]$. We then run the generalized quantal response model with the estimated $\{\eta_t^*\}_{t \in [T-1]}$ on another *unseen test set*, and compare it with the behavior of the actual LLM agents. We use all the interaction data from Section 3.2 and split it in half for training and testing.

We also use the same framework to understand the regrettable behaviors in Section 3.4. This analysis uses all the data from Section 3.4. We first find that such fitting procedures do not yield good predictions for LLMs on those counter-examples. Therefore, we resort to a more expressive model by directly fitting each η_t as

$$\eta_t^* \in \arg \min_{\eta_t > 0} \sum_{j \in [n]} \left\| \pi_{t+1}^{(j)} - P_{quantal}^{\eta_t} \left(\cdot \mid \ell_{1:t}^{(j)} \right) \right\|_1$$

separately for each $t \in [T-1]$. Even under the expressive model, LLMs fail to follow the generalized quantal response for the counter-examples with noisy alternating or adaptive loss sequences, as Figure 4.1 shows the gap between GPT-4 (dynamic) regret and the our model’s (dynamic) regret.

E Deferred Results and Proofs in Section 5

E.1 Basic Lemmas

Lemma 5 (Double iterated limit). *For a sequence $(a_{mn})_{m,n \in \mathbb{N}^+}$, suppose that $\lim_{m,n \rightarrow \infty} a_{mn} = L$. Then the*

following are equivalent:

- For each m , $\lim_{n \rightarrow \infty} a_{mn}$ exists;
- $\lim_{m \rightarrow \infty} \lim_{n \rightarrow \infty} a_{mn} = L$.

Lemma 6 (Hoeffding's inequality). Let X_1, X_2, \dots, X_n be independent random variables bounded by the intervals $[a_i, b_i]$, respectively. Define $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ and let $\mu = \mathbb{E}[\bar{X}]$ be the expected value of \bar{X} . Then, for any $t > 0$,

$$\mathbb{P}(|\bar{X} - \mu| \geq t) \leq 2 \exp\left(-\frac{2n^2 t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right).$$

Lemma 7 (Uniform convergence \implies Interchanging limit and infimum). If $(f_n : X \rightarrow \mathbb{R})_{n \in \mathbb{N}^+}$ is a sequence of continuous functions that uniformly converge to a function $f : X \rightarrow \mathbb{R}$ on the domain X , then $\lim_{n \rightarrow \infty} \inf_{x \in X} f_n(x) = \inf_{x \in X} f(x)$ holds.

E.2 Deferred Proof for the Arguments in Section 5.1

In this section, we prove some properties of $\mathcal{L}(\theta, k, N)$ under certain regularity conditions of f, h . Throughout this subsection, we will assume the following condition holds.

Condition 1. For $h : \mathbb{R} \rightarrow \mathbb{R}^+$ and $f : \mathbb{R} \times \mathbb{N}^+ \rightarrow \mathbb{R}^+$, suppose $h(\cdot)$ and $f(\cdot, k)$ are both continuous and non-decreasing functions for any $k \in \mathbb{N}^+$. The derivative $h' : \mathbb{R} \rightarrow \mathbb{R}$ is also a continuous function. Moreover, f satisfies that $\log f(R_1, k_1) - \log f(R_1, k_2) \geq \log f(R_2, k_1) - \log f(R_2, k_2)$ for $R_1 \geq R_2$ and $k_1 \geq k_2$, i.e., $\log f$ is supermodular. Lastly, f is a function such that $\lim_{k \rightarrow \infty} \frac{f(R_1, k)}{f(R_2, k)} = \infty \cdot \mathbb{1}(R_1 > R_2) + \mathbb{1}(R_1 = R_2)$, with the convention of $\infty \cdot 0 = 0$. Lastly, $(\ell_t^{(j)})_{t \in [T], j \in [N]}$ are continuous random variables supported on $[-B, B]^{T \times N}$.

Claim 1 (Iterated limit of $\mathcal{L}(\theta, k, N)$ is the same as double limit of $\mathcal{L}(\theta, k, N)$). It holds that:

$$\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \mathcal{L}(\theta, k, N) = \lim_{N, k \rightarrow \infty} \mathcal{L}(\theta, k, N) = \lim_{k \rightarrow \infty} \lim_{N \rightarrow \infty} \mathcal{L}(\theta, k, N) = h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right).$$

Proof. **Step 1. Proving** $\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \mathcal{L}(\theta, k, N) = h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right)$.

Firstly, as both h and f are non-negative (Condition 1), and $\mathbb{E}_{(\ell_t^{(j)})_{t \in [T], j \in [N]}} \left[h(\max_{j \in [N]} \text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) \right]$ exists, we have by dominated convergence theorem that

$$\begin{aligned} \lim_{k \rightarrow \infty} \mathcal{L}(\theta, k, N) &= \mathbb{E} \lim_{k \rightarrow \infty} \left[\frac{\sum_{j \in [N]} h(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) f(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)}{\sum_{j \in [N]} f(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)} \right] \\ &= \mathbb{E}_{(\ell_t^{(j)})_{t \in [T], j \in [N]}} \left[h(\max_{j \in [N]} R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) \right] \end{aligned}$$

where R_{LLM_θ} denotes an abbreviation of $\text{Regret}_{\text{LLM}_\theta}$. By (Ahsanullah et al., 2013, Chapter 11), we have $h(\max_{j \in [N]} \text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) \xrightarrow{p} h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$ when $N \rightarrow \infty$. Hence, we have $\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \mathcal{L}(\theta, k, N) = h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$ holds.

Step 2. Proving $\lim_{N, k \rightarrow \infty} \mathcal{L}(\theta, k, N) = h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$.

Now, we will calculate $\lim_{N, k \rightarrow \infty} \mathcal{L}(\theta, k, N)$.

Lemma 8. For any $0 < \epsilon < 1$, it follows that

$$\lim_{N,k \rightarrow \infty} \frac{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} = 0$$

and

$$\lim_{N,k \rightarrow \infty} \frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} = 0$$

hold with probability 1, where X_i 's are i.i.d. random variables, $\text{esssup}(H(X_i)) = 1$, and $H : \mathbb{R} \rightarrow \mathbb{R}^+$ is a continuous non-decreasing function.

Proof of Lemma 8. Since $f(\cdot, k), H$ are non-negative and non-decreasing functions, we have

$$\frac{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} \leq \frac{(1 - \epsilon) f(H^{-1}(1 - \epsilon), k) |\{i \in [N] \mid (H(X_i) < 1 - \epsilon)\}|}{(1 - \epsilon/2) f(H^{-1}(1 - \epsilon/2), k) |\{i \in [N] \mid (H(X_i) > 1 - \epsilon/2)\}|}$$

and we know that

$$\frac{|\{i \in [N] \mid (H(X_i) < 1 - \epsilon)\}|}{|\{i \in [N] \mid (H(X_i) > 1 - \epsilon/2)\}|} \xrightarrow{a.s.} \frac{F(1 - \epsilon)}{1 - F(1 - \epsilon/2)}$$

as $N \rightarrow \infty$, where F is the cumulative distribution function of random variable $H(X)$. Therefore, we have

$$\begin{aligned} 0 &\leq \lim_{N,k \rightarrow \infty} \frac{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} \\ &\leq \lim_{N,k \rightarrow \infty} \frac{(1 - \epsilon) f(H^{-1}(1 - \epsilon), k) |\{i \in [N] \mid (H(X_i) < 1 - \epsilon)\}|}{(1 - \epsilon/2) f(H^{-1}(1 - \epsilon/2), k) |\{i \in [N] \mid (H(X_i) > 1 - \epsilon/2)\}|} \\ &\leq \lim_{a.s. N,k \rightarrow \infty} \frac{(1 - \epsilon) f(H^{-1}(1 - \epsilon), k)}{(1 - \epsilon/2) f(H^{-1}(1 - \epsilon/2), k)} \frac{F(1 - \epsilon)}{1 - F(1 - \epsilon/2)} = 0. \end{aligned}$$

By a similar argument, we have

$$\lim_{N,k \rightarrow \infty} \frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} = 0$$

with probability 1. □

One key idea in the proof above is the use of some *truncation* level ϵ for $H(X)$ with $\text{esssup}(H(X)) = 1$. By Lemma 8, we have

$$\lim_{N,k \rightarrow \infty} \frac{\sum_{i=1}^N f(X_i, k) H(X_i) \mathbb{1}(H(X_i) > 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) H(X_i)} = \lim_{N,k \rightarrow \infty} \frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k)} = 1,$$

since

$$0 \leq \frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k)} \leq \frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon/2)}$$

holds with probability 1. Therefore, for any $0 < \epsilon < 1$, we have

$$\begin{aligned}
\lim_{N, k \rightarrow \infty} \mathcal{L}(\theta, k, N) &= \mathbb{E} \lim_{N, k \rightarrow \infty} \left[\frac{\sum_{j \in [N]} h(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) f(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)}{\sum_{j \in [N]} f(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)} \right] \\
&= h \left(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right) \\
&\quad \times \mathbb{E} \lim_{N, k \rightarrow \infty} \left[\frac{\sum_{j \in [N]} \frac{h(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} f(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k) \mathbb{1}(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} > 1 - \epsilon)}{\sum_{j \in [N]} f(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k) \mathbb{1}(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} > 1 - \epsilon)} \right] \\
&\geq (1 - \epsilon) h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))
\end{aligned}$$

which implies $\lim_{N, k \rightarrow \infty} \mathcal{L}(\theta, k, N) = h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$ since

$$\mathcal{L}(\theta, k, N) \leq h \left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right)$$

by definition of \mathcal{L} , the fact that h is non-decreasing, and by setting $\epsilon \rightarrow 0$ to obtain

$$\mathcal{L}(\theta, k, N) \geq h \left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right).$$

Here, we used the fact that $(\ell_t)_{t \in [T]}$ has a continuous distribution, $\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})$ is a continuous function, and the non-decreasing property and continuity of h (Condition 1), which lead to:

$$\text{esssup} \left(h \left(\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right) \right) = \max_{\ell_1, \dots, \ell_T} h \left(\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right) = h \left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}) \right). \quad (\text{E.1})$$

Equation (E.1) will be used frequently in the overall proof in Appendix E.2.

Step 3. Proving $\lim_{k \rightarrow \infty} \lim_{N \rightarrow \infty} \mathcal{L}(\theta, k, N) = h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$.

Lastly, if $N \rightarrow \infty$, similarly by dominated convergence theorem we have

$$\begin{aligned}
\lim_{N \rightarrow \infty} \mathcal{L}(\theta, k, N) &= \mathbb{E} \lim_{N \rightarrow \infty} \left[\frac{\sum_{j \in [N]} h \left(R_{\text{LLM}_\theta} \left((\ell_t^{(j)})_{t \in [T]} \right) \right) f \left(R_{\text{LLM}_\theta} \left((\ell_t^{(j)})_{t \in [T]} \right), k \right)}{\sum_{j \in [N]} f \left(R_{\text{LLM}_\theta} \left((\ell_t^{(j)})_{t \in [T]} \right), k \right)} \right] \\
&= \frac{\mathbb{E} \left[h \left(R_{\text{LLM}_\theta} \left((\ell_t^{(j)})_{t \in [T]} \right) \right) f \left(R_{\text{LLM}_\theta} \left((\ell_t^{(j)})_{t \in [T]} \right), k \right) \right]}{\mathbb{E} \left[f \left(R_{\text{LLM}_\theta} \left((\ell_t^{(j)})_{t \in [T]} \right), k \right) \right]}.
\end{aligned}$$

Thus, $\lim_{N \rightarrow \infty} \mathcal{L}(\theta, k, N)$ always exists for every k . Now, we use the known property of double iterated limit (Lemma 5), and obtain that $\lim_{k \rightarrow \infty} \lim_{N \rightarrow \infty} \mathcal{L}(\theta, k, N) = h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$. \square

Claim 2 (Uniform convergence of $\mathcal{L}(\theta, k, N)$ (with respect to k and N)). $\mathcal{L}(\theta, k, N)$ uniformly converges to $h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$ on the domain Θ .

Proof. We will provide a similar analysis as Lemma 8 as follows:

Lemma 9. For any $0 < \epsilon < 1$, $0 < \delta < 1$, and $k \in \mathbb{N}^+$, we have

$$\frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon)} = \tilde{\mathcal{O}}\left(A(k, H, \epsilon) \left(\frac{1}{1 - F_{H,X}(1 - \epsilon/2)} + \frac{1}{\sqrt{N}}\right)\right)$$

with probability at least $1 - \delta$, where X_i 's are i.i.d. random variables, $\text{esssup}(H(X_i)) = 1$, $H : \mathbb{R} \rightarrow \mathbb{R}^+$ is a continuous non-decreasing function, $A(k, t, \epsilon) := \frac{(1-\epsilon)f((t/\text{esssup}(t(X)))^{-1}(1-\epsilon), k)}{(1-\epsilon/2)f((t/\text{esssup}(t(X)))^{-1}(1-\epsilon/2), k)}$, for any non-decreasing function $t : \mathbb{R} \rightarrow \mathbb{R}^+$, and $F_{t,X}$ is a cumulative distribution function of random variable $t(X)/\text{esssup}(t(X))$.

Proof of Lemma 9. With the same argument as the proof of Lemma 8, we have

$$\frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} \leq \frac{f(H^{-1}(1 - \epsilon), k) |\{i \in [N] \mid (H(X_i) < 1 - \epsilon)\}|}{f(H^{-1}(1 - \epsilon/2), k) |\{i \in [N] \mid (H(X_i) > 1 - \epsilon/2)\}|}.$$

It holds that $\frac{1}{N} |\{i \in [N] \mid (H(X_i) < 1 - \epsilon)\}| = F_{H,X}(1 - \epsilon) + \tilde{\mathcal{O}}(1/\sqrt{N})$ with probability at least $1 - \delta/2$ due to Hoeffding's inequality (Lemma 6). Similarly, we have $\frac{1}{N} |\{i \in [N] \mid (H(X_i) > 1 - \epsilon/2)\}| = 1 - F_{H,X}(1 - \epsilon/2) + \tilde{\mathcal{O}}(1/\sqrt{N})$ with probability at least $1 - \delta/2$. Therefore,

$$\frac{|\{i \in [N] \mid (H(X_i) < 1 - \epsilon)\}|}{|\{i \in [N] \mid (H(X_i) > 1 - \epsilon/2)\}|} = \frac{F_{H,X}(1 - \epsilon)}{1 - F_{H,X}(1 - \epsilon/2)} + \tilde{\mathcal{O}}(\sqrt{1/N}) \leq \frac{1}{1 - F_{H,X}(1 - \epsilon/2)} + \tilde{\mathcal{O}}(\sqrt{1/N}),$$

with probability at least $1 - \delta$. Finally, we have

$$\frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon)} < \frac{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) < 1 - \epsilon)}{\sum_{i=1}^N f(X_i, k) \mathbb{1}(H(X_i) > 1 - \epsilon/2)} \leq A(k, H, \epsilon) \left(\frac{1}{1 - F_{H,X}(1 - \epsilon/2)} + \tilde{\mathcal{O}}\left(\frac{1}{\sqrt{N}}\right)\right).$$

□

Note that $\lim_{k \rightarrow \infty} A(k, H, \epsilon) = 0$, since $\lim_{k \rightarrow \infty} \frac{f(R_1, k)}{f(R_2, k)} = \infty \cdot \mathbb{1}(R_1 > R_2) + \mathbb{1}(R_1 = R_2)$. By Lemma 9 with $H(R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})) = \frac{h(R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))}$, we have

$$\begin{aligned} & \frac{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k) \mathbb{1}\left(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} \geq 1 - \epsilon\right)}{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k)} \\ &= \frac{1}{\frac{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k) \mathbb{1}\left(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} < 1 - \epsilon\right)}{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k) \mathbb{1}\left(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} \geq 1 - \epsilon\right)}} \\ &\geq \frac{1}{1 + A(k, H, \epsilon) \left(\frac{1}{1 - F_{H, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})}(1 - \epsilon/2)} + \tilde{\mathcal{O}}(\sqrt{1/N})\right)}, \end{aligned}$$

where we recall the shorthand notation of $R_{\text{LLM}_\theta} = \text{Regret}_{\text{LLM}_\theta}$. Note that $A(k, H, \epsilon) = A(k, h, \epsilon)$ and

$F_{H, R_{\text{LLM}_\theta}} = F_{h, R_{\text{LLM}_\theta}}$ hold by the definitions of $F_{t, X}$ and $A(k, t, \epsilon)$ in Lemma 9. Therefore,

$$\begin{aligned}
1 &\geq \frac{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k) \frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))}}{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k)} \\
&\geq \frac{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k) \frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} \mathbb{1}(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} \geq 1 - \epsilon)}{\sum_{i=1}^N f(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}), k) \mathbb{1}(\frac{h(R_{\text{LLM}_\theta}((\ell_t^{(i)})_{t \in [T]}))}{h(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))} \geq 1 - \epsilon)} \\
&\quad \times \frac{1}{1 + A(k, h, \epsilon)(\frac{1}{1 - F_{h, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}})(1 - \epsilon/2)} + \widetilde{\mathcal{O}}(\sqrt{1/N}))} \\
&\geq \frac{1 - \epsilon}{1 + A(k, h, \epsilon)(\frac{1}{1 - F_{h, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}})(1 - \epsilon/2)} + \widetilde{\mathcal{O}}(\sqrt{1/N}))}
\end{aligned}$$

with probability at least $1 - \delta$.

Now, for any $\epsilon > 0$ and $\delta > 0$, we have

$$\begin{aligned}
0 &\leq h\left(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) - \mathcal{L}(\theta, k, N) \\
&\leq h\left(\max_{\ell_1, \dots, \ell_T} R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) \left(1 - \frac{(1 - \delta)(1 - \epsilon)}{1 + A(k, h, \epsilon)(\frac{1}{1 - F_{h, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}})(1 - \epsilon/2)} + \widetilde{\mathcal{O}}(\sqrt{1/N}))}\right).
\end{aligned}$$

Note that

$$1 - F_{h, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})}(1 - \epsilon/2) = \mathbb{P}\left(h(\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})) > (1 - \epsilon/2)h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right)\right)$$

is a continuous function of θ , since we assume LLM_θ is a continuous function of θ , $(\ell_t)_{t \in [T]}$ has a continuous distribution, and $\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})$ is a continuous function of LLM_θ and $(\ell_t)_{t \in [T]}$. Since we consider a compact Θ (as several recent works on analyzing Transformers (Bai et al., 2023; Lin et al., 2024)), we have $p(\epsilon) := \min_{\theta \in \Theta} 1 - F_{h, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})}(1 - \epsilon/2) > 0$. Therefore,

$$\left(1 - \frac{(1 - \delta)(1 - \epsilon)}{1 + A(k, h, \epsilon)(\frac{1}{1 - F_{h, R_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}})(1 - \epsilon/2)} + \widetilde{\mathcal{O}}(\sqrt{1/N}))}\right) \leq \left(1 - \frac{(1 - \delta)(1 - \epsilon)}{1 + A(k, h, \epsilon)(\frac{1}{p(\epsilon)} + \widetilde{\mathcal{O}}(\sqrt{1/N}))}\right), \quad (\text{E.2})$$

and we know that $\lim_{N, k \rightarrow \infty} 1 + A(k, h, \epsilon)(\frac{1}{p(\epsilon)} + \widetilde{\mathcal{O}}(\sqrt{1/N})) = 1$, which is not dependent on θ . Thus, we can conclude that $\lim_{N, k \rightarrow \infty} \sup_{\theta \in \Theta} |h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})) - \mathcal{L}(\theta, k, N)| = 0$, as we can choose arbitrarily small ϵ, δ . \square

Claim 3 (Double iterated limit of supremum). *It holds that:*

$$\lim_{N \rightarrow \infty} \limsup_{k \rightarrow \infty} \sup_{\theta \in \Theta} \left| \mathcal{L}(\theta, k, N) - h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) \right| = 0.$$

Proof. Since $h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})) \geq \mathcal{L}(\theta, k, N)$, we will prove

$$\lim_{N \rightarrow \infty} \limsup_{k \rightarrow \infty} \sup_{\theta \in \Theta} h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) - \mathcal{L}(\theta, k, N) = 0.$$

Lemma 10. $\frac{\sum_{i=1}^N f(X_i, k_1)h(X_i)}{\sum_{i=1}^N f(X_i, k_1)} \leq \frac{\sum_{i=1}^N f(X_i, k_2)h(X_i)}{\sum_{i=1}^N f(X_i, k_2)}$ holds if $0 < k_1 \leq k_2$ for any real-valued $(X_i)_{i \in [N]}$.

Proof. By multiplying $(\sum_{i=1}^N f(X_i, k_1))(\sum_{i=1}^N f(X_i, k_2))$ on both sides of the formula, we know that it is equivalent to $\sum_{1 \leq i \neq j \leq N} f(X_i, k_1)h(X_i)f(X_j, k_2) \leq \sum_{1 \leq i \neq j \leq N} f(X_i, k_1)h(X_j)f(X_j, k_2)$. This is equivalent to

$$\sum_{1 \leq i \neq j \leq N} (f(X_i, k_1)f(X_j, k_2) - f(X_j, k_1)f(X_i, k_2))(h(X_i) - h(X_j)) \leq 0,$$

which is true since if $X_i \geq X_j$, $(f(X_i, k_1)f(X_j, k_2) - f(X_j, k_1)f(X_i, k_2)) \leq 0$ due to the log-increasing difference of f (Condition 1), as $\log f(X_j, k_1) - \log f(X_j, k_2) \geq \log f(X_i, k_1) - \log f(X_i, k_2)$ if $X_i \geq X_j$. \square

Therefore, $\mathcal{L}(\theta, k, N)$ is a non-decreasing function of k if N is fixed, which indicates that

$$\lim_{k \rightarrow \infty} \sup_{\theta \in \Theta} h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) - \mathcal{L}(\theta, k, N)$$

exists, as $\mathcal{L}(\theta, k, N)$ is also bounded. Therefore, by Lemma 5 and Claim 2, we know that

$$\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \sup_{\theta \in \Theta} \left| \mathcal{L}(\theta, k, N) - h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) \right|$$

exists and this value should be 0. \square

Claim 4. *It holds that*

$$\lim_{N, k \rightarrow \infty} \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) = \lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) = \inf_{\theta \in \Theta} h\left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right).$$

Proof. Firstly, by Lemma 7, we have $\lim_{N, k \rightarrow \infty} \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) = \inf_{\theta \in \Theta} h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$. Plus, we already know that $\mathcal{L}(\theta, k, N)$ is a monotonically non-decreasing function of k for any fixed N (Lemma 10), and it is bounded, $\lim_{k \rightarrow \infty} \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N)$ always exists. Therefore, by Lemma 5, we also have $\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) = \inf_{\theta \in \Theta} h(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]}))$. \square

E.3 Definition of the Empirical Loss Function

Definition E.1 (Empirical loss function). *We define the empirical loss $\widehat{\mathcal{L}}$ computed with N_T samples as follows:*

$$\widehat{\mathcal{L}}(\theta, k, N, N_T) := \frac{1}{N_T} \sum_{s=1}^{N_T} \left[\frac{\sum_{j \in [N]} h\left(\text{Regret}_{\text{LLM}_\theta}((\ell_{s,t}^{(j)})_{t \in [T]})\right) f\left(\text{Regret}_{\text{LLM}_\theta}((\ell_{s,t}^{(j)})_{t \in [T]}), k\right)}{\sum_{j \in [N]} f\left(\text{Regret}_{\text{LLM}_\theta}((\ell_{s,t}^{(j)})_{t \in [T]}), k\right)} \right] \quad (\text{E.3})$$

where $(\ell_{s,t}^{(j)})_{j \in [N], t \in [T]}$ denotes the s -th sample of $(\ell_t^{(j)})_{j \in [N], t \in [T]}$ for estimating $\mathcal{L}(\theta, k, N)$.

E.4 Deferred Proofs of Theorem E.1 and Theorem 5.1

Theorem E.1. (Generalization gap). *Suppose LLM_θ is Lipschitz-continuous with respect to the model parameter θ , then for any $0 < \epsilon < 1/2$, with probability at least $1 - \epsilon$, we have*

$$\mathcal{L}(\widehat{\theta}_{k, N, N_T}, k, N) - \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) \leq \widetilde{\mathcal{O}}\left(\sqrt{\frac{d_\theta + \log(1/\epsilon)}{N_T}}\right), \quad (\text{E.4})$$

for any N and sufficiently large k , where d_θ is the dimension of the parameter θ .

Through a careful use of Berge’s Maximum Theorem (Berge, 1877), we prove that the right-hand side of Equation (E.4) does *not* depend on k and N , which allows us to take the limit of $\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty}$ without affecting the generalization bound.

Before proving the theorem, we remark on what LLM structure enjoys the Lipschitz-continuity. We provide two auxiliary results in the following proposition. The first result is from (Bai et al., 2023, Section J.1), which is about the Lipschitzness of Transformers. The second result is regarding processing the output of Transformers. In particular, the output of Transformers is usually not directly used, but passed through some matrix multiplication (by some matrix A), followed by some projection Operator (to be specified later).

Proposition 2. *The L -layer Transformer TF_θ as defined in Section 2.2 is C_{TF} -Lipschitz continuous with respect to θ with $C_{TF} := L \left((1 + B_{TF}^2)(1 + B_{TF}^2 R^3) \right)^L B_{TF} R (1 + B_{TF} R^2 + B_{TF}^3 R^2)$, i.e.,*

$$\|TF_{\theta_1}(Z) - TF_{\theta_2}(Z)\|_{2,\infty} \leq C_{TF} \|\theta_1 - \theta_2\|_{TF}$$

where $\|\cdot\|_{TF}$ is as defined in Equation (2.1), and R, Z, B_{TF} are as introduced in Section 2.2. Moreover, the function $Operator(A \cdot TF_\theta(\cdot)_{-1})$ is $\|A\|_{op} C_{TF}$ -Lipschitz continuous with respect to θ , i.e.,

$$\|Operator(A \cdot TF_{\theta_1}(Z)_{-1}) - Operator(A \cdot TF_{\theta_2}(Z)_{-1})\|_2 \leq \|A\|_{op} C_{TF} \|\theta_1 - \theta_2\|_{TF}.$$

Here, $Operator$ is either the projection operator onto some convex set, or the Softmax function.

Proof. The first result is from (Bai et al., 2023, Section J.1). The second result comes from

- If $Operator$ is a projection onto the convex set, then $\|Operator(x) - Operator(y)\|_2 \leq \|x - y\|_2$;
- If $Operator$ is Softmax, then $\|Softmax(x) - Softmax(y)\|_2 \leq \|x - y\|_2$ (Gao & Pavel, 2017, Corollary 3).

Note that the only condition that we require for $Operator$ is its non-expansiveness. \square

Proof of Theorem E.1. Let C_{LLM} denote the Lipschitz-continuity constant for LLM_θ with respect to some norm $\|\cdot\|_{LLM}$, where $\|\cdot\|_{LLM}$ denotes any norm defined on the parameter space of LLM (e.g., the norm $\|\cdot\|_{TF}$ above in Proposition 2). Now, we prove that regret is also a Lipschitz-continuous function with respect to the LLM’s parameter.

Lemma 11 (Lipschitzness of regret). *The function $Regret_{LLM_\theta}$ is $C_{Reg} := BC_{LLM}T$ -Lipschitz continuous with respect to θ , i.e.,*

$$\left| Regret_{LLM_{\theta_1}}((\ell_t)_{t \in [T]}) - Regret_{LLM_{\theta_2}}((\ell_t)_{t \in [T]}) \right| \leq C_{Reg} \|\theta_1 - \theta_2\|_{LLM}.$$

Proof. By definition, we have

$$\begin{aligned} \left| Regret_{LLM_{\theta_1}}((\ell_t)_{t \in [T]}) - Regret_{LLM_{\theta_2}}((\ell_t)_{t \in [T]}) \right| &= \left| \sum_{t=1}^T \langle \ell_t, LLM_{\theta_1}(Z_{t-1}) - LLM_{\theta_2}(Z_{t-1}) \rangle \right| \\ &= B \sum_{t=1}^T \|LLM_{\theta_1}(Z_{t-1}) - LLM_{\theta_2}(Z_{t-1})\| \\ &\leq BC_{LLM}T \|\theta_1 - \theta_2\|_{LLM} \end{aligned}$$

where $Z_t := (\ell_1, \dots, \ell_t, c)$ for all $t \in [T]$ and $Z_0 = (c)$ where c is a d -dimensional vector. \square

Now, we will prove the Lipschitzness of

$$C\left((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta\right) := \frac{\sum_{j \in [N]} h(\text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) f(\text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)}{\sum_{j \in [N]} f(\text{Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)} \quad (\text{E.5})$$

with respect to the model parameter θ .

Claim 5. For any $R > 0$, there exists $\beta_R > 0$ such that if $\beta > \beta_R$, we have

$$\left| \frac{\sum_{n \in [N]} x_n f(x_n, \beta)}{\sum_{n \in [N]} f(x_n, \beta)} - \frac{\sum_{n \in [N]} y_n f(y_n, \beta)}{\sum_{n \in [N]} f(y_n, \beta)} \right| \leq 2\|x - y\|_\infty$$

for every $x, y \in \mathbb{R}^n$ such that $|x_i| \leq R, |y_i| \leq R$ for all $i \in [N]$.

Proof. If $\beta = \infty$, we have

$$\lim_{\beta \rightarrow \infty} \left(\left| \frac{\sum_{n \in [N]} x_n f(x_n, \beta)}{\sum_{n \in [N]} f(x_n, \beta)} - \frac{\sum_{n \in [N]} y_n f(y_n, \beta)}{\sum_{n \in [N]} f(y_n, \beta)} \right| / \|x - y\|_\infty \right) = \frac{|\max_{n \in [N]} x_n - \max_{n \in [N]} y_n|}{\|x - y\|_\infty} \leq 1$$

holds. Moreover, consider the following constrained optimization problem:

$$\begin{aligned} \max_{x, y \in \mathbb{R}^n} \quad & \left(\left| \frac{\sum_{n \in [N]} x_n f(x_n, \beta)}{\sum_{n \in [N]} f(x_n, \beta)} - \frac{\sum_{n \in [N]} y_n f(y_n, \beta)}{\sum_{n \in [N]} f(y_n, \beta)} \right| / \|x - y\|_\infty \right) \\ \text{subject to} \quad & |x_i| \leq R, \quad |y_i| \leq R \quad \text{for all } i \in [N], \end{aligned}$$

whose optimum is denoted as $F(R, \beta)$. Then, since $\|x\|_\infty \leq R$ and $\|y\|_\infty \leq R$ is a compact set, by Berge's Maximum Theorem (Berge, 1977), we have that $F(R, \beta)$ is a continuous function for β . Moreover, we know that $F(R, \infty) \leq 1$, which indicates that we can find a large enough β_R such that if $\beta > \beta_R$, $F(R, \beta) \leq 2$. \square

Note that Claim 5 does not hold if either x_i or y_i is unbounded. Now, we will apply Claim 5 to Equation (E.5). We can guarantee that $|\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})| \leq \text{diam}(\Pi, \|\cdot\|_2)TB$.

Also, note that the domain of $h : \mathbb{R} \rightarrow \mathbb{R}^+$ is effectively constrained to the range that $\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})$ can achieve, which means that we can regard h as $h : [-\text{diam}(\Pi, \|\cdot\|_2)TB, \text{diam}(\Pi, \|\cdot\|_2)TB] \rightarrow \mathbb{R}^+$. Due to the continuity of h' , and the fact that h has a compact domain, we know that $h(\cdot)$ is C_h -Lipschitz continuous for some $C_h > 0$ on this interval of $[-\text{diam}(\Pi, \|\cdot\|_2)TB, \text{diam}(\Pi, \|\cdot\|_2)TB]$.

Lemma 12 (Lipschitzness of C in Equation (E.5)). The function C in Equation (E.5) is $C_{\text{cost}} := 2C_h C_{\text{Reg}}$ -Lipschitz continuous with respect to θ , if $k > k_{\text{diam}(\Pi, \|\cdot\|_2)TB}$ for some $k_{\text{diam}(\Pi, \|\cdot\|_2)TB} > 0$, i.e.,

$$\left| C\left((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta_1\right) - C\left((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta_2\right) \right| \leq C_{\text{cost}} \|\theta_1 - \theta_2\|_{\text{LLM}}.$$

Proof.

$$\begin{aligned} & \left| C((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta_1) - C((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta_2) \right| \\ & \stackrel{(i)}{\leq} 2 \|h(\text{Regret}_{\text{LLM}_{\theta_1}}((\ell_t^{(j)})_{t \in [T]})) - h(\text{Regret}_{\text{LLM}_{\theta_2}}((\ell_t^{(j)})_{t \in [T]}))\|_\infty \\ & \stackrel{(ii)}{\leq} 2C_h \|\text{Regret}_{\text{LLM}_{\theta_1}}((\ell_t^{(j)})_{t \in [T]}) - \text{Regret}_{\text{LLM}_{\theta_2}}((\ell_t^{(j)})_{t \in [T]})\|_\infty \\ & \stackrel{(iii)}{\leq} 2C_h C_{\text{Reg}} \|\theta_1 - \theta_2\|_{\text{LLM}} = C_{\text{cost}} \|\theta_1 - \theta_2\|_{\text{LLM}}. \end{aligned}$$

Here, (i) holds due to Claim 5, (ii) holds since h is C_h -Lipschitz continuous on the range of $\text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})$, and (iii) holds due to Lemma 11. \square

For completeness of the paper, we provide the definition of covering set and covering number.

Definition E.2 (Covering set and covering number). *For $\delta > 0$, a metric space $(X, \|\cdot\|)$, and subset $Y \subseteq X$, set $C \subset Y$ is a δ -covering of Y when $Y \subseteq \cup_{c \in C} B(c, \delta, \|\cdot\|)$ holds. δ -covering number $N(\delta; Y, \|\cdot\|)$ is defined as the minimum cardinality of any covering set.*

By (Wainwright, 2019, Example 5.8), for any $r > 0$, we can verify that the δ -covering number $N(\delta; B(0, r, \|\cdot\|_{\text{LLM}}), \|\cdot\|_{\text{LLM}})$ can be bounded by

$$\log N(\delta; B(0, r, \|\cdot\|_{\text{LLM}}), \|\cdot\|_{\text{LLM}}) \leq d_\theta \log(1 + 2r/\delta),$$

where d_θ is the dimension of the LLM's whole parameter. For example, if we use the $\|\cdot\|_{\text{TF}}$ and consider the Transformer model as defined in Section 2.2, for any $r > 0$,

$$\log N(\delta; B(0, r, \|\cdot\|_{\text{LLM}}), \|\cdot\|_{\text{LLM}}) \leq L(3Md^2 + 2d(dd' + 3md^2)) \log(1 + 2r/\delta).$$

Since we consider a compact Θ (as several recent works on analyzing Transformers (Bai et al., 2023; Lin et al., 2024)), let $R_\Theta := \text{diam}(\Theta, \|\cdot\|_{\text{LLM}})$ (which corresponds to B_{TF} for the Transformer models as defined in Section 2.2, with $\|\cdot\|_{\text{LLM}} = \|\cdot\|_{\text{TF}}$), then there exists a set Θ_0 with $\log|\Theta_0| = d_\theta \log(1 + 2R_\Theta/\delta)$ such that for any $\theta \in \Theta$, there exists a $\theta_0 \in \Theta_0$ with

$$\left| C\left((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta\right) - C\left((\ell_t^{(j)})_{t \in [T], j \in [N]}, k, \theta_0\right) \right| \leq C_{\text{cost}} \delta.$$

Then, by the standard result from statistical learning theory (Wainwright, 2019, Chapter 5), when trained with N_T samples, for every $0 < \epsilon < 1/2$, with probability at least $1 - \epsilon$, we have

$$\mathcal{L}(\widehat{\theta}_{k, N, N_T}, k, N) - \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) \leq \sqrt{\frac{2(\log|\Theta_0| + \log(2/\epsilon))}{N_T}} + 2C_{\text{cost}} \delta.$$

Setting $\delta = \Omega(\sqrt{\log(\epsilon)/N_T})$, we further obtain

$$\mathcal{L}(\widehat{\theta}_{k, N, N_T}, k, N) - \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) \leq \widetilde{\mathcal{O}}\left(\sqrt{\frac{\log|\Theta_0| + \log(1/\epsilon)}{N_T}}\right)$$

with probability at least $1 - \epsilon$, completing the proof. \square

Theorem 5.1. (Regret). *Suppose³ for any $k \in \mathbb{N}^+$, $h, f(\cdot, k)$ are non-decreasing, and $\log f$ is a supermodular function (i.e., $\log f(R_1, k_1) - \log f(R_1, k_2) \geq \log f(R_2, k_1) - \log f(R_2, k_2)$ for $R_1 \geq R_2$ and $k_1 \geq k_2$). Then, with high probability, we have*

$$h\left(\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_{\widehat{\theta}_{k, N, N_T}}}((\ell_t)_{t \in [T]})\right) \leq h\left(\inf_{\theta \in \Theta} \max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_\theta}((\ell_t)_{t \in [T]})\right) + \widetilde{\mathcal{O}}\left(\sqrt{\frac{d_\theta}{N_T}}\right).$$

Proof. The limit on the right-hand side of Equation (E.4) remains as $\widetilde{\mathcal{O}}\left(\sqrt{\frac{d_\theta + \log(1/\epsilon)}{N_T}}\right)$, since we firstly take $\lim_{k \rightarrow \infty}$ and then take $\lim_{N \rightarrow \infty}$, thanks to the fact that Theorem E.1 holds for large enough k

³Note that these conditions on h, f are in addition to those specified after Equation (5.2).

and any N . Next, we have

$$\begin{aligned}
& \lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \left| \mathcal{L}(\widehat{\theta}_{k,N,N_T}, k, N) - h \left(\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_{\widehat{\theta}_{k,N,N_T}}} ((\ell_t)_{t \in [T]}) \right) \right| \\
& \leq \lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \left| \mathcal{L}(\widehat{\theta}_{k,N,N_T}, k, N) - h \left(\max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_{\widehat{\theta}_{k,N,N_T}}} ((\ell_t)_{t \in [T]}) \right) \right| + \\
& \quad \lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \left| h \left(\max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_{\widehat{\theta}_{k,N,N_T}}} ((\ell_t)_{t \in [T]}) \right) - h \left(\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_{\widehat{\theta}_{k,N,N_T}}} ((\ell_t)_{t \in [T]}) \right) \right| \\
& \leq \lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \sup_{\theta \in \Theta} \left| \mathcal{L}(\theta, k, N) - h \left(\max_{\|\ell_t\|_\infty \leq B} \text{Regret}_{\text{LLM}_\theta} ((\ell_t)_{t \in [T]}) \right) \right| + 0 = 0,
\end{aligned}$$

due to the continuity of h and Claim 3. Finally, we have

$$\lim_{N \rightarrow \infty} \lim_{k \rightarrow \infty} \inf_{\theta \in \Theta} \mathcal{L}(\theta, k, N) = \inf_{\theta \in \Theta} h \left(\max_{\ell_1, \dots, \ell_T} \text{Regret}_{\text{LLM}_\theta} ((\ell_t)_{t \in [T]}) \right)$$

due to Claim 4, which, combined with the fact that h is non-decreasing, completes the proof. \square

As a result, the coarse correlated equilibrium will emerge as the long-term interactions of multiple such learned LLMs, as stated in the following corollary.

Corollary 1. (Emerging behavior: Coarse correlated equilibrium). *For a sufficiently large N_T , if each agent in the matrix game plays according to $\text{LLM}_{\widehat{\theta}_{k,N,N_T}}$, then the time-averaged policy for each agent will constitute an approximate coarse correlated equilibrium of the game.*

Remark E.1 (Dynamic-regret loss). *So far, we have focused on the canonical online learning setting with regret being the metric. One can also generalize the results to the non-stationary setting, with dynamic regret being the metric. Specifically, one can define the dynamic-regret-loss function as follows:*

$$\mathcal{L}(\theta, k, N) := \mathbb{E} \left[\frac{\sum_{j \in [N]} h(\text{D-Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]})) f(\text{D-Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)}{\sum_{j \in [N]} f(\text{D-Regret}_{\text{LLM}_\theta}((\ell_t^{(j)})_{t \in [T]}), k)} \right].$$

Then, one can also establish similar results as before, since the analysis does not utilize other properties of the regret except its boundedness, and the Lipschitz-continuity of LLM with respect to θ . To be specific, Lemma 11 holds due to the reason that we can bound the difference of the regret with the term

$$\left| \sum_{t=1}^T \langle \ell_t, (\text{LLM}_{\theta_1}(Z_{t-1}) - \text{LLM}_{\theta_2}(Z_{t-1})) \rangle \right|,$$

as well as the fact that $\inf_{\pi_i \in \Pi} \langle \ell_i, \pi_i \rangle$ will be canceled. One can verify that all the arguments in Appendix E.2 also hold for similar reasons.

E.5 Detailed Explanation of Optimizing Equation (5.2) with Single-layer Self-attention Model

We consider the following structure of single-layer self-attention model g (see a formal introduction in Section 2.2):

$$g(Z_t; V, K, Q, v_c, k_c, q_c) := (V \ell_{1:t} + v_c \mathbf{1}_t^\top) \text{Softmax} \left((K \ell_{1:t} + k_c \mathbf{1}_t^\top)^\top \cdot (Qc + q_c) \right), \quad (\text{E.6})$$

where $Z_t = (\ell_1, \dots, \ell_t, c)$ and $V, K, Q \in \mathbb{R}^{d \times d}$ correspond to the value, key, and query matrices, respectively, $v_c, k_c, q_c \in \mathbb{R}^d$ correspond to the bias terms associated with V, K, Q , and $c \neq \mathbf{0}_d$ is a constant vector. We then have the following result.

Theorem E.2. Consider the policy space $\Pi = B(0, R_\Pi, \|\cdot\|)$ for some $R_\Pi > 0$. The configuration of a single-layer self-attention model in Equation (E.6) (V, K, Q, v_c, k_c, q_c) such that $K^\top(Qc + q_c) = v_c = \mathbf{0}_d$ and $V = -R_\Pi \frac{T}{\sum_{t=1}^{T-1} 1/t} \Sigma^{-1} \mathbb{E} \left[\left\| \sum_{t=1}^T \ell_t \right\| \ell_1 \ell_2^\top \right] \Sigma^{-1}$ is a first-order stationary point of Equation (5.2) with $N = 1$, $h(x) = x^2$. Moreover, if Σ is a diagonal matrix, then plugging this configuration into Equation (E.6), and projecting the output with $\text{Proj}_{\Pi, \|\cdot\|}$ would perform FTRL with an L_2 -regularizer for the loss vectors $(\ell_t)_{t \in [T]}$.

In practical training, such stationary points of the loss may be attained by first-order optimization algorithms of (stochastic) gradient descent, the workhorse in machine learning.

E.6 Deferred Proof of Theorem E.2

Theorem E.2. Consider the policy space $\Pi = B(0, R_\Pi, \|\cdot\|)$ for some $R_\Pi > 0$. The configuration of a single-layer self-attention model in Equation (E.6) (V, K, Q, v_c, k_c, q_c) such that $K^\top(Qc + q_c) = v_c = \mathbf{0}_d$ and $V = -R_\Pi \frac{T}{\sum_{t=1}^{T-1} 1/t} \Sigma^{-1} \mathbb{E} \left[\left\| \sum_{t=1}^T \ell_t \right\| \ell_1 \ell_2^\top \right] \Sigma^{-1}$ is a first-order stationary point of Equation (5.2) with $N = 1$, $h(x) = x^2$. Moreover, if Σ is a diagonal matrix, then plugging this configuration into Equation (E.6), and projecting the output with $\text{Proj}_{\Pi, \|\cdot\|}$ would perform FTRL with an L_2 -regularizer for the loss vectors $(\ell_t)_{t \in [T]}$.

Proof. We will locally use $\mathcal{A} = [d]$ without losing generality as \mathcal{A} is finite with $|\mathcal{A}| = d$, and will interchangeably use $\ell_i(j)$ and ℓ_{ij} for notational convenience. Define $a := K^\top(Qc + q_c) \in \mathbb{R}^d$ and $b_{t-1} := \beta \mathbf{1}_{t-1} := k_c^\top(Qc + q_c) \mathbf{1}_{t-1} \in \mathbb{R}^{t-1}$. With $N = 1$, $h(x) = x^2$, and the choice of Π , the loss function (Equation (5.2)) can be written as follows:

$$f(V, a, (b_t)_{t \in [T-1]}, v_c) := \mathbb{E} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2,$$

where for $t = 1$, we use the output of the single-layer self-attention as v_c and we will write it as $(V \ell_{1:0} + v_c \mathbf{1}_0^\top) \text{Softmax}(\ell_{1:0}^\top a + b_0)$ for notational consistency with $t \geq 2$. Also, we will define empty sum $\sum_{i=1}^0 a_i = 0$ for any sequence $(a_i)_{i \in \mathbb{N}^+}$.

Step 1. Calculating $\frac{\partial f}{\partial a}$.

For $x \in [d]$, we calculate the corresponding directional derivative with the following equation for $t \geq 2$:

$$\begin{aligned} & \frac{\partial}{\partial a_x} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) \\ &= \frac{\partial}{\partial a_x} \sum_{i=1}^{t-1} \ell_i^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) e_i \frac{\exp(e_i^\top (\ell_{1:t-1}^\top a + b_{t-1}))}{\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1}))} \\ &= \frac{\sum_{i=1}^{t-1} \ell_i^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) e_i \exp(e_i^\top (\ell_{1:t-1}^\top a + b_{t-1})) \frac{\partial e_i^\top (\ell_{1:t-1}^\top a + b_{t-1})}{\partial a_x} (\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1})))}{(\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1})))^2} \\ & \quad - \frac{\sum_{i=1}^{t-1} \ell_i^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) e_i \exp(e_i^\top (\ell_{1:t-1}^\top a + b_{t-1})) \left(\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1})) \frac{\partial e_s^\top (\ell_{1:t-1}^\top a + b_{t-1})}{\partial a_x} \right)}{(\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1})))^2}. \end{aligned}$$

Plugging $a = \mathbf{0}_d$ and $v_c = \mathbf{0}_d$, and $(b_t = \beta \mathbf{1}_t)_{t \in [T-1]}$ provides

$$\begin{aligned} & \left. \frac{\partial}{\partial a_x} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) \right|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\ &= \sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \ell_{ix}}{(t-1)} - \sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \left(\sum_{s=1}^{t-1} \ell_{sx} \right)}{(t-1)^2}. \end{aligned}$$

For $t = 1$, as $\ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) = \ell_1^\top v_c$, $\left. \frac{\partial}{\partial a_x} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) \right|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} = 0$, so we can use the same formula as $t \geq 2$ with empty sum $\sum_{i=1}^{t-1}$. Using

the above calculation, we can further compute $\left. \frac{\partial f}{\partial a_x} \right|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}}$ as follows:

$$\begin{aligned} & \left. \frac{\partial f(V, a, (b_t)_{t \in [T-1]}, v_c)}{\partial a_x} \right|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\ &= \mathbb{E} \left[\frac{\partial}{\partial a_x} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \right]_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\ &= \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \right]_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\ & \quad \left. \frac{\partial}{\partial a_x} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \right]_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\ &= \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top V \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_i + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \sum_{t=1}^T \left(\sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \ell_{ix}}{(t-1)} - \sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \left(\sum_{s=1}^{t-1} \ell_{sx} \right)}{(t-1)^2} \right) \right] \quad (\text{E.7}) \\ &= 0, \end{aligned}$$

where we used the fact that ℓ_i is drawn from a symmetric distribution, and flipping the sign of the variable as $-\ell_i$ yields the same distribution, which leads to the following:

$$\begin{aligned} & \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top V \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_i + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \sum_{t=1}^T \left(\sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \ell_{ix}}{(t-1)} - \sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \left(\sum_{s=1}^{t-1} \ell_{sx} \right)}{(t-1)^2} \right) \right] \\ &= \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top V \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_i + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \sum_{t=1}^T \left(- \sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \ell_{ix}}{(t-1)} + \sum_{i=1}^{t-1} \frac{\ell_t^\top V \ell_i \left(\sum_{s=1}^{t-1} \ell_{sx} \right)}{(t-1)^2} \right) \right]. \end{aligned}$$

This yields Equation (E.7)=0.

Step 2. Calculating $\frac{\partial f}{\partial v_c}$.

We will use the following equation for $t \geq 2$:

$$\begin{aligned} & \frac{\partial}{\partial v_c} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) \\ &= \frac{\partial}{\partial v_c} \sum_{i=1}^{t-1} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) e_i \frac{\exp(e_i^\top (\ell_{1:t-1}^\top a + b_{t-1}))}{\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1}))} = \ell_t. \end{aligned}$$

For $t = 1$, we define $\frac{\partial}{\partial v_c} \ell_1^\top (V \ell_{1:0} + v_c \mathbf{1}_0^\top) \text{Softmax}(\ell_{1:0}^\top a + b_0) = \ell_1$, so that we can use the same formula as $t \geq 2$. Therefore, we can calculate $\frac{\partial f}{\partial v_c} \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}}$ as follows:

$$\begin{aligned}
& \frac{\partial f(V, a, (b_t)_{t \in [T-1]}, v_c)}{\partial v_c} \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \mathbb{E} \frac{\partial}{\partial v_c} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \right. \\
&\quad \left. \frac{\partial}{\partial v_c} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \right] \\
&= \mathbb{E} \left[\left(\sum_{t=2}^T \ell_t^\top V \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_i + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \sum_{t=1}^T \ell_t \right] = 0.
\end{aligned}$$

The last line is due to the same reason as the last part of Step 1.

Step 3. Calculating $\frac{\partial f}{\partial V}$.

We calculate the following equation, which will be used to calculate $\frac{\partial f}{\partial V} \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}}$ for $t \geq 2$:

$$\begin{aligned}
& \frac{\partial}{\partial V} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \frac{\partial}{\partial V} \sum_{i=1}^{t-1} \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) e_i \frac{\exp(e_i^\top (\ell_{1:t-1}^\top a + b_{t-1}))}{\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1}))} \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \sum_{i=1}^{t-1} \ell_t \ell_i^\top \frac{\exp(e_i^\top (\ell_{1:t-1}^\top a + b_{t-1}))}{\sum_{s=1}^{t-1} \exp(e_s^\top (\ell_{1:t-1}^\top a + b_{t-1}))} \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} = \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_t \ell_i^\top.
\end{aligned}$$

For $t = 1$, note that $\frac{\partial}{\partial V} \ell_1^\top v_c = \mathbf{0}_{d \times d}$, so we can use the same formula as $t \geq 2$ with empty sum $\sum_{i=1}^{t-1}$.

Therefore, we have

$$\begin{aligned}
& \left. \frac{\partial f(V, a, (b_t)_{t \in [T-1]}, v_c)}{\partial V} \right|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \mathbb{E} \frac{\partial}{\partial V} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2 \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \right]_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&\quad \frac{\partial}{\partial V} \left(\sum_{t=1}^T \ell_t^\top (V \ell_{1:t-1} + v_c \mathbf{1}_{t-1}^\top) \text{Softmax}(\ell_{1:t-1}^\top a + b_{t-1}) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \ell_t^\top V \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_i + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right) \sum_{t=1}^T \sum_{i=1}^{t-1} \frac{1}{t-1} \ell_t \ell_i^\top \right] \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{i=1}^{t-1} \left(\frac{1}{t-1} \ell_t^\top V \ell_i \right) \left(\frac{1}{t-1} \ell_t \ell_i^\top \right) + R_\Pi T \left\| \sum_{t'=1}^T \ell_{t'} \right\|_2 \ell_t \ell_i^\top \right) \right] \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{i=1}^{t-1} \sum_{x=1}^d \sum_{y=1}^d v_{xy} \ell_{tx} \ell_{iy} \left(\frac{1}{t-1} \right)^2 [\ell_{tz} \ell_{iw}]_{(z,w)} + R_\Pi T \left\| \sum_{t'=1}^T \ell_{t'} \right\|_2 \ell_t \ell_i^\top \right) \right] \\
&= \sum_{t=1}^T \sum_{i=1}^{t-1} \sum_{x=1}^d \sum_{y=1}^d \frac{1}{(t-1)^2} [\sigma_{xz} v_{xy} \sigma_{yw}]_{(z,w)} + \mathbb{E} \left[R_\Pi T \left\| \sum_{t'=1}^T \ell_{t'} \right\|_2 \ell_t \ell_i^\top \right] \\
&= \left(\sum_{t=1}^{T-1} \frac{1}{t} \right) \Sigma V \Sigma + \mathbb{E} \left[R_\Pi T \left\| \sum_{t'=1}^T \ell_{t'} \right\|_2 \ell_t \ell_i^\top \right].
\end{aligned}$$

Therefore, if $V^\star = R_\Pi \frac{T}{\sum_{t=1}^{T-1} 1/t} \Sigma^{-1} \mathbb{E} \left[\left\| \sum_{t=1}^T \ell_t \right\|_2 \ell_t \ell_i^\top \right] \Sigma^{-1}$, then $\left. \frac{\partial f}{\partial V} \right|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}, V=V^\star} = \mathbf{0}_{d \times d}$.

Lastly, we have

$$\begin{aligned}
\left. \frac{\partial f}{\partial K} \right|_{K^\top(Qc+q_c)=v_c=\mathbf{0}_d, V=V^\star} &= \left(\frac{\partial f}{\partial a} \frac{\partial a}{\partial K} \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}, V=V^\star} = \mathbf{0}_{d \times d} \\
\left. \frac{\partial f}{\partial Q} \right|_{K^\top(Qc+q_c)=v_c=\mathbf{0}_d, V=V^\star} &= \left(\frac{\partial f}{\partial a} \frac{\partial a}{\partial Q} \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}, V=V^\star} = \mathbf{0}_{d \times d} \\
\left. \frac{\partial f}{\partial q_c} \right|_{K^\top(Qc+q_c)=v_c=\mathbf{0}_d, V=V^\star} &= \left(\frac{\partial f}{\partial a} \frac{\partial a}{\partial q_c} \right) \Big|_{a=\mathbf{0}_d, v_c=\mathbf{0}_d, (b_t=\beta \mathbf{1}_t)_{t \in [T-1]}, V=V^\star} = \mathbf{0}_d
\end{aligned}$$

which means that such configurations are first-order stationary points of Equation (5.2) with $N = 1$, $h(x) = x^2$, and $\Pi = B(0, R_\Pi, \|\cdot\|)$. \square

E.7 Deferred Proof of Theorem 5.2

Theorem 5.2. Consider the policy space $\Pi = B(0, R_\Pi, \|\cdot\|)$ for some $R_\Pi > 0$. The configuration of a single-layer linear self-attention model in Equation (5.3) (V, K, Q, v_c, k_c, q_c) such that $K^\top(Qc + q_c) = v_c = \mathbf{0}_d$ and $V = -2R_\Pi \Sigma^{-1} \mathbb{E} \left(\left\| \sum_{t=1}^T \ell_t \right\|_1 \ell_2^\top \right) \Sigma^{-1}$ is a global optimal solution of Equation (5.2) with $N = 1$, $h(x) =$

x^2 . Moreover, every global optimal configuration of Equation (5.2) within the parameterization class of Equation (5.3) has the same output function g . Additionally, if Σ is a diagonal matrix, then plugging any global optimal configuration into Equation (5.3), and projecting the output with $\text{Proj}_{\Pi, \|\cdot\|}$ is equivalent to FTRL with an L_2 -regularizer.

This theorem involves the analysis of a *non-convex optimization* problem through stationary point analysis. We identified the set of stationary points. By constructing the optimization problem as shown in Equation (E.13), we significantly reduced the candidate set for optimal points using our novel argument on the expected value of a nonnegative definite matrix. The main challenge here was to address the global optimization problem in a non-convex setting, which required the exploitation of the particular Transformer architecture.

Proof. The output of the single-layer linear self-attention structure is as follows:

$$\begin{aligned} g(Z_t; V, K, Q, v_c, k_c, q_c) \\ = \sum_{i=1}^t \left(V \ell_i \ell_i^\top (K^\top (Qc + q_c)) + \left(V k_c^\top (Qc + q_c) + v_c (Qc + q_c)^\top K \right) \ell_i + v_c k_c^\top (Qc + q_c) \right), \end{aligned} \quad (\text{E.8})$$

which can be expressed with a larger class

$$g(Z_t, \mathbb{A}, \beta, \mathbb{C}, \delta) := \sum_{i=1}^t (\mathbb{A} \ell_i \ell_i^\top \beta + \mathbb{C} \ell_i + \delta), \quad (\text{E.9})$$

where $\mathbb{A} \in \mathbb{R}^{d \times d}$, $\beta, \mathbb{C}, \delta \in \mathbb{R}^d$. Then, if a minimizer of

$$f(\mathbb{A}, \beta, \mathbb{C}, \delta) := \mathbb{E} \left(\sum_{t=1}^T \langle \ell_t, \sum_{i=1}^{t-1} (\mathbb{A} \ell_i \ell_i^\top \beta + \mathbb{C} \ell_i + \delta) \rangle - \inf_{\pi \in \Pi} \left\langle \sum_{t=1}^T \ell_t, \pi \right\rangle \right)^2$$

can be expressed as $\mathbb{A} = V, \beta = K^\top (Qc + q_c), \mathbb{C} = V k_c^\top (Qc + q_c) + v_c (Qc + q_c)^\top K, \beta = v_c k_c^\top (Qc + q_c)$, then we can conclude that the corresponding V, Q, K, v_c, q_c, k_c are also a minimizer of

$$\mathbb{E} \left(\sum_{t=1}^T \langle \ell_t, g(Z_{t-1}) \rangle - \inf_{\pi \in \Pi} \left\langle \sum_{t=1}^T \ell_t, \pi \right\rangle \right)^2,$$

since the corresponding V, Q, K, v_c, q_c, k_c constitute a minimizer among a larger class. Now, since $\Pi = B(\mathbf{0}_d, R_\Pi, \|\cdot\|)$, we can rewrite f as

$$f(\mathbb{A}, \beta, \mathbb{C}, \delta) = \mathbb{E} \left(\sum_{t=1}^T \langle \ell_t, \sum_{i=1}^{t-1} (\mathbb{A} \ell_i \ell_i^\top \beta + \mathbb{C} \ell_i + \delta) \rangle + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2. \quad (\text{E.10})$$

Step 1. Finding condition for $\frac{\partial f}{\partial \delta} = 0$.

Due to the Leibniz rule, if we calculate the partial derivative of Equation (E.10) w.r.t. δ , we have

$$\begin{aligned} \frac{\partial f(\mathbb{A}, \beta, \mathbb{C}, \delta)}{\partial \delta} &= \frac{\partial}{\partial \delta} \mathbb{E} \left(\sum_{t=1}^T \langle \ell_t, \sum_{i=1}^{t-1} (\mathbb{A} \ell_i \ell_i^\top \beta + \mathbb{C} \ell_i + \delta) \rangle + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2 \\ &= \mathbb{E} \frac{\partial}{\partial \delta} \left(\sum_{t=1}^T \langle \ell_t, \sum_{i=1}^{t-1} (\mathbb{A} \ell_i \ell_i^\top \beta + \mathbb{C} \ell_i + \delta) \rangle + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2 \\ &= \mathbb{E} \sum_{t=1}^T \ell_t \left(\sum_{t=1}^T \sum_{i=1}^{t-1} (t-1) \ell_t^\top (\mathbb{A} \ell_i \ell_i^\top \beta + \mathbb{C} \ell_i + \delta) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right). \end{aligned} \quad (\text{E.11})$$

Since the expectation of either odd-order polynomial or even-order polynomial times $\|\cdot\|_2$ is 0, due to that ℓ_t follows a symmetric distribution, we have

$$\mathbb{E} \sum_{t=1}^T (t-1) \ell_t R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 = 0, \quad \mathbb{E} \sum_{t=1}^T (t-1) \ell_t \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i = 0.$$

Now, we calculate

$$\begin{aligned} \mathbb{E} \sum_{t=1}^T (t-1) \ell_t \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{A} \ell_i \ell_i^\top \beta &= \mathbb{E} \sum_{t_1=1}^T \sum_{t=1}^T \sum_{i=1}^{t-1} (t_1-1) \ell_{t_1} \ell_t^\top \mathbb{A} \ell_i \ell_i^\top \beta \\ &\stackrel{(i)}{=} \mathbb{E} \sum_{t=1}^T \sum_{i=1}^{t-1} (t-1) \ell_t \ell_t^\top \mathbb{A} \ell_i \ell_i^\top \beta = \mathbb{E} \sum_{t=1}^T (t-1)^2 \ell_t \ell_t^\top \mathbb{A} \Sigma \beta = \frac{1}{6} T(2T^2 - 3T + 1) \Sigma \mathbb{A} \Sigma \beta, \end{aligned}$$

where (i) holds since if $t_1 \neq t$, due to the independence of ℓ_t, ℓ_{t_1} , we can use $\mathbb{E} \ell_t = 0$. Lastly,

$$\mathbb{E} \sum_{t=1}^T (t-1) \ell_t \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \delta = \mathbb{E} \sum_{t_1=1}^T \sum_{t=1}^T (t_1-1)(t-1) \ell_{t_1} \ell_t^\top \delta = \frac{1}{6} T(2T^2 - 3T + 1) \Sigma \delta.$$

Plugging the above equations into Equation (E.11), we have

$$\frac{\partial f(\mathbb{A}, \beta, \mathbb{C}, \delta)}{\partial \delta} = \frac{1}{6} T(2T^2 - 3T + 1) (\Sigma \mathbb{A} \Sigma \beta + \Sigma \delta).$$

Due to the optimality condition, we have

$$\mathbb{A} \Sigma \beta + \delta = 0. \tag{E.12}$$

Step 2. Plugging the optimality condition for $\frac{\partial f}{\partial \delta}$ into Equation (E.10).

Plugging Equation (E.12) to Equation (E.10), f can be written as

$$\begin{aligned} f(\mathbb{A}, \beta, \mathbb{C}, -\mathbb{A} \Sigma \beta) &= \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top (\mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta + \mathbb{C} \ell_i) + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2 \\ &= \underbrace{\mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right)^2}_{(i)} + \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \right)^2 + \mathbb{E} \left(R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)^2 \\ &\quad + \underbrace{2 \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right) \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \right)}_{(ii)} \\ &\quad + \underbrace{2 \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right) \left(R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right)}_{(iii)} \\ &\quad + 2 \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \right) \left(R_\Pi \left\| \sum_{t=1}^T \ell_t \right\|_2 \right). \end{aligned}$$

For the part (i), we have

$$\begin{aligned}
\mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right)^2 &= \mathbb{E} \left[\sum_{t=1}^T \sum_{i_1=1}^{t_1-1} \sum_{t=1}^T \sum_{i=1}^{t-1} \beta^\top (\ell_{i_1} \ell_{i_1}^\top - \Sigma) \mathbb{A}^\top \ell_{t_1} \ell_t^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right] \\
&\stackrel{(1)}{=} \mathbb{E} \left[\sum_{t=1}^T \sum_{i_1=1}^{t-1} \sum_{i=1}^{t-1} \beta^\top (\ell_{i_1} \ell_{i_1}^\top - \Sigma) \mathbb{A}^\top \ell_{t_1} \ell_t^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right] \\
&\stackrel{(2)}{=} \mathbb{E} \left[\sum_{t=1}^T \sum_{i=1}^{t-1} \beta^\top (\ell_i \ell_i^\top - \Sigma) \mathbb{A}^\top \ell_i \ell_i^\top \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta \right] \\
&= \frac{(T-1)T}{2} \beta^\top \mathbb{E} \left[(\ell_i \ell_i^\top - \Sigma) \mathbb{A}^\top \Sigma \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \right] \beta \\
&= \frac{(T-1)T}{2} \beta^\top \mathbb{E} \left[(\sqrt{\Sigma} \mathbb{A}(\ell_i \ell_i^\top - \Sigma))^\top (\sqrt{\Sigma} \mathbb{A}(\ell_i \ell_i^\top - \Sigma)) \right] \beta.
\end{aligned} \tag{E.13}$$

Here, (1) holds because if $t_1 \neq t$, we know that $\mathbb{E} \ell_{t_1} = \mathbb{E} \ell_t = 0$, and they are independent, and (2) holds because if $i_1 \neq i$, we can calculate $\mathbb{E}(\ell_{i_1} \ell_{i_1}^\top - \Sigma) = \mathbf{0}_{d \times d}$. In addition, we can easily check that (ii) and (iii) are 0 as they are polynomials of odd degrees and we have $Z \stackrel{d}{=} -Z$. Note that Equation (E.13) is minimized when $\mathbb{P}(\sqrt{\Sigma} \mathbb{A}(\ell_i \ell_i^\top - \Sigma) \beta = \mathbf{0}_d) = 1$.

If $\mathbb{A} \neq \mathbf{0}_{d \times d}$, suppose that the singular value decomposition of $A = U \Lambda V$ yields that Λ is a diagonal matrix whose first diagonal element is non-zero, and U, V are orthogonal matrices. Then, we want to find β that $\sqrt{\Sigma} U \Lambda V(\ell_i \ell_i^\top - \Sigma) \beta = \mathbf{0}_d$ for any ℓ_i such that $p(\ell_i) \neq 0$, where p indicates the probability density function of loss vectors. Since Σ and U are invertible, we only need to consider $\Lambda V(\ell_i \ell_i^\top - \Sigma) \beta = \mathbf{0}_d$. Since Λ 's first diagonal component is non-zero, we will consider equation $e_1^\top \Lambda V(\ell_i \ell_i^\top - \Sigma) \beta = 0$. This is equivalent to $V_1(\ell_i \ell_i^\top - \Sigma) \beta = 0$, where V_1 is the first row of V , and is a non-zero vector.

Now, we will generally consider $a_{x,y}(v) := vv^\top x - y$ where $x, y, v \in \mathbb{R}^d$ and $a_{x,y} : B(\mathbf{0}_d, 2\epsilon_1, \|\cdot\|) \rightarrow \mathbb{R}^d$ function. Then, we can check that the Jacobian of $a_{x,y}(v)$ is $vx^\top + (v \cdot x)I$, and we can find that the determinant of the Jacobian is nonzero when $v = \epsilon_1 x$ if $x \neq \mathbf{0}_d$. Therefore, the volume of $(V_1(\ell_i \ell_i^\top - \Sigma))$ for $\ell_i \in B(\mathbf{0}_d, c_z, \|\cdot\|)$ is greater than the volume of $(V_1(vv^\top - \Sigma))$ for $v \in B(\epsilon_1 V_1^\top, \epsilon_2, \|\cdot\|)$, where c_z is a constant such that $B(\mathbf{0}_d, c_z, \|\cdot\|) \subseteq \text{supp}(Z)$, and $\epsilon_1, \epsilon_2 > 0$ satisfy that $\epsilon_1 |V_1| + \epsilon_2 < c_z$. Here, we define $\epsilon_2 > 0$ sufficiently small so that the determinant of Jacobian $(vv^\top V_1^\top - \Sigma V_1^\top) > 0$ for $v \in B(\epsilon_1 V_1^\top, \epsilon_2, \|\cdot\|)$, and $v \rightarrow vv^\top V_1^\top - \Sigma V_1^\top$ is a one-to-one correspondence, by inverse function theorem. Therefore, the volume of $(V_1(vv^\top - \Sigma))$ for $v \in B(\epsilon_1 V_1^\top, \epsilon_2, \|\cdot\|)$ can be calculated as

$$\begin{aligned}
&[\text{Volume}(V_1(vv^\top - \Sigma)) \text{ for } v \in B(\epsilon_1 V_1^\top, \epsilon_2, \|\cdot\|)] \\
&= \int_{v \in B(\epsilon_1 V_1^\top, \epsilon_2, \|\cdot\|)} |\det(\text{Jacobian}(V_1(vv^\top - \Sigma)))| dv > 0.
\end{aligned}$$

Therefore, $\text{Volume}(V_1(vv^\top - \Sigma))$ where $v \in B(\epsilon_1 V_1^\top, \epsilon_2, \|\cdot\|)$ is non-zero, so that we can find d loss vectors $\{\ell_i\}_{i \in [d]}$ such that the vectors $\{V_1(\ell_i \ell_i^\top - \Sigma)\}_{i \in [d]}$ are linearly independent. Hence, if we want to minimize Equation (E.13), either $A = \mathbf{0}_{d \times d}$ or $\beta = \mathbf{0}_d$ should hold. In both cases, Equation (E.9) can be re-written as

$$g(Z_t; \mathbb{A}, \beta, \mathbb{C}, \delta) := \sum_{i=1}^t \mathbb{C} \ell_i,$$

and this is covered by the original parametrization (Equation (E.8)) with $K^\top(Qc + q_c) = v_c = \mathbf{0}_d$.

Step 3. Calculating $\frac{\partial f}{\partial \mathbb{C}}$.

Now, we optimize over \mathbb{C} , by minimizing the following objective:

$$\begin{aligned}
f(\mathbb{C}) &:= \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i + R_\Pi \left\| \sum_{t=1}^T \ell_t \right\| \right)^2 \\
&= \underbrace{\mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \right)^2}_{(i)} + 2 \mathbb{E} \left(\left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \right) R_\Pi \left\| \sum_{t=1}^T \ell_t \right\| \right) + \mathbb{E} \left(R_\Pi \left\| \sum_{t=1}^T \ell_t \right\| \right)^2 \\
&= \frac{T(T-1)}{2} \text{Tr}(\mathbb{C}^\top \Sigma \mathbb{C} \Sigma) + 2 \mathbb{E} \left(B \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \left\| \sum_{j=1}^T \ell_j \right\| \right) + \mathbb{E} \left(R_\Pi \left\| \sum_{t=1}^T \ell_t \right\| \right)^2.
\end{aligned}$$

Here, (i) can be calculated as follows:

$$\begin{aligned}
\mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C} \ell_i \right)^2 &= \mathbb{E} \left(\sum_{t_1=1}^T \sum_{i_1=1}^{t_1-1} \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_{t_1}^\top \mathbb{C}^\top \ell_{t_1} \ell_i^\top \mathbb{C} \ell_i \right) \\
&\stackrel{(1)}{=} \mathbb{E} \left(\sum_{t=1}^T \sum_{i_1=1}^{t-1} \sum_{i=1}^{t-1} \ell_{t_1}^\top \mathbb{C}^\top \ell_{i_1} \ell_i^\top \mathbb{C} \ell_i \right) = \mathbb{E} \left(\sum_{t=1}^T \sum_{i_1=1}^{t-1} \sum_{i=1}^{t-1} \ell_{t_1}^\top \mathbb{C}^\top \Sigma \mathbb{C} \ell_i \right) \\
&\stackrel{(2)}{=} \mathbb{E} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t^\top \mathbb{C}^\top \Sigma \mathbb{C} \ell_i \right) \stackrel{(3)}{=} \mathbb{E} \text{Tr} \left(\sum_{t=1}^T \sum_{i=1}^{t-1} \mathbb{C}^\top \Sigma \mathbb{C} \ell_i \ell_i^\top \right) = \frac{T(T-1)}{2} \text{Tr}(\mathbb{C}^\top \Sigma \mathbb{C} \Sigma),
\end{aligned}$$

since (1) holds because if $t_1 \neq t$, we already know that $\mathbb{E} \ell_t = \mathbb{E} \ell_{t_1} = 0$, (2) holds due to a similar reason, and (3) comes from $\text{Tr}(AB) = \text{Tr}(BA)$.

We calculate $\frac{\partial f(\mathbb{C})}{\partial \mathbb{C}}$:

$$\frac{\partial f(\mathbb{C})}{\partial \mathbb{C}} = T(T-1) \Sigma \mathbb{C} \Sigma + 2 R_\Pi \mathbb{E} \left(\left\| \sum_{j=1}^T \ell_j \right\| \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t \ell_i^\top \right).$$

Hence, the optimal $\mathbb{C} = -\frac{2R_\Pi}{T(T-1)} \Sigma^{-1} \mathbb{E} \left(\left\| \sum_{j=1}^T \ell_j \right\| \sum_{t=1}^T \sum_{i=1}^{t-1} \ell_t \ell_i^\top \right) \Sigma^{-1}$.

Now, we see that for the special case of $\Sigma = I$, we have $\mathbb{C} = -R_\Pi \mathbb{E} \left(\left\| \sum_{j=1}^T \ell_j \right\| \ell_t \ell_i^\top \right)$. If we calculate the (a, b) -coordinate of \mathbb{C} , we need to calculate

$$\mathbb{E}_\ell \left[\sqrt{\sum_{o=1}^d \left(\sum_{s=1}^T \ell_{so} \right)^2} \ell_{ia} \ell_{kb} \right].$$

If $a \neq b$, then since Z is symmetric, the term above becomes zero. Therefore, we only need to consider the case when $a = b$, which is $\mathbb{E}_\ell \left[\sqrt{\sum_{o=1}^d \left(\sum_{s=1}^T \ell_{so} \right)^2} \ell_{ia} \ell_{ka} \right]$, and it will be the same value for all $a \in [d]$ since ℓ_i 's coordinates are independent.

Now, we calculate the scale of $\mathbb{E}_\ell \left[\sqrt{\sum_{o=1}^d \left(\sum_{s=1}^T \ell_{so} \right)^2} \ell_{i1} \ell_{k1} \right]$. We have $Z := \frac{\sum_{o=1}^{d-1} \left(\sum_{s=1}^T \ell_{so} \right)^2}{T(d-1)} \xrightarrow{a.s.} 1$ as $d \rightarrow \infty$ (by the law of large numbers) and we define $W := \sum_{s \neq i, k} \ell_{s1} / \sqrt{T}$ which is independent of ℓ_{i1}

and ℓ_{k1} .

$$\begin{aligned}
\mathbb{E}_\ell \left[\sqrt{\sum_{o=1}^d \left(\sum_{s=1}^T \ell_{so} \right)^2} \ell_{i1} \ell_{k1} \right] &= \mathbb{E}_{Z,W,\ell_{i1},\ell_{k1}} \left[\sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} + \ell_{k1})^2} \ell_{i1} \ell_{k1} \right] \\
&= \mathbb{E}_{Z,W,\ell_{i1},\ell_{k1} \geq 0} \left[\sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} + \ell_{k1})^2} \ell_{i1} \ell_{k1} - \sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} - \ell_{k1})^2} \ell_{i1} \ell_{k1} \right] \\
&= \mathbb{E}_{Z,W,\ell_{i1},\ell_{k1} \geq 0} \left[\frac{4(\sqrt{T}W + \ell_{i1})\ell_{k1}}{\sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} + \ell_{k1})^2} + \sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} - \ell_{k1})^2}} \ell_{i1} \ell_{k1} \right].
\end{aligned}$$

Taking $d \rightarrow \infty$, we have

$$\frac{\sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} + \ell_{k1})^2} + \sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} - \ell_{k1})^2}}{2\sqrt{Td}} \xrightarrow{d} 1,$$

which further implies

$$\begin{aligned}
&\sqrt{Td} \frac{4(\sqrt{T}W + \ell_{i1})\ell_{k1}}{\sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} + \ell_{k1})^2} + \sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} - \ell_{k1})^2}} \ell_{i1} \ell_{k1} \\
&\xrightarrow{d} \sqrt{Td} \frac{4(\sqrt{T}W + \ell_{i1})\ell_{k1}}{2\sqrt{Td}} \ell_{i1} \ell_{k1} = 2(\sqrt{T}W + \ell_{i1})\ell_{i1} \ell_{k1}
\end{aligned}$$

as $d \rightarrow \infty$. Therefore,

$$\begin{aligned}
&\lim_{d \rightarrow \infty} \mathbb{E}_{Z,W,\ell_{i1},\ell_{k1} \geq 0} \left[\sqrt{Td} \frac{4(\sqrt{T}W + \ell_{i1})\ell_{k1}}{\sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} + \ell_{k1})^2} + \sqrt{T(d-1)Z + (\sqrt{T}W + \ell_{i1} - \ell_{k1})^2}} \ell_{i1} \ell_{k1} \right] \\
&= \mathbb{E}_{Z,W,\ell_{i1},\ell_{k1} \geq 0} [2(\sqrt{T}W + \ell_{i1})\ell_{i1} \ell_{k1}] = \mathbb{E}_{\ell_{i1},\ell_{k1} \geq 0} [\ell_{i1}^2 \ell_{k1}]
\end{aligned}$$

which is a constant. The last equality came from the fact that W, ℓ_{i1}, ℓ_{k1} are independent random variables, and expectation of ℓ_{i1} is zero. Therefore, the output of the single-layer linear self-attention provides us with online gradient descent with step-size $\Theta(R_\Pi/\sqrt{Td})$. In the online learning literature, we usually set the gradient step size as $\Theta(R_\Pi/\sqrt{Td})$ (Hazan, 2016, Theorem 3.1), which is consistent with the result above. \square

E.8 Empirical Validation of Theorem E.2 and Theorem 5.2

We now provide empirical validations for Theorem E.2 and Theorem 5.2. We provide the training details and the results as follows.

E.8.1 Empirical Validation of Theorem E.2

Our model architecture is defined as follows: the number of layers T is set to 30 and the dimensionality d to 32, with the loss vector ℓ_i 's distribution Z following a standard normal distribution $\mathcal{N}(0,1)$. During training, we conducted 40,000 epochs with a batch size of 512. We employed the Adam optimizer, setting the learning rate to 0.001. We initialized the value, query, and key vectors (v_c, q_c, k_c) as zero vectors.

Our empirical analysis aims to demonstrate that the optimized model inherently emulates online gradient descent. To illustrate this, we will focus on two key convergence properties: $K^\top Q$ approaching the zero matrix $\mathbf{O}_{d \times d}$ and V converging to $a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$, where a and b are constants in \mathbb{R} . The conditions $K^\top Q = \mathbf{O}_{d \times d}$ and $V = a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$ imply that the function $g(Z_t; V, Q, K) = \sum_{i=1}^t (b-a)\ell_i$, effectively emulating the process of an online gradient descent method. We repeated 10 times of the experiments. For verifying $K^\top Q = \mathbf{O}_{d \times d}$, we will measure Frobenius norm ($\|\cdot\|_F$) of $K^\top Q$. Also for measuring the closeness of V and $a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$, we will measure $\min_{a,b \in \mathbb{R}} \|V - (a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d})\|_F / b$. The results are demonstrated in the first plot of Figure E.1.

E.8.2 Empirical Validation of Theorem 5.2

We now focus on two key convergence properties: $K^\top(Q\mathbf{1}_d + q_c)$ approaching the zero vector $\mathbf{0}_d$ and V converging to $a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$, where a and b are constants in \mathbb{R} . The conditions $K^\top(Q\mathbf{1}_d + q_c) = \mathbf{0}_d$ and $V = a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$ imply that the function $g(Z_t; V, Q, K) = \sum_{i=1}^t (b-a)\ell_i$, effectively emulating the process of an online gradient descent method. We repeated 10 times. For verifying $K^\top(Q\mathbf{1}_d + q_c) = \mathbf{0}_d$, we will measure 2-norm of $K^\top(Q\mathbf{1}_d + q_c)$. Also for measuring the closeness of V and $a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$, we will measure $\min_{a,b \in \mathbb{R}} \|V - (a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d})\|_F / b$. The results are demonstrated in the second plot of Figure E.1.




Figure E.1: Empirical validation of Theorem E.2 (top), Theorem 5.2 (middle), and Conjecture 3 (bottom). The observed convergence in Theorem E.2 and Conjecture 3’s result suggests that configuration in Theorem E.2 and Conjecture 3 are not only the local optimal point, but it has the potential as being the global optimizer.

E.9 Discussions on the Production of FTRL with Entropy Regularization

Now, we will consider projecting a single-layer linear self-attention model into a constrained domain such as a simplex, which is more amenable to the Experts Problem setting. To this end, we consider the following parameterization by adding an additional *non-linear* structure for the single-layer linear self-attention:

$$g(Z_t; V, K, Q, v_c, k_c, q_c) = \text{Operator} \left(\sum_{i=1}^t (V\ell_i + v_c)((K\ell_i + k_c)^\top \cdot (Qc + q_c)) \right), \quad (\text{E.14})$$

where the Operator denotes projection to the convex set.

Conjecture 3. Assume $\Sigma = I$. Then, the configuration that $K^\top(Qc + q_c) = v_c = \mathbf{0}_d$ and $V = \tilde{\Omega} \left(-\frac{1}{\sqrt{nd}} \right) I_{d \times d}$ is a first-order stationary point of Equation (5.2) with $N = 1$ and $h(x) = x^2$ when LLM_θ is parameterized with Equation (E.14), Operator = Softmax, and $\Pi = \Delta(\mathcal{A})$. This configuration performs FTRL with an entropy regularizer which is a no-regret algorithm.

We provide an idea for proving the conjecture, together with its numerical validation. Also, we have observed in Figure E.1 that Theorem E.2 and Conjecture 3 might also be a global optimizer, as training results have provided the configuration that Theorem E.2 and Conjecture 3 have suggested. To be specific, we will consider

$$f(V, a, \beta, v_c) = \mathbb{E} \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right)^2$$

and will try to prove that $a = \mathbf{0}_d, v_c = v \mathbf{1}_d, V = kI$ is a first-order stationary point.

Step 1. Calculating $\frac{\partial f}{\partial v_c}$.

We use the following formula: for $x \in [d]$ and $t \geq 2$, we have

$$\begin{aligned} & \frac{\partial}{\partial v_{cx}} \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} \\ &= \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \frac{\partial}{\partial v_{cx}} \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} \\ &= \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \sum_{i=1}^t (a^\top \ell_i \ell_i^\top e_x + \beta) \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} \\ &= t \beta \exp(v \beta) \exp(\beta k \sum_{i=1}^t \ell_{iy}), \end{aligned}$$

and for $t = 1$, $\frac{\partial}{\partial v_{cx}} \exp(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta)) \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} = 0$, so we can use the same formula with $t \geq 2$. Thus, we have

$$\begin{aligned} & \frac{\partial}{\partial v_{cx}} \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right) \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} \\ &= \beta \exp(v \beta) \\ & \quad \sum_{t=1}^T t \sum_{s=1}^d \ell_{ts} \frac{\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) - \sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy})}{(\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} \beta V \ell_j))^2} \\ &= 0. \end{aligned}$$

Therefore,

$$\begin{aligned} & \frac{\partial f(V, a, \beta, v_c)}{\partial v_{cx}} \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} \\ &= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\ & \quad \left. \frac{\partial}{\partial v_{cx}} \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right] \Big|_{a=\mathbf{0}_d, v_c=v \mathbf{1}_d, V=kI} \\ &= 0. \end{aligned}$$

Step 2. Calculating $\frac{\partial f}{\partial V}$.

The following formula will be used for calculating $\left. \frac{\partial f}{\partial V} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$: for $r, c \in [d]$, we have

$$\begin{aligned}
& \left. \frac{\partial}{\partial V_{rc}} \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \left. \frac{\partial}{\partial V_{rc}} \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \exp \left(\sum_{i=1}^t k \beta \ell_{iy} + v \beta \right) \sum_{i=1}^t \beta \mathbf{1}(y=r) \ell_{ic}.
\end{aligned}$$

Therefore,

$$\begin{aligned}
& \left. \frac{\partial f(V, a, \beta, v_c)}{\partial V_{rc}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp \left(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta) \right)}{\sum_{y=1}^d \exp \left(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta) \right)} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\
& \quad \left. \frac{\partial}{\partial V_{rc}} \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp \left(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta) \right)}{\sum_{y=1}^d \exp \left(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta) \right)} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right] \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp \left(\sum_{j=1}^{t-1} \beta k \ell_{js} + v \beta \right)}{\sum_{y=1}^d \exp \left(\sum_{j=1}^{t-1} \beta V \ell_{jy} + v \beta \right)} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\
& \quad \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\sum_{j=1}^{t-1} \beta \mathbf{1}(s=r) \ell_{jc} \exp \left(\sum_{j=1}^{t-1} \beta k \ell_{js} + v \beta \right) \sum_{y=1}^d \exp \left(\sum_{j=1}^{t-1} \beta k \ell_{jy} + v \beta \right)}{\left(\sum_{y=1}^d \exp \left(\sum_{j=1}^{t-1} \beta k \ell_{jy} + v \beta \right) \right)^2} \right. \\
& \quad \left. \left. - \sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp \left(\sum_{j=1}^{t-1} \beta k \ell_{js} + v \beta \right) \sum_{y=1}^d \left(\sum_{j=1}^{t-1} \beta \mathbf{1}(y=r) \ell_{jc} \exp \left(\sum_{j=1}^{t-1} \beta k \ell_{jy} + v \beta \right) \right)}{\left(\sum_{y=1}^d \exp \left(\sum_{j=1}^{t-1} \beta k \ell_{jy} + v \beta \right) \right)^2} \right) \right] \\
&= \beta \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp \left(\sum_{j=1}^{t-1} \beta k \ell_{js} \right)}{\sum_{y=1}^d \exp \left(\sum_{j=1}^{t-1} \beta V \ell_{jy} \right)} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\
& \quad \left(\underbrace{\frac{\sum_{t=1}^T \sum_{j=1}^{t-1} \sum_{y=1}^d \ell_{tr} \ell_{jc} \exp \left(\beta k \sum_{j=1}^{t-1} \ell_{jr} \right) \exp \left(\beta k \sum_{j=1}^{t-1} \ell_{jy} \right)}{\left(\sum_{y=1}^d \exp \left(\beta k \sum_{j=1}^{t-1} \ell_{jy} \right) \right)^2}}_{(i)} \right. \\
& \quad \left. \left. - \underbrace{\frac{\sum_{t=1}^T \sum_{j=1}^{t-1} \sum_{y=1}^d \ell_{ty} \ell_{jc} \exp \left(\beta k \sum_{j=1}^{t-1} \ell_{jr} \right) \exp \left(\beta k \sum_{j=1}^{t-1} \ell_{jy} \right)}{\left(\sum_{y=1}^d \exp \left(\beta k \sum_{j=1}^{t-1} \ell_{jy} \right) \right)^2}}_{(ii)} \right) \right].
\end{aligned}$$

We can observe the followings: 1) if $r_1 \neq c_1$ and $r_2 \neq c_2$, $\left. \frac{\partial f}{\partial V_{r_1 c_1}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} = \left. \frac{\partial f}{\partial V_{r_2 c_2}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$

holds, and 2) $\left. \frac{\partial f}{\partial V_{r_1 r_1}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} = \left. \frac{\partial f}{\partial V_{r_2 r_2}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$.

Step 3. Calculating $\frac{\partial f}{\partial \beta}$.

The following formula will be used for calculating $\left. \frac{\partial f}{\partial \beta} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$:

$$\begin{aligned}
& \left. \frac{\partial}{\partial \beta} \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \left. \frac{\partial}{\partial \beta} \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= tv\beta \exp \left(\sum_{i=1}^t k\beta \ell_{iy} + v\beta \right).
\end{aligned}$$

Further, we have

$$\begin{aligned}
& \left. \frac{\partial}{\partial \beta} \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= v\beta \exp(v\beta) \\
& \quad \sum_{t=1}^T t \sum_{s=1}^d \ell_{ts} \frac{\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) - \sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy})}{\left(\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} \beta V \ell_j) \right)^2} \\
&= 0.
\end{aligned}$$

Step 4. Calculating $\frac{\partial f}{\partial a}$.

Note that

$$\begin{aligned}
& \left. \frac{\partial}{\partial a_x} \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \left. \frac{\partial}{\partial a_x} \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \exp \left(e_y^\top \sum_{i=1}^t (V \ell_i \ell_i^\top a + (\beta V + v_c a^\top) \ell_i + v_c \beta) \right) \sum_{i=1}^t \left(e_y^\top V \ell_i \ell_i^\top e_x + e_y^\top v_c \ell_i^\top e_x \right) \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \exp \left(\sum_{i=1}^t \beta k \ell_{iy} + v\beta \right) \sum_{i=1}^t (k \ell_{iy} \ell_{ix} + v \ell_{ix}).
\end{aligned}$$

Therefore,

$$\begin{aligned}
& \left. \frac{\partial f(V, a, \beta, v_c)}{\partial a_x} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\
&\quad \left. \frac{\partial}{\partial a_x} \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(e_s^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))}{\sum_{y=1}^d \exp(e_y^\top \sum_{j=1}^{t-1} (V \ell_j \ell_j^\top a + (\beta V + v_c a^\top) \ell_j + v_c \beta))} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right] \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} \\
&= \mathbb{E} \left[\left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(\sum_{j=1}^{t-1} \beta k \ell_{js})}{\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy})} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\
&\quad \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\sum_{j=1}^{t-1} (k \ell_{js} \ell_{jx} + v \ell_{jx}) \exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) \sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy})}{\left(\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \right)^2} \right. \\
&\quad \left. \left. - \sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) \sum_{y=1}^d \left(\sum_{j=1}^{t-1} (k \ell_{jy} \ell_{jx} + v \ell_{jx}) \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \right)}{\left(\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \right)^2} \right) \right] \\
&= \mathbb{E} \left[k \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(\sum_{j=1}^{t-1} \beta k \ell_{js})}{\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy})} - \min_s \sum_{t=1}^T \ell_{ts} \right) \right. \\
&\quad \left(\sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\sum_{j=1}^{t-1} \ell_{js} \ell_{jx} \exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) \sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy})}{\left(\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \right)^2} \right. \\
&\quad \left. \left. - \sum_{t=1}^T \sum_{s=1}^d \ell_{ts} \frac{\exp(\sum_{j=1}^{t-1} \beta k \ell_{js}) \sum_{y=1}^d \left(\sum_{j=1}^{t-1} \ell_{jy} \ell_{jx} \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \right)}{\left(\sum_{y=1}^d \exp(\sum_{j=1}^{t-1} \beta k \ell_{jy}) \right)^2} \right) \right]
\end{aligned}$$

Note that the value does not depend on x , which means that $\left. \frac{\partial f}{\partial a} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} = \widetilde{c} \mathbf{1}_d$ for some constant \widetilde{c} .

E.9.1 Numerical Analysis of Step 2 and Step 4

In Steps 2 and 4 above, we were not able to show that a k whose value becomes zero exists. We hence provide some empirical evidence here. First, we attach the estimated $\left. \frac{\partial f}{\partial V_{rc}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ ($r \neq c$),

$\left. \frac{\partial f}{\partial V_{rr}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$, $\left. \frac{\partial f}{\partial a_x} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ and $\left. \frac{\partial f}{\partial a_x} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ graph with respect to k value when

$\ell_{ts} \sim \text{Unif}([0, 1])$ for all $t \in [T], s \in [d]$. While the graph of $\left. \frac{\partial f}{\partial V} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ is not stable, we can see

that k for $\left. \frac{\partial f}{\partial V_{rc}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} = 0$, $\left. \frac{\partial f}{\partial V_{rr}} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} = 0$ and $\left. \frac{\partial f}{\partial a_x} \right|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI} = 0$ is very similar in

Figure E.2. We used the Monte Carlo estimation of 1,000,000 times.

E.9.2 Empirical Validation

Our model architecture is defined as follows: the number of layers T is set to 30 and the dimensionality d to 32, with the loss vector l_i 's distribution Z following a standard normal distribution

$\mathcal{N}(0,1)$. During training, we conducted 40,000 epochs with a batch size of 512. We employed the Adam optimizer, setting the learning rate to 0.001. We focus on two key convergence properties: $K^\top(Q\mathbf{1} + q_c)$ approaching the zero vector $\mathbf{0}_d$ and V converging to $a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$, where a and b are constants in \mathbb{R} . The conditions $K^\top(Q\mathbf{1} + q_c) = \mathbf{0}_d$ and $V = a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$ imply that the function $g(Z_t; V, Q, K) = \sum_{i=1}^t (b-a)l_i$, effectively emulating the process of an online gradient descent method. We repeated 10 times. For verifying $K^\top(Q\mathbf{1} + q_c) = \mathbf{0}_d$, we will measure 2-norm of $K^\top(Q\mathbf{1} + q_c)$. Also for measuring the closeness of V and $a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d}$, we will measure $\min_{a,b \in \mathbb{R}} \|V - (a\mathbf{1}_d\mathbf{1}_d^\top + bI_{d \times d})\|_{2,2}/b$. The results are demonstrated in the third plot of Figure E.1.

E.10 Comparison with In-Context-Learning Analyses in Supervised Learning

The very recent studies by Ahn et al. (2023); Zhang et al. (2023a); Mahankali et al. (2023) have demonstrated that if $Z_t = ((x_1, y_1), \dots, (x_t, y_t), (x_{t+1}, 0))$ and the “instruction tuning” loss (i.e., $\mathbb{E}[\|\widehat{y}_{t+1} - y_{t+1}\|^2]$) is being minimized with a single-layer linear self-attention model, then a global optimizer among single-layer linear self-attention models yields the output $\widehat{y}_{n+1} = \eta \sum_{i=1}^n y_i x_i^\top x_{n+1}$. This output can be interpreted as a *gradient descent* algorithm, indicating that a single-layer linear self-attention model **implicitly** performs gradient descent. However, in the online learning setting where there are no y -labels, such an implicit gradient descent update-rule is hard to define. Compared to the previous studies, our global optimizer among single-layer linear self-attention models is an *explicit* and *online* gradient descent update for online learning. With a different loss (regret-loss v.s. instruction-tuning-loss), the techniques to obtain the seemingly similar results are also fundamentally different.

E.11 Training Details of Section 5.4

We provide the training details of Section 5.4. For the multi-layer Transformer training, we used 4 layers, 1 head Transformer. For both single-layer and multi-layer, we employed the Adam optimizer, setting the learning rate to 0.001. During training, we conducted 2,000 epochs with a batch size 512. Moreover, when we trained for the loss sequences with the predictable trend, we used 4 layers, 1 head Transformer. For both single-layer and multi-layer, we employed the Adam optimizer, setting the learning rate to 0.001. During training, we conducted 9,000 epochs with a batch size of 512.

E.12 Ablation Study on Training Equation (5.2)

In this section, we provide an ablation study that changes N and k in Equation (5.2). To be specific, we will set $N = 1, 2, 4$, $f(x, k) = \max(x, 0)^k$, $h(x) = \max(x, 0)^2$, and $k = 1, 2$. For the multi-layer Transformer training, we used 4 layers and 1 head Transformer. For both single-layer and multi-layer, we employed the Adam optimizer, setting the learning rate to 0.001. During training, we conducted 2,000 epochs with a batch size of 512. We experimented on the randomly generated loss sequences. Especially, we used the uniform loss sequence ($\ell_t \sim \text{Unif}([0, 10]^2)$), with the results in Figure E.3 and Figure E.4; and the Gaussian loss sequence ($\ell_t \sim \mathcal{N}(5 \cdot \mathbf{1}_2, I)$), with the results in Figure E.5 and Figure E.6.

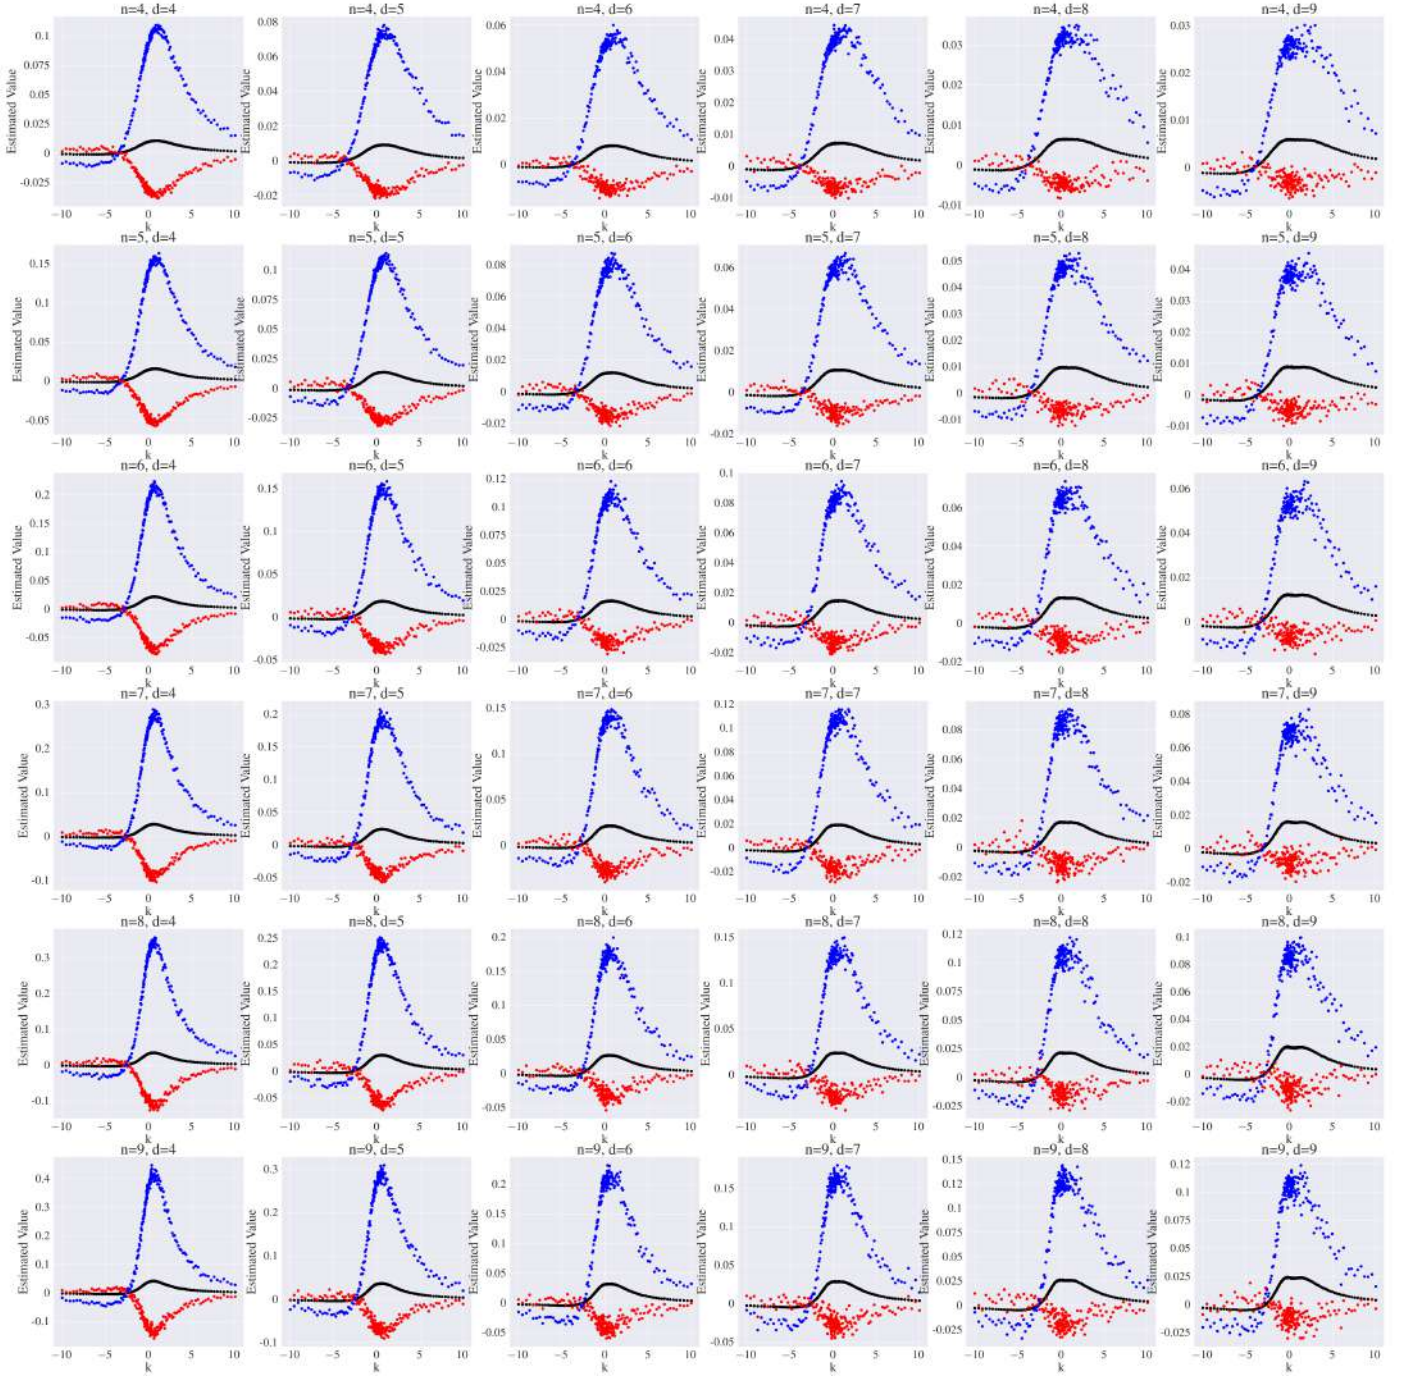


Figure E.2: Calculation of $20 \frac{\partial f}{\partial V_{rc}} \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ ($r \neq c$)(red), $20 \frac{\partial f}{\partial V_{rr}} \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ (blue), and $\frac{\partial f}{\partial a_x} \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ (black). We experimented with $n \in [4, 9]$ and $d \in [4, 9]$. The figure might indicate that βk that makes the derivative zero of $\frac{\partial f}{\partial V_{rc}} \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ ($r \neq c$), $\frac{\partial f}{\partial V_{rr}} \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$, and $\frac{\partial f}{\partial a_x} \Big|_{a=\mathbf{0}_d, v_c=v\mathbf{1}_d, V=kI}$ would coincide.

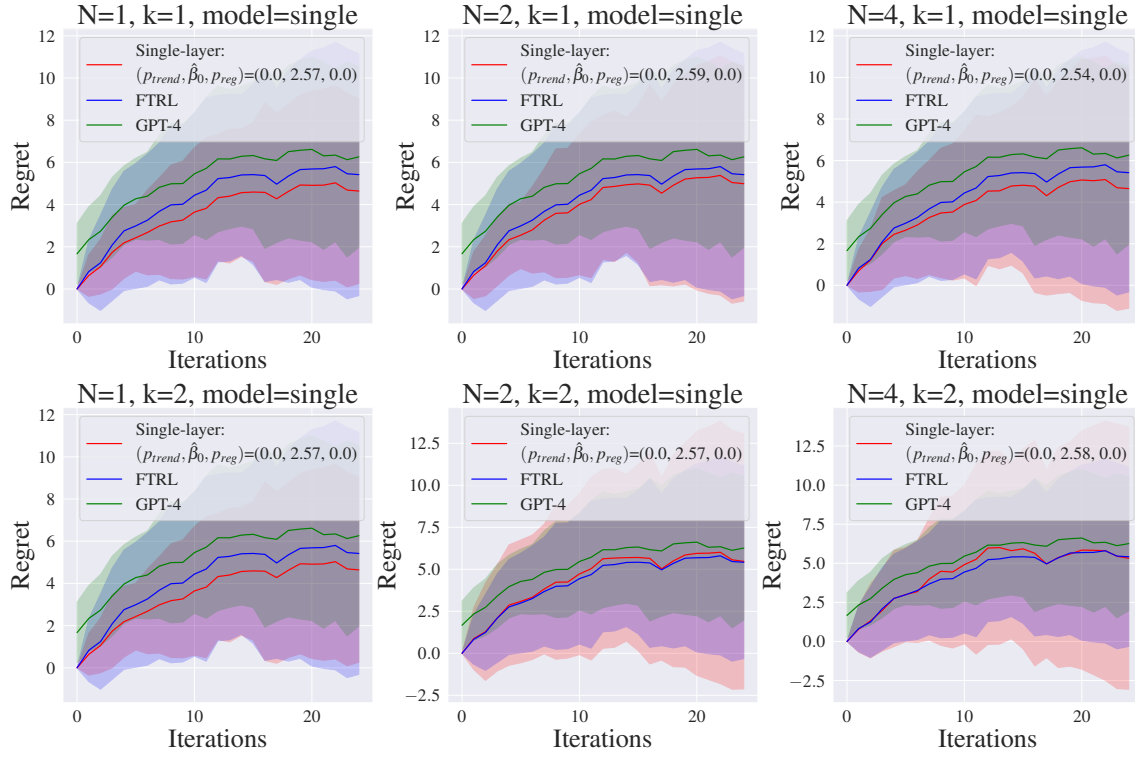


Figure E.3: Ablation study for the uniform loss sequence trained with single-layer self-attention layer and Softmax projection.

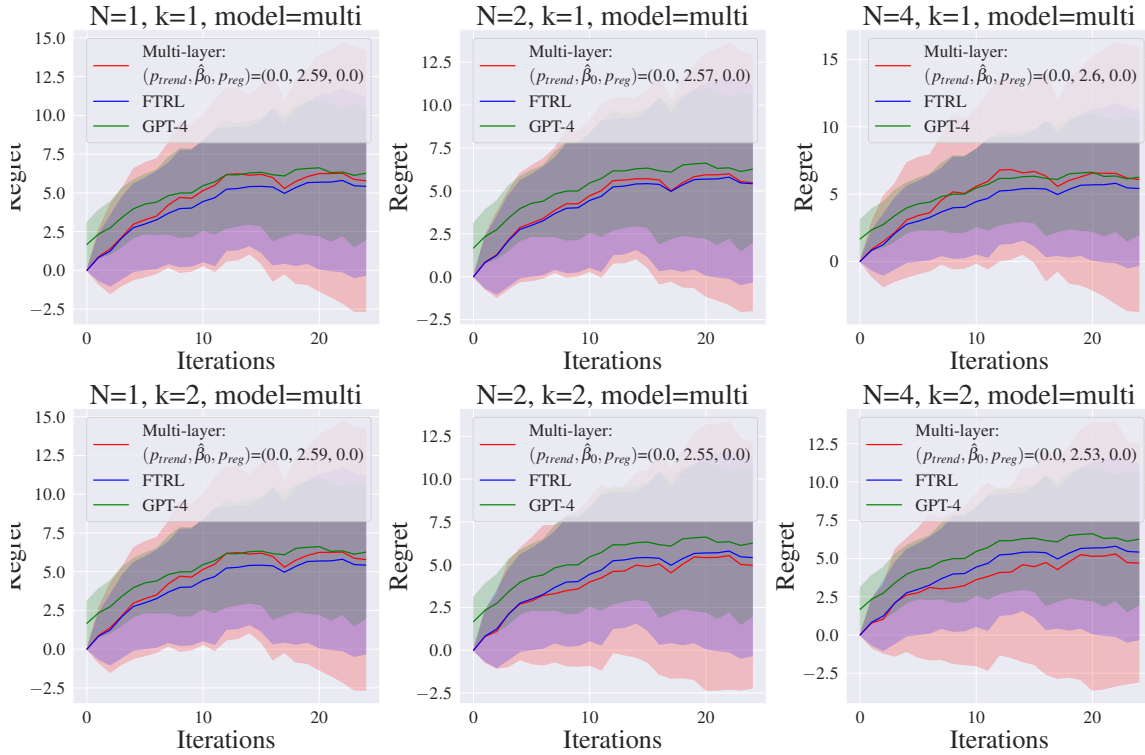


Figure E.4: Ablation study for the uniform loss sequence trained with multi-layer self-attention layer and Softmax projection.

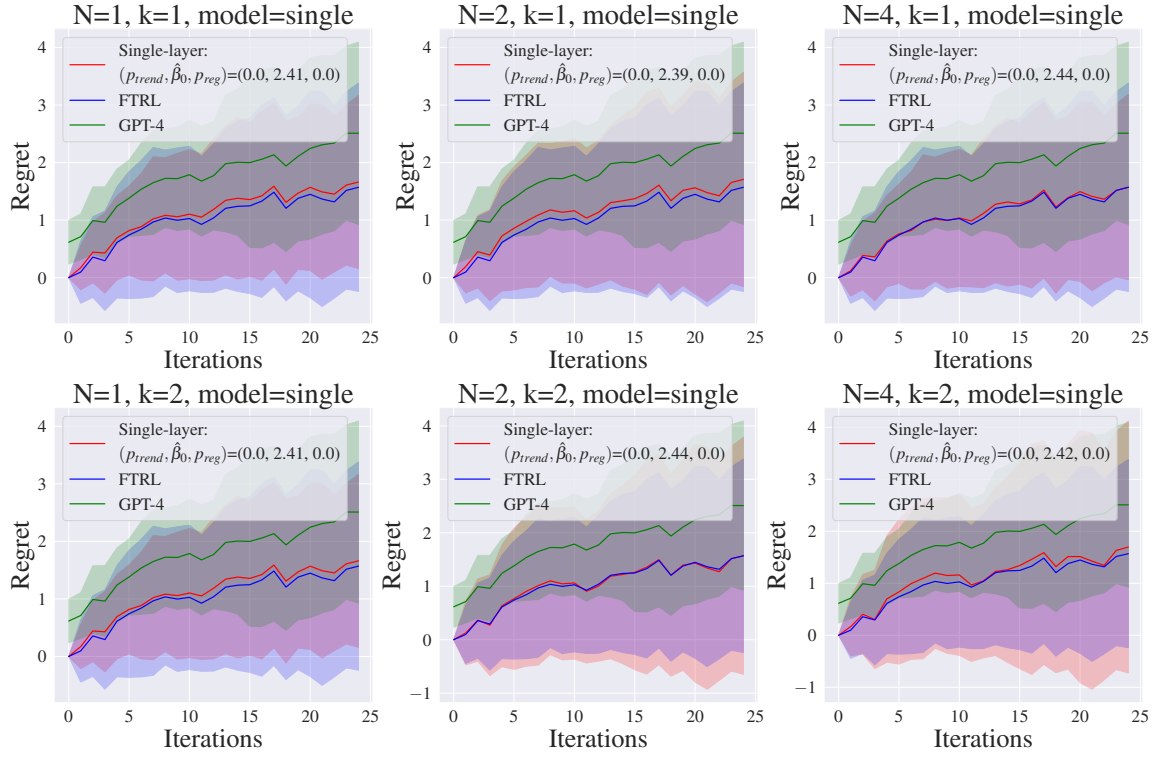


Figure E.5: Ablation study for the Gaussian loss sequence trained with single-layer self-attention layer and Softmax projection.

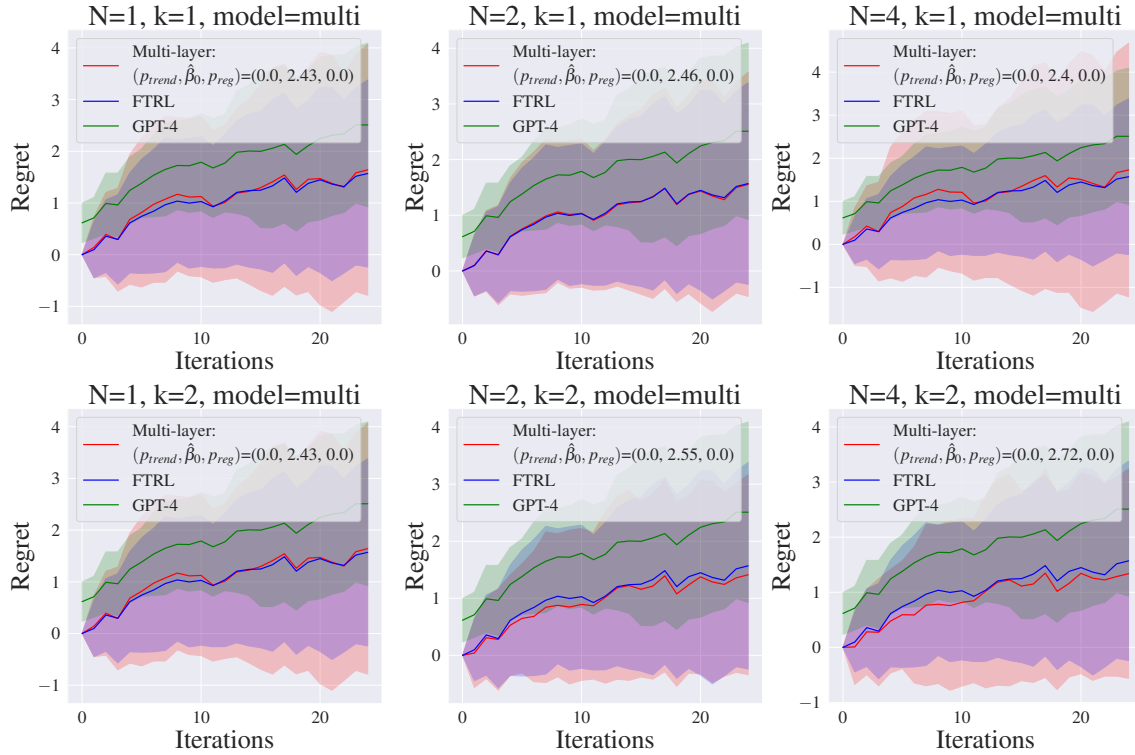


Figure E.6: Ablation study for the Gaussian loss sequence trained with single-layer self-attention layer and Softmax projection.